



异常控制流

100076202: 计算机系统导论

任课教师:

计卫星 宿红毅 张艳

原作者:

Randal E. **Bryant** and David R. O'Hallaron



**Carnegie
Mellon
University**



内容提纲

- **异常控制流** Exceptional Control Flow
- **异常** Exceptions
- **进程** Processes
- **进程控制** Process Control




控制流/Control Flow

■ 处理器只做一件事/Processors do only one thing:

- 从开机到关机，CPU只是读入和执行（解释）指令序列，每次一条
From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
- 这个序列就是CPU的控制流 This sequence is the CPU's *control flow* (or *flow of control*)

物理控制流/Physical control flow





改变控制流/Altering the Control Flow

- **目前的两种方式/Up to now: two mechanisms for changing control flow:**
 - 跳转分支指令/Jumps and branches
 - 调用和返回指令/Call and return

是对程序状态改变的响应/React to changes in *program state*
- **对有用系统来说还不够/Insufficient for a useful system:**

难以对系统状态的改变进行响应/Difficult to react to changes in *system state*

 - 从磁盘或者网络获取的数据到达/Data arrives from a disk or a network adapter
 - 指令除零/Instruction divides by zero
 - 用户按下了Ctrl-C/User hits Ctrl-C at the keyboard
 - 时钟超时触发/System timer expires
- **系统需要异常控制流处理机制/System needs mechanisms for “exceptional control flow”**



异常控制流 / Exceptional Control Flow

- 存在系统的每个层次 / Exists at all levels of a computer system
- 低层次机制 / Low level mechanisms
 - 1. 异常 / **Exceptions**
 - 为响应系统事件改变控制流（例如，系统状态改变） / Change in control flow in response to a system event (i.e., change in system state)
 - 硬件和OS软件组合实现 / Implemented using combination of hardware and OS software
- 高层次机制 / Higher level mechanisms
 - 2. 进程上下文切换 / **Process context switch**
 - 硬件时钟和OS软件实现 / Implemented by OS software and hardware timer
 - 3. 信号 **Signals**
 - OS软件实现 / Implemented by OS software
 - 4. 非局部跳转 / **Nonlocal jumps**: `setjmp()` and `longjmp()`
 - C运行时库实现 / Implemented by C runtime library



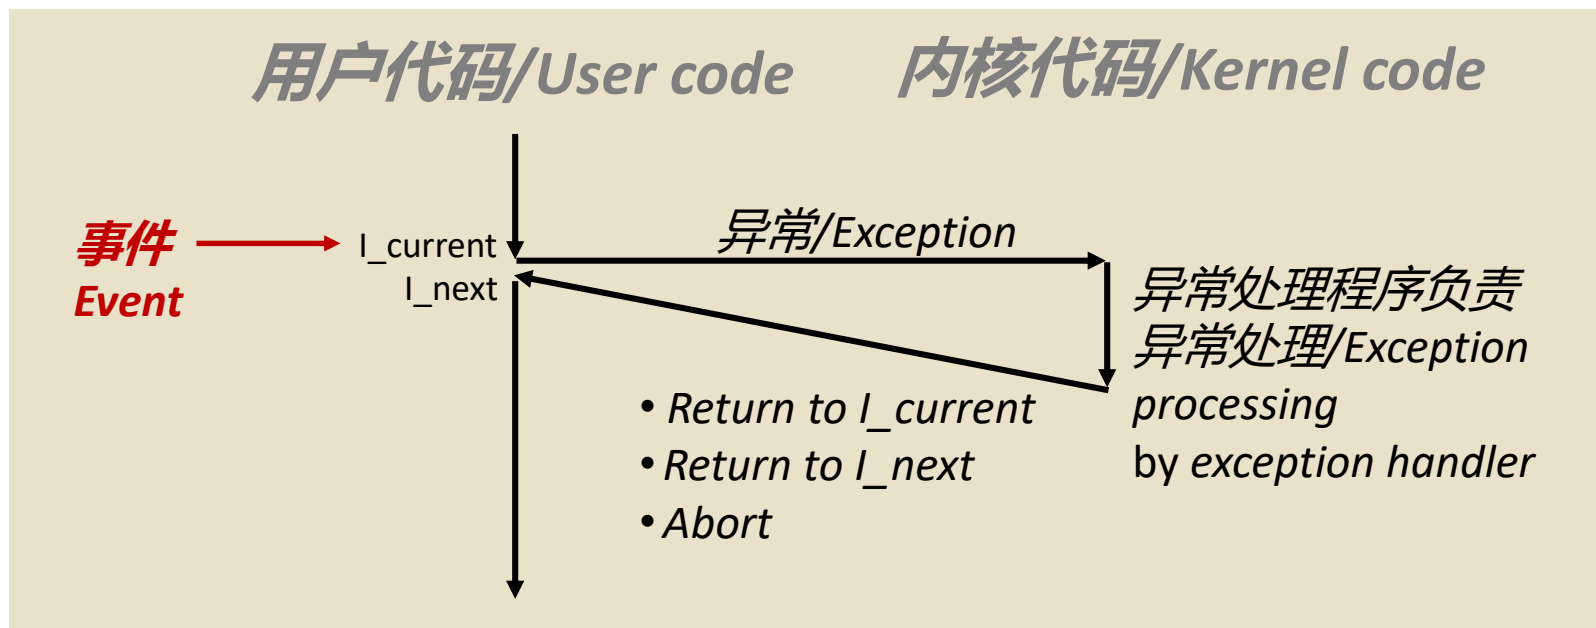
内容提纲

- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control



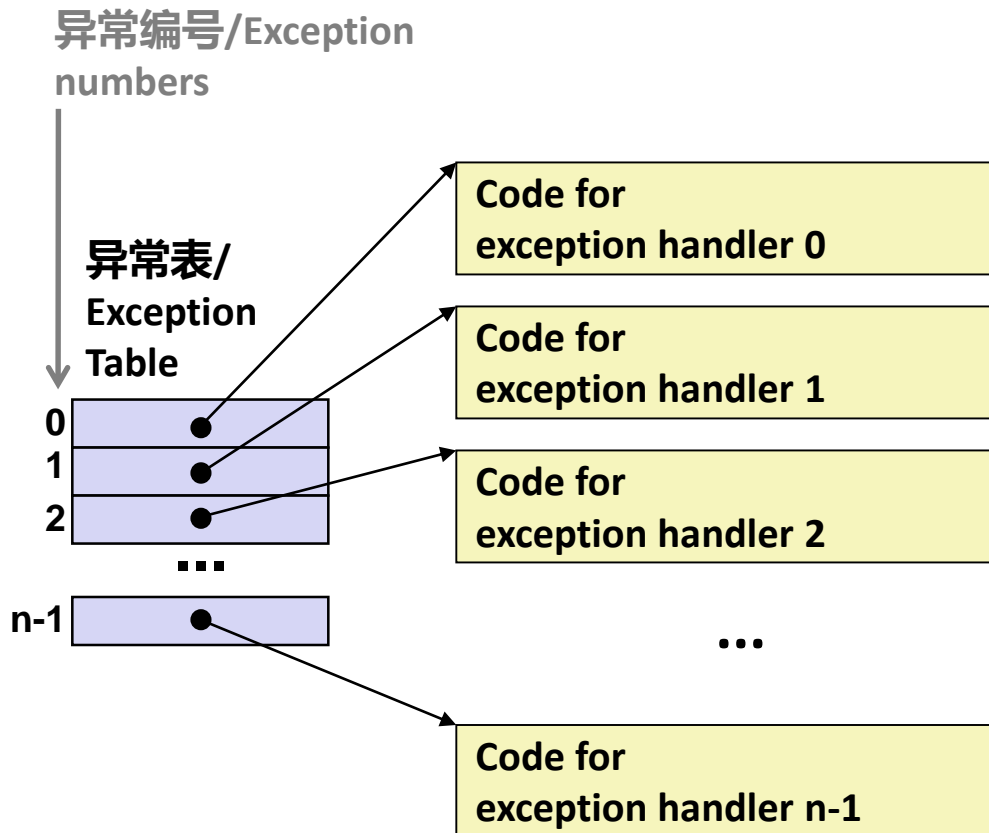
异常/Exceptions

- 异常是为了响应某些事件而将控制转移到OS内核（例如，处理器状态改变）
/An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
 - 内核是操作系统的内存驻留/Kernel is the memory-resident part of the OS
 - 事件举例：除零错误、算术溢出、缺页中断、I/O请求完成、Ctrl-C输入/Ctrl+C
Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C





异常表格 Exception Tables



- 每个事件类型有一个编号
Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Handler k is called each time exception k occurs



异步异常（中断/Asynchronous Exceptions (Interrupts)

- **由处理器外部事件引起/Caused by events external to the processor**
 - 通过处理器的中断管脚给出/Indicated by setting the processor's *interrupt pin*
 - 中断处理程序返回后执行下一条指令/Handler returns to “next” instruction
- **举例/Examples:**
 - 时钟中断/Timer interrupt
 - 大约几毫秒，外部时钟芯片触发/Every few ms, an external timer chip triggers an interrupt
 - 将控制权从用户切换到内核/Used by the kernel to take back control from user programs
 - 外部设备的I/O中断/ I/O interrupt from external device
 - 键盘输入Ctrl-C/Hitting Ctrl-C at the keyboard
 - 网络数据包到达/Arrival of a packet from a network
 - 磁盘数据到达/Arrival of data from a disk



同步异常/Synchronous Exceptions

- **指令执行导致的异常事件/Caused by events that occur as a result of executing an instruction:**
 - **陷入/陷阱 Traps**
 - 人为的/Intentional
 - 例如：系统调用、断点、特殊指令等/Examples: **system calls**, breakpoint traps, special instructions
 - 控制流返回下一条指令/Returns control to “next” instruction
 - **故障 Faults**
 - 不是有意的但是大概率可恢复/Unintentional but possibly recoverable
 - 例如：缺页异常、保护异常、浮点异常/Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - 重新执行或者终止执行/Either re-executes faulting (“current”) instruction or aborts
 - **终止 Aborts**
 - 非故意且不可恢复/Unintentional and unrecoverable
 - 例如：非法指令、校验错误、机器检查/Examples: illegal instruction, parity error, machine check
 - 终止当前程序执行/Aborts current program



系统调用/System Calls

- 每个x86-64系统调用都有一个唯一的ID编号/Each x86-64 system call has a unique ID number
- 例如：/Examples:

编号 /Number	名称 /Name	描述 /Description
0	read	读文件/Read file
1	write	写文件/Write file
2	open	打开文件/Open file
3	close	关闭文件/Close file
4	stat	获取文件信息/Get info about file
57	fork	创建进程/Create process
59	execve	执行程序/Execute a program
60	_exit	终止进程/Terminate process
62	kill	给进程发送信号/Send signal to process



系统调用举例：打开文件/System Call Example: Opening File


- 调用接口/User calls: **open(filename, options)**
- 调用__open函数，会执行系统调用指令syscall / Calls **__open** function, which invokes system call instruction **syscall**

```
0000000000e5d70 <__open>:
```

```
...  
e5d79:  b8 02 00 00 00      mov  $0x2,%eax  # open is syscall #2  
e5d7e:  0f 05               syscall          # Return value in %rax  
e5d80:  48 3d 01 f0 ff ff    cmp  $0xffffffffffffffff001,%rax  
...  
e5dfa:  c3                  retq
```

用户代码/
User code

内核代码/
Kernel code



- `%rax`包含了系统调用编号/`%rax` contains syscall number
- 其他的参数在/Other arguments in `%rdi, %rsi, %rdx, %r10, %r8, %r9`
- 返回值在`%rax`/Return value in `%rax`
- 负数表示出错对应`errno`/Negative value is an error corresponding to negative `errno`

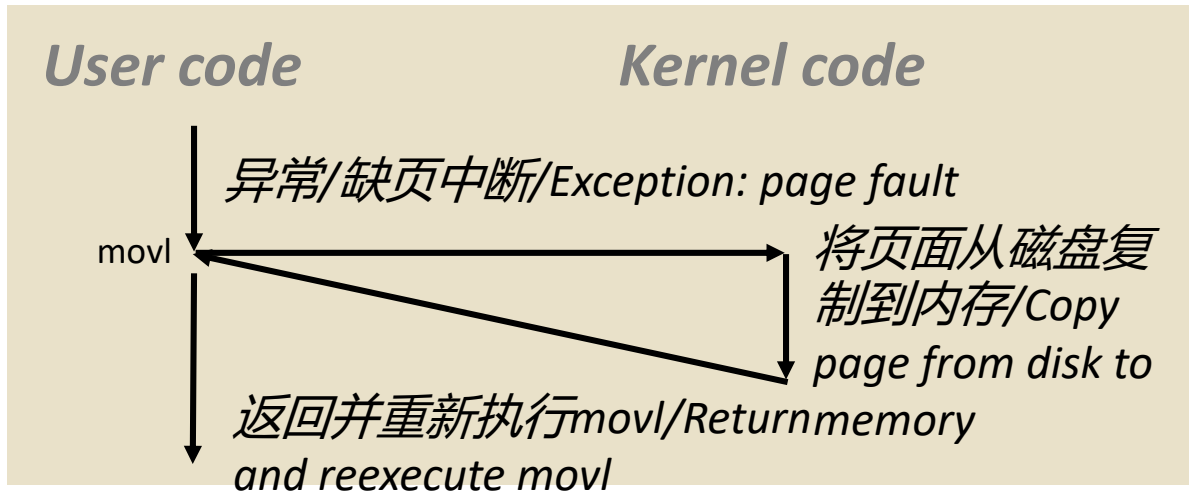


故障举例：缺页异常 / Fault Example: Page Fault

- 用户写内存/User writes to memory location
- 对应的用户内存页面在磁盘上/ That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

80483b7:	c7 05 10 9d 04 08 0d	movl	\$0xd,0x8049d10
----------	----------------------	------	-----------------

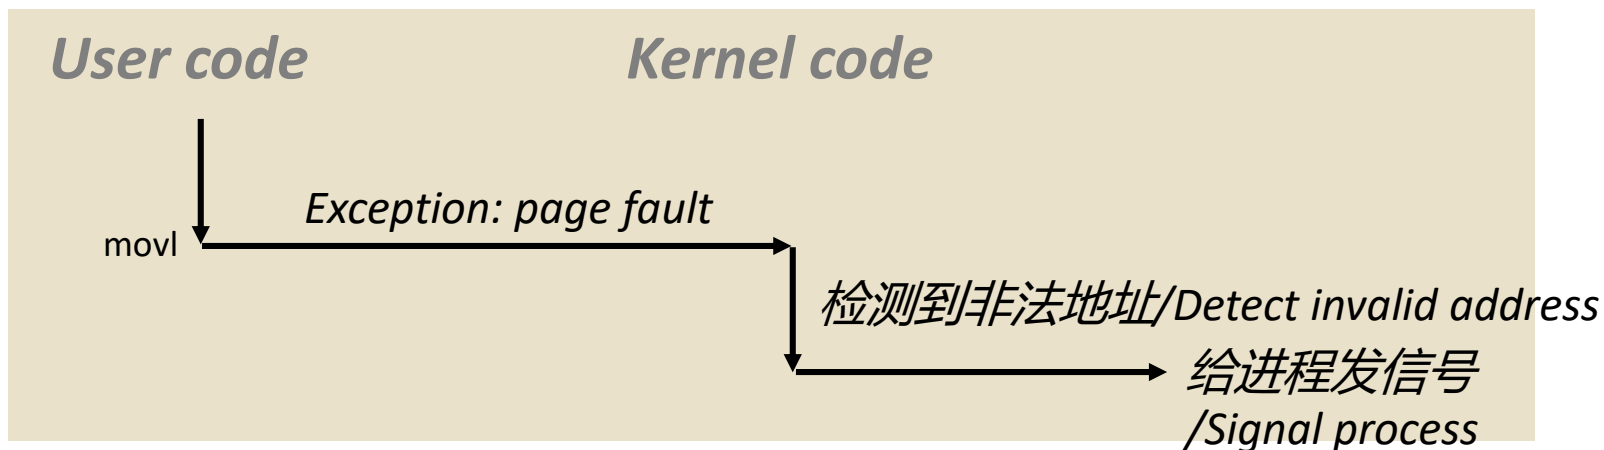




故障举例：非法内存引用/ Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360



- 发送SIGSEGV信号给用户进程/Sends **SIGSEGV** signal to user process
- 用户进程会异常退出/User process exits with “segmentation fault”




内容提纲

- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control




进程/Processes

- **定义：进程是程序的运行的实例/Definition: A *process* is an instance of a running program.**
 - 计算机科学最重要的概念之一/One of the most profound ideas in computer science
 - 与程序和处理器不同/Not the same as “program” or “processor”
- **进程为每个程序提供了两个关键抽象/Process provides each program with two key abstractions:**
 - **逻辑控制流/Logical control flow**
 - 每个程序看起来独占CPU/Each program seems to have exclusive use of the CPU
 - 内核支持的上下文切换/Provided by kernel mechanism called *context switching*
 - **私有地址空间/Private address space**
 - 每个程序看起来独占主存空间/Each program seems to have exclusive use of main memory.
 - 系统支持的虚拟内存/Provided by kernel mechanism called *virtual memory*





多进程幻象/Multiprocessing: The Illusion



■ 计算机同时运行很多进程/Computer runs many processes simultaneously

- 单个或多个用户的应用/Applications for one or more users
 - 网页浏览器、邮件客户端、编辑器/Web browsers, email clients, editors, ...
- 后台任务/Background tasks
 - 监控网络和I/O设备/Monitoring network & I/O devices



多进程举例/Multiprocessing Example

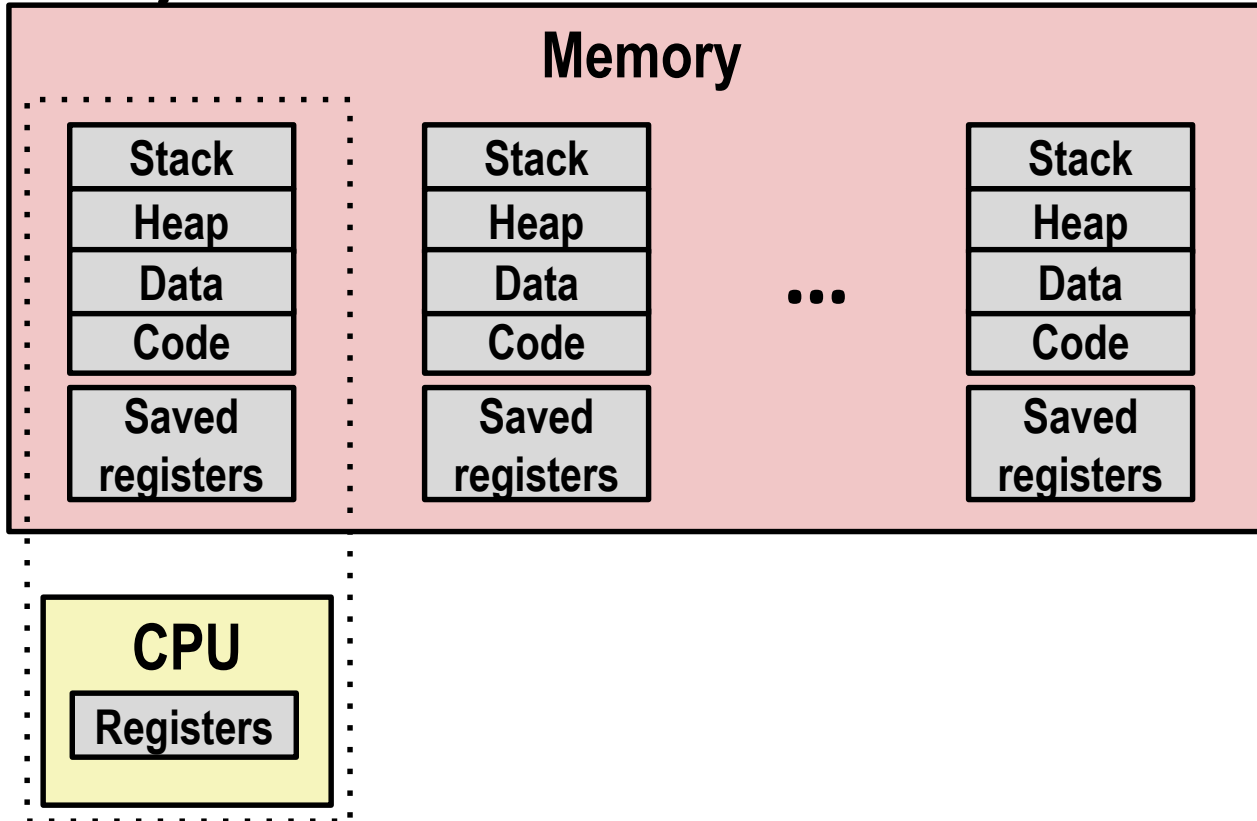
```
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14  CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

PID    COMMAND    %CPU TIME    #TH    #WQ    #PORT    #MREG    RPRVT    RSHRD    RSIZE    VPRVT    VSIZE
99217-  Microsoft Of 0.0 02:28.34 4    1    202    418    21M    24M    21M    66M    763M
99051   usbmuxd     0.0 00:04.10 3    1    47     66     436K    216K    480K    60M    2422M
99006   iTunesHelper 0.0 00:01.23 2    1    55     78     728K    3124K    1124K    43M    2429M
84286   bash        0.0 00:00.11 1    0    20     24     224K    732K    484K    17M    2378M
84285   xterm       0.0 00:00.83 1    0    32     73     656K    872K    692K    9728K    2382M
55939-  Microsoft Ex 0.3 21:58.97 10   3    360    954    16M    65M    46M    114M    1057M
54751   sleep       0.0 00:00.00 1    0    17     20     92K     212K    360K    9632K    2370M
54739   launchdadd  0.0 00:00.00 2    1    33     50     488K    220K    1736K    48M    2409M
54737   top         6.5 00:02.53 1/1  0    30     29     1416K    216K    2124K    17M    2378M
54719   automountd  0.0 00:00.02 7    1    53     64     860K    216K    2184K    53M    2413M
54701   ocspd       0.0 00:00.05 4    1    61     54     1268K    2644K    3132K    50M    2426M
54691   Grab        0.6 00:02.75 6    3    222+   388+   15M+    26M+    40M+    75M+    2556M+
54681   mdworker    0.0 00:01.67 4    1    52     91     7628K    7412K    16M    48M    2438M
50478   xterm       0.0 00:00.13 1    0    32     73     280K    872K    532K    9700K    2382M
50410   xterm       0.0 00:00.70 1    0    20     35     52K     216K    88K     18M    2392M
```

- 运行top命令 Running program “top” on Mac
- 系统有123个进程，5个是活跃状态/System has 123 processes, 5 of which are active
- 使用PID标识/Identified by Process ID (PID)



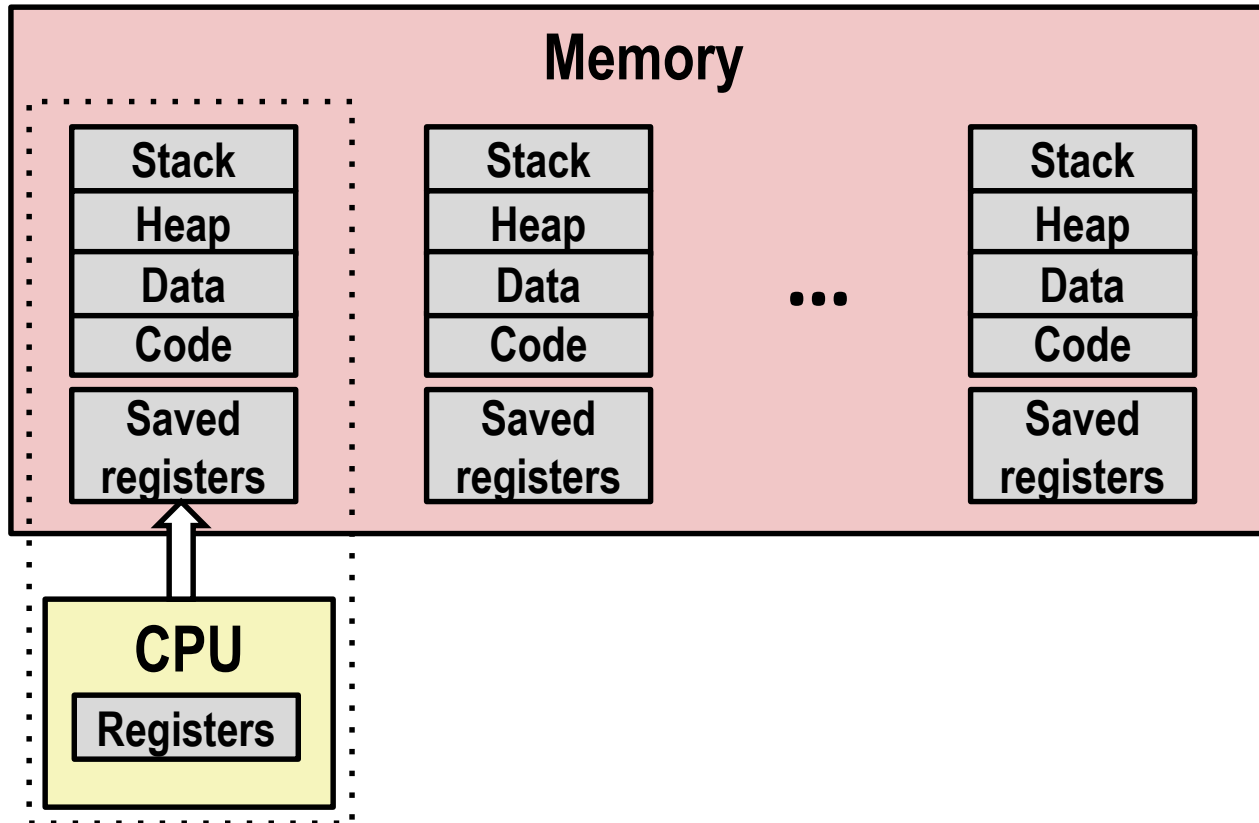
多进程(传统)真像/Multiprocessing: The (Traditional) Reality



- **单个处理器并发执行多个进程/Single processor executes multiple processes concurrently**
 - 进程交替执行（多任务）/ Process executions interleaved (multitasking)
 - 地址空间由虚拟内存系统管理 / Address spaces managed by virtual memory system (later in course)
 - 非激活进程的寄存器值存储在内存中 / Register values for nonexecuting processes saved in memory



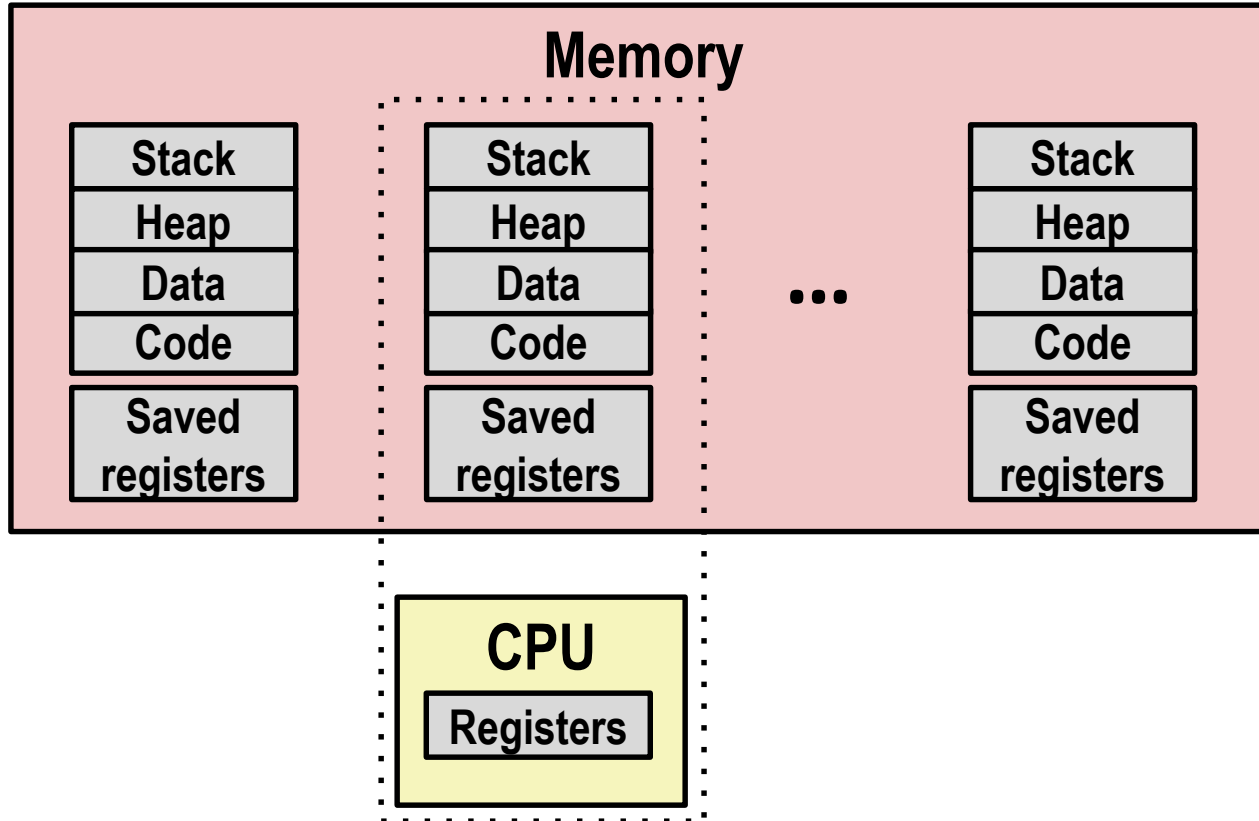
多进程真像/Multiprocessing: The (Traditional) Reality



- 将当前寄存器存储在内存里/Save current registers in memory



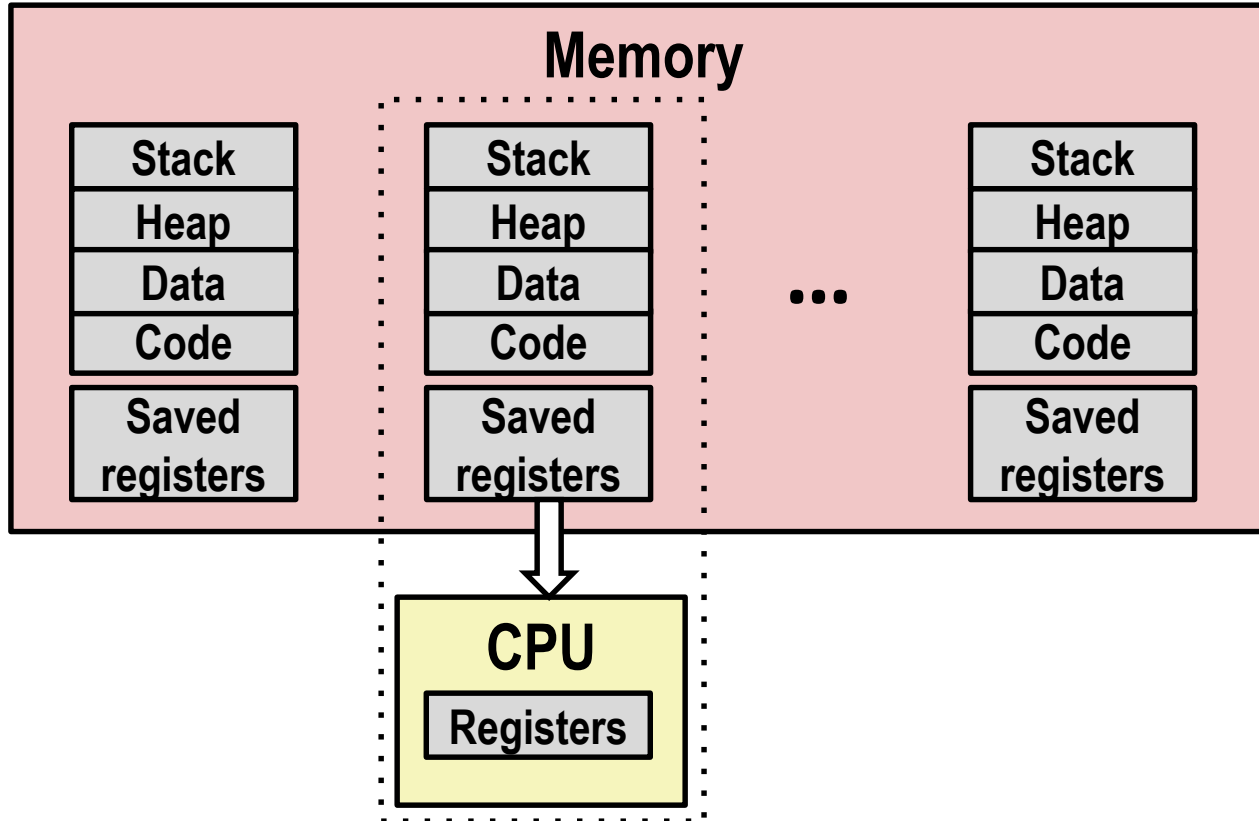
多进程真像 Multiprocessing: The (Traditional) Reality



- 调度下一个进程执行/Schedule next process for execution



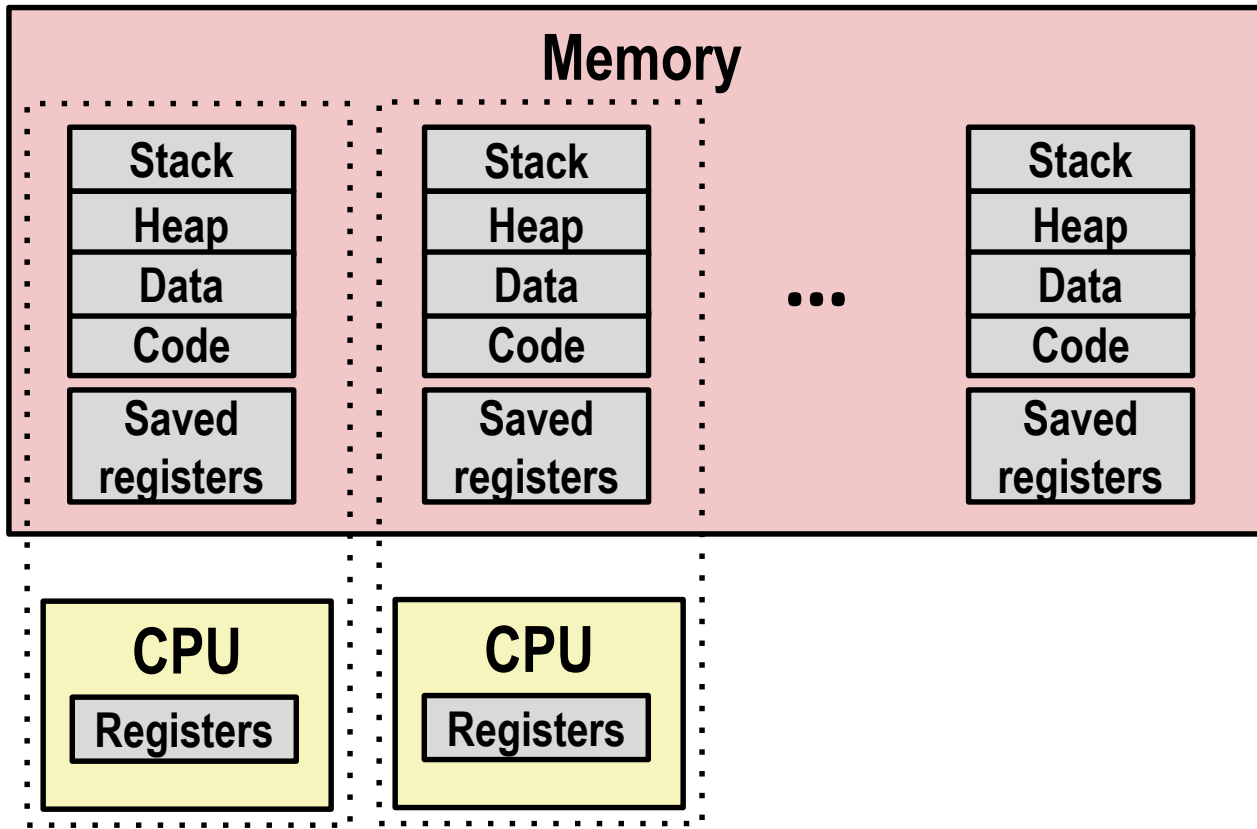
多进程真像/Multiprocessing: The (Traditional) Reality



- 加载寄存器并切换地址空间（上下文切换） Load saved registers and switch address space (context switch)



多进程真像 Multiprocessing: The (Modern) Reality



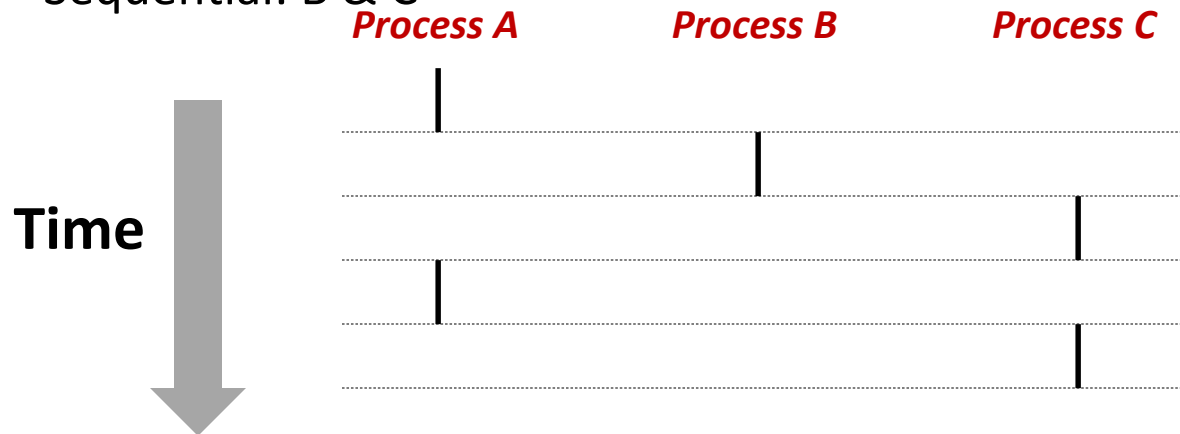
■ 多处理器/Multicore processors

- 一个芯片上有多个CPU/Multiple CPUs on single chip
- 共享主存储器（部分Cache）/Share main memory (and some of the caches)
- 每个可以执行一个进程/Each can execute a separate process
 - 调度是由系统内核完成的/Scheduling of process onto cores done by kernel



并发进程/Concurrent Processes


- 每个进程是一个逻辑控制流/Each process is a logical control flow.
- 两个进程并发运行如果在时间上重叠/Two processes *run concurrently* (are concurrent) if their flows overlap in time
- 否则是顺序执行/Otherwise, they are *sequential*
- 例如 (在单个核上运行) /Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



并发进程的用户视图/User View of Concurrent Processes



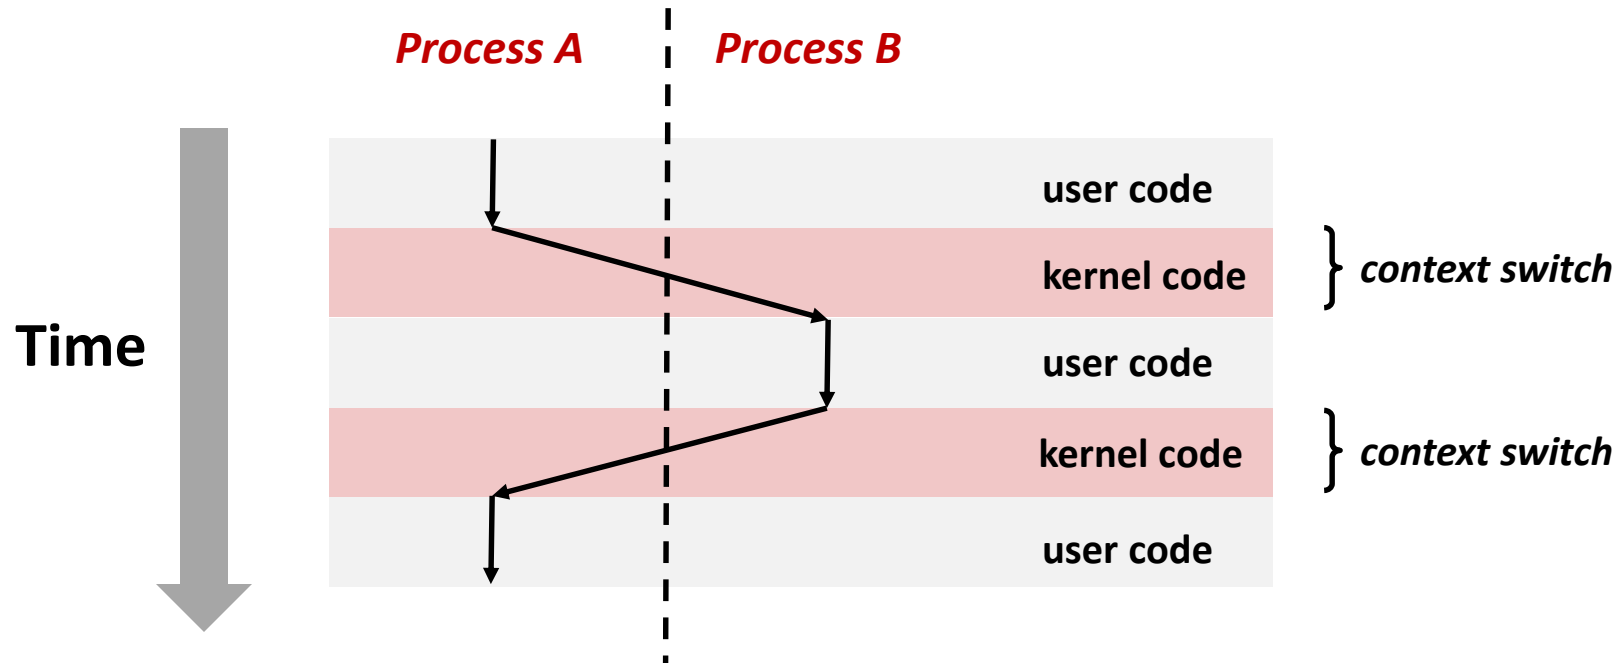
- 每个并发进程的控制流在时间上并不重叠/Control flows for concurrent processes are physically disjoint in time
- 但是我们可以认为所有的并发进程都是并行执行的/However, we can think of concurrent processes as running in parallel with each other





上下文切换/Context Switching

- 进程是由操作系统内核管理的/Processes are managed by a shared chunk of memory-resident OS code called the **kernel**
 - 重点：内核不是一个独立的进程，而是作为某些进程的一部分运行
Important: the kernel is not a separate process, but rather runs as part of some existing process.
- 上下文切换使得控制流从一个进程切换到另一个进程 Control flow passes from one process to another via a **context switch**





内容提纲

- 异常控制流 Exceptional Control Flow
- 异常 Exceptions
- 进程 Processes
- 进程控制 Process Control



系统调用错误处理/System Call Error Handling

- 出错时，系统函数返回-1并通过全局变量`errno`给出出错原因/On error, Linux system-level functions typically return -1 and set global variable `errno` to indicate cause.
- 硬性规定/Hard and fast rule:
 - 每次系统调用要检查调用结果/You must check the return status of every system-level function
 - 返回值为`void`的除外/Only exception is the handful of functions that return `void`
- 例如/Example:

```
if ((pid = fork()) < 0) {  
    fprintf(stderr, "fork error: %s\n", strerror(errno));  
    exit(0);  
}
```



错误报告函数/Error-reporting functions

- 可以使用错误报告函数简化处理/Can simplify somewhat using an *error-reporting function*:

```
void unix_error(char *msg) /* Unix-style error */
{
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));
    exit(0);
}
```

```
if ((pid = fork()) < 0)
    unix_error("fork error");
```



错误处理包装函数/Wrappers Error-handling Wrappers

- 采用Stevens-style 的错误处理方式简化代码/We simplify the code we present to you even further by using Stevens-style error-handling wrappers:

```
pid_t Fork(void)
{
    pid_t pid;

    if ((pid = fork()) < 0)
        unix_error("Fork error");
    return pid;
}
```

```
pid = Fork();
```

Stevens-style : <https://www.superfrink.net/docs/debugging-intro.html>



获得进程/PID Obtaining Process IDs

- `pid_t getpid(void)`
 - 返回当前进程的PID/Returns PID of current process
- `pid_t getppid(void)`
 - 返回父进程的PID/Returns PID of parent process



创建和终止进程/Creating and Terminating Processes

从程序员的角度，可以认为一个进程处于三种状态之一/

From a programmer's perspective, we can think of a process as being in one of three states

■ 运行/Running

- 进程或者正在执行，或者等待被执行，最终会被内核调度
/Process is either executing, or waiting to be executed and will eventually be *scheduled* (i.e., chosen to execute) by the kernel

■ 停止/Stopped

- 进程执行被挂起，直到被触发重新调度执行 Process execution is *suspended* and will not be scheduled until further notice (next lecture when we study signals)

■ 终止 Terminated

- 进程永远停止运行 Process is stopped permanently



进程终止 Terminating Processes

- 进程由于以下三个原因终止 Process becomes terminated for one of three reasons:
 - 收到一个信号，这个信号的默认动作是终止/ Receiving a signal whose default action is to terminate (next lecture)
 - 从主函数返回 Returning from the `main` routine
 - 调用`exit`函数 Calling the `exit` function
- `void exit(int status)`
 - 以某个状态`status`终止/Terminates with an *exit status* of `status`
 - 惯例：返回0表示正常退出，其他表示出错/Convention: normal return status is 0, nonzero on error
 - 也可以通过`main`函数的返回值显式设置进程退出状态/Another way to explicitly set the exit status is to return an integer value from the main routine
- `exit`只会被调用一次且会不返回/`exit` is called **once** but **never** returns.



创建进程/Creating Processes

- **父进程通过fork创建一个新的进程** / Parent process creates a new running *child process* by calling `fork`
- `int fork(void)`
 - 对子进程返回0, 对父进程返回子进程的PID / Returns 0 to the child process, child's PID to parent process
 - 子进程几乎和父进程是一样的 / Child is *almost* identical to parent:
 - 子进程具有和父进程一样但是独立的虚拟地址空间 / Child get an identical (but separate) copy of the parent's virtual address space.
 - 子进程会有和父进程一样的文件描述符拷贝 / Child gets identical copies of the parent's open file descriptors
 - 子进程和父进程的PID不同 / Child has a different PID than the parent
- **fork是很有意思的 (让人困惑的) , 因为只调用了一次但是返回两次** / `fork` is interesting (and often confusing) because it is called *once* but returns *twice*



fork举例 / fork Example

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

fork.c

```
linux> ./fork
parent: x=0
child : x=2
```

- 一次调用，两次返回/Call once, return twice
- 并发执行/Concurrent execution
 - 不能预测父进程和子进程之间的执行顺序/Can't predict execution order of parent and child
- 重复的但是独立的地址空间/Duplicate but separate address space
 - fork时x的值是1/x has a value of 1 when fork returns in parent and child
 - 后续对x的修改都是独立的/Subsequent changes to x are independent
- 共享打开的文件/Shared open files
 - `stdio`在子进程和父进程中都是一样的/`stdout` is the same in both parent and child



使用进程图描述fork /Modeling fork with Process Graphs

- **进程图是描述并发程序中语句偏序关系的有用工具/A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program:**
 - 每个顶点是一个语句/Each vertex is the execution of a statement
 - $a \rightarrow b$ 表示a在b之前发生/ $a \rightarrow b$ means a happens before b
 - 每个边可以用当前值或者变量进行标注/Edges can be labeled with current value of variables
 - `printf`节点可以用标注为output/`printf` vertices can be labeled with output
 - 每个图的开始节点没有入边/Each graph begins with a vertex with no inedges
- **图的任何拓扑排序都对应一个可行的全局序/Any *topological sort* of the graph corresponds to a feasible total ordering.**
 - 全局序节点之间的边从左指向右/Total ordering of vertices where all edges point from left to right



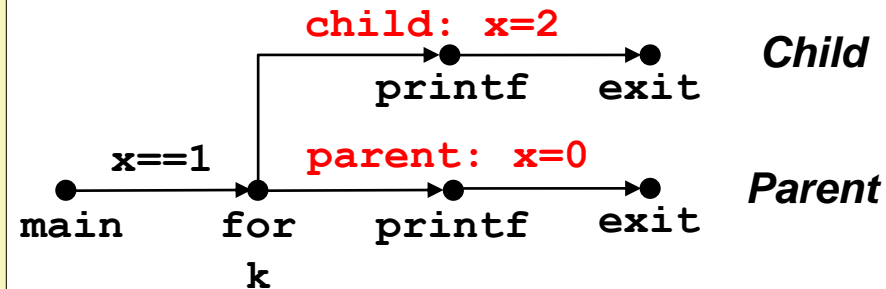
进程图举例/Process Graph Example

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

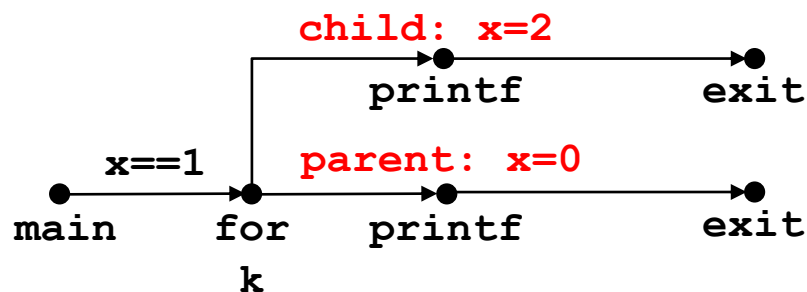
fork.c






解释进程图/Interpreting Process Graphs

■ 原图/Original graph:




■ 标注后的图/Relabled graph:




可能的全局序/

Feasible total ordering:



不可能的全局序/

Infeasible total ordering:

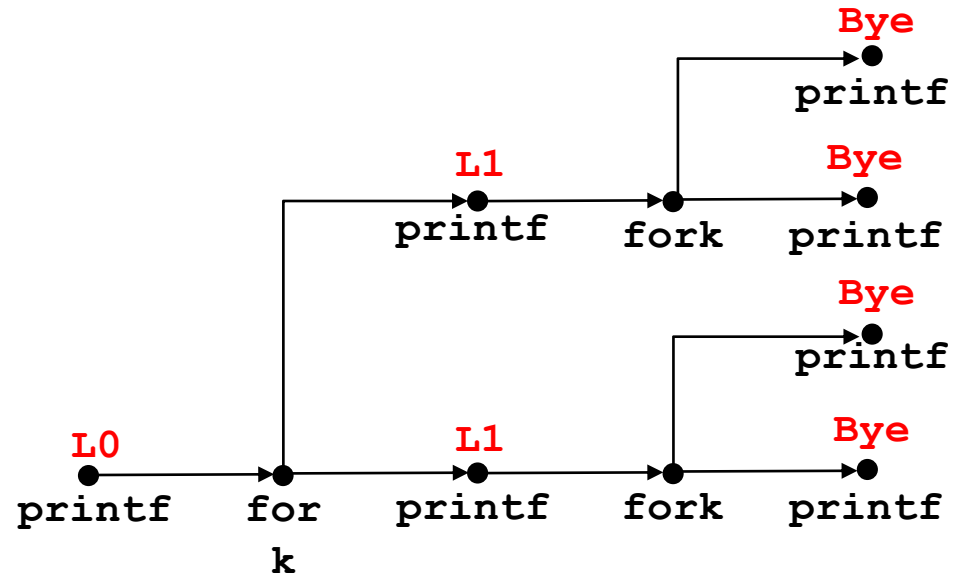




fork举例：两个连续的fork/fork Example: Two consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

forks.c



可能的输出/

Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye

不可能的输出/

Infeasible output:

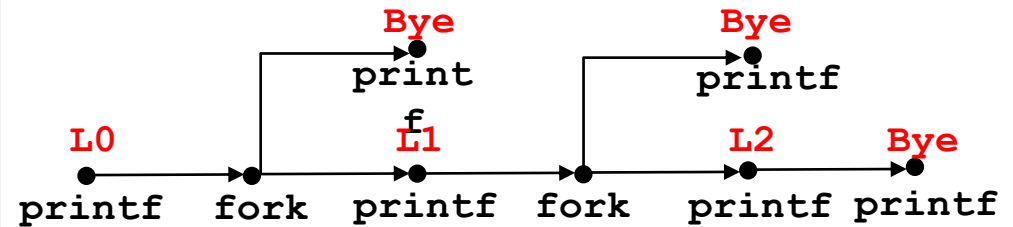
L0
Bye
L1
Bye
L1
Bye
Bye



fork举例：父类进程中的嵌套forks/ fork Example: Nested forks in parent

```
void fork4()  
{  
    printf("L0\n");  
    if (fork() != 0) {  
        printf("L1\n");  
        if (fork() != 0) {  
            printf("L2\n");  
        }  
    }  
    printf("Bye\n");  
}
```

forks.c



可能的输出/
Feasible output:

L0
L1
Bye
Bye
L2
Bye

不可能的输出/
Infeasible output:


L0
Bye
L1
Bye
Bye
L2



fork举例，子进程嵌套fork/fork Example: Nested forks in children

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

forks.c



可能的输出/
Feasible output:

L0
Bye
L1
L2
Bye
Bye

不可能的输出/
Infeasible output:

L0
Bye
L1
Bye
Bye
L2



捕获子进程/Reaping Child Processes

■ Idea

- 进程终止后仍然消耗系统资源/When process terminates, it still consumes system resources
 - 例如：退出状态，各种OS表格/Examples: Exit status, various OS tables
- 僵尸 Called a “zombie”
 - Living corpse, half alive and half dead

■ 捕获 Reaping

- 父类进程等待子进程终止/Performed by parent on terminated child (using `wait` or `waitpid`)
- 父类进程获得退出状态信息/Parent is given exit status information
- 内核删掉僵尸子进程 Kernel then deletes zombie child process

■ 如果父类进行没有回收会怎么样？ What if parent doesn't reap?

- 如果父类不回收，则由init进程回收 If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
- 所以只需要显式回收长时间运行的进程 So, only need explicit reaping in long-running processes
 - e.g., shells and servers

僵尸进程举例/Zombie Example

```
void fork7() {
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

forks.c

```
linux> ./forks 7 &
```

```
[1] 6639
```

```
Running Parent, PID = 6639
```

```
Terminating Child, PID = 6640
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6639	ttyp9	00:00:03	forks
6640	ttyp9	00:00:00	forks <defunct>
6641	ttyp9	00:00:00	ps

```
linux> kill 6639
```

```
[1] Terminated
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6642	ttyp9	00:00:00	ps

■ 使用ps命令显示子进程是“defunct”/ps shows child process as “defunct” (i.e., a zombie)

■ 杀死父进程允许子进程通过init被重新捕获/Killing parent allows child to be reaped by init

非终止子进程举例 Non-terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
            getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
            getpid());
        exit(0);
    }
}
```

forks.c

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6676 tttyp9        00:00:06 forks
 6677 tttyp9        00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9        00:00:00 tcsh
 6678 tttyp9        00:00:00 ps
```

■ 父进程终止后子进程仍然活跃
/Child process still active even
though parent has terminated

■ 必须显式杀死子进程，否则子进程
将永远运行下去/Must kill child
explicitly, or else will keep running
indefinitely



wait: 子进程同步 / wait: Synchronizing with Children


- **父进程通过wait函数重新对接子进程 / Parent reaps a child by calling the wait function**
- **int wait(int *child_status)**
 - 挂起当前进程直到其中一个子进程退出 / Suspends current process until one of its children terminates
 - 返回的是终止的子进程的PID / Return value is the **pid** of the child process that terminated
 - 如果**child_status**不为空, 该指针指向的值表示的是子进程终止和退出的原因 / If **child_status != NULL**, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - 可以使用wait.h中定义的下列宏进行检查 / Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`
 - See textbook for details



wait: 子进程同步 / wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

forks.c



Feasible output:

HC
HP
CT
Bye

Infeasible output:

HP
CT
Bye
HC



另一个wait的例子/Another wait Example

- 多个子进程可能按照任意顺序完成/If multiple children completed, will take in arbitrary order
- 可以使用WIFEXITED 和WEXITSTATUS获取退出的信息/Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10() {  
    pid_t pid[N];  
    int i, child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0) {  
            exit(100+i); /* Child */  
        }  
    for (i = 0; i < N; i++) { /* Parent */  
        pid_t wpid = wait(&child_status);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

forks.c



waitpid: 等待特定进程/waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - 挂起当前进程直到特定进程终止/Suspends current process until specific process terminates
 - 很多不同选项（见教材）/Various options (see textbook)

```
void fork11() {
    pid_t pid[N];
    int i;
    int child_status;

    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = N-1; i >= 0; i--) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```

forks.c

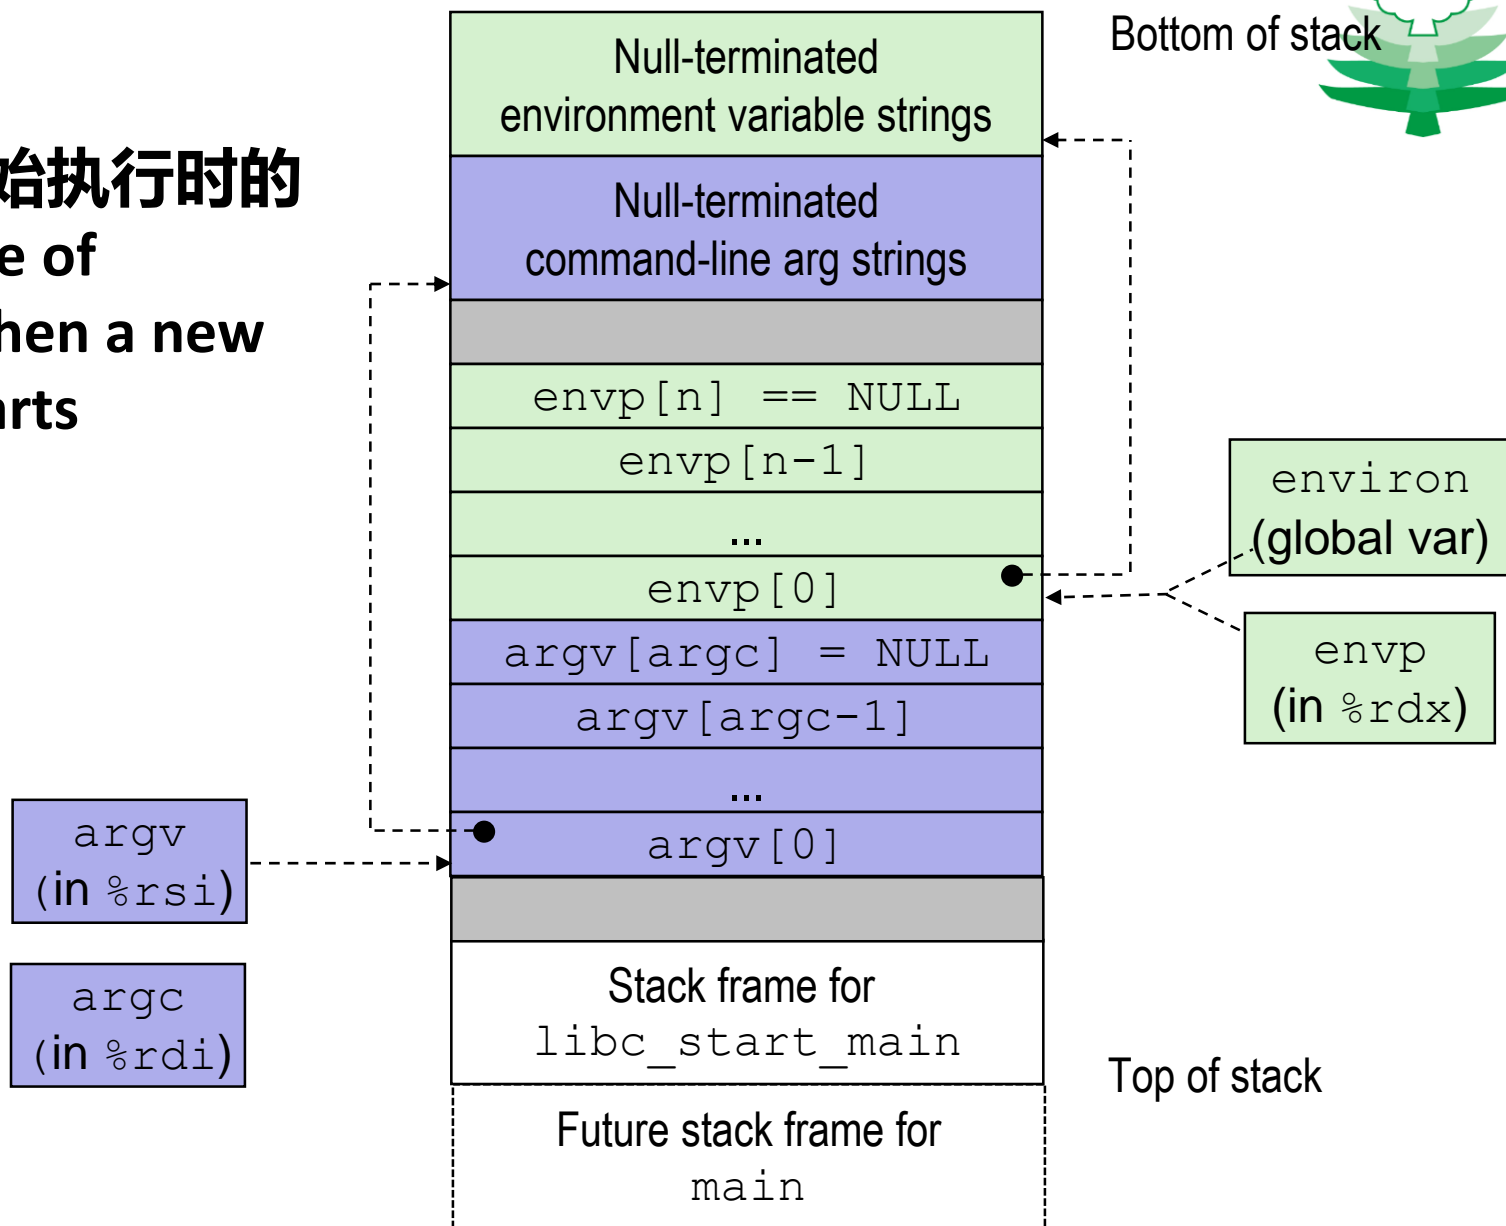


execve: 加载运行程序/execve: Loading and Running Programs

- **int execve(char *filename, char *argv[], char *envp[])**
- **在当前进程加载运行/Loads and runs in the current process:**
 - filename: Executable file **filename** / **可执行文件名字**
 - 目标代码文件或者以`#!/interpreter`开头的脚本文件 Can be object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
 - 带有参数列表/ ...with argument list **argv**
 - 按惯例`argv[0]`是第一个参数filename/By convention **argv[0]==filename**
 - 带有环境变量列表/ ...and environment variable list **envp**
 - “name=value” strings (e.g., `USER=droh`)
 - `getenv`, `putenv`, `putenv`
- **覆盖代码、数据和堆栈/Overwrites code, data, and stack**
 - 保留PID、打开的文件和信号上下文/Retains PID, open files and signal context
- **一次调用，并无返回/Called **once** and **never** returns**
 - ...除非有错误/...except if there is an error



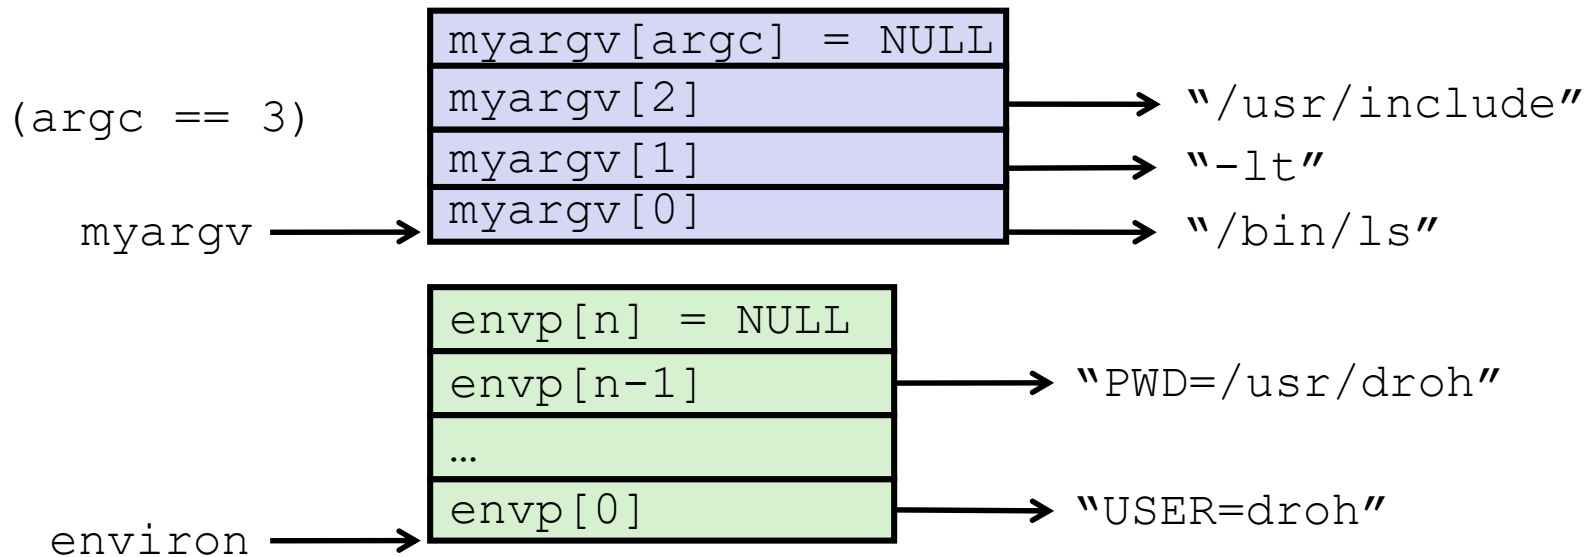
一个程序开始执行时的栈/Structure of the stack when a new program starts





execve举例 / execve Example

- 使用当前环境在子进程中执行“/bin/ls -lt /usr/include” /Executes “/bin/ls -lt /usr/include” in child process using current environment:



```
if ((pid = Fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```



总结 Summary

■ 异常/Exceptions

- 需要非标准控制流的事件/Events that require nonstandard control flow
- 由外部（中断）或内部（陷阱和故障）生成 / Generated externally (interrupts) or internally (traps and faults)

■ 进程/Processes

- 任意时刻，系统有多个进程/At any given time, system has multiple active processes
- 每个时刻单个核只能执行一个进程/Only one can execute at a time on a single core, though
- 每个进程看起来是独占处理器和私有内存空间/Each process appears to have total control of processor + private memory space



总结/Summary (cont.)

■ 生成新进程/Spawning processes

- 调用fork/Call `fork`
- 一次调用，两次返回/One call, two returns

■ 结束进程 Process completion

- 调用exit/Call `exit`
- 一次调用，无返回/One call, no return

■ 等待进程 Reaping and waiting for processes

- 调用wait或waitpid/Call `wait` or `waitpid`

■ 加载运行程序 Loading and running programs

- 调用execve/Call `execve` (or variant)
- 一次调用，一般情况下无返回/One call, (normally) no return