



# Virtual Memory: Systems

## 虚拟内存:系统

100076202: 计算机系统导论

**任课教师:**

**计卫星 宿红毅 张艳**

**原作者:**

Randal E. Bryant and David R. O'Hallaron



**Carnegie  
Mellon  
University**



# 内容提纲/Today

- 简单内存系统示例/Simple memory system example
- 实例：Core i7/Linux内存系统/Case study: Core i7/Linux memory system
- 内存映射/Memory mapping



# 符号回顾/Review of Symbols

## ■ 基本参数/Basic Parameters

- $N = 2^n$ : Number of addresses in virtual address space/虚拟内存空间的地址数量
- $M = 2^m$ : Number of addresses in physical address space/物理内存空间的地址数量
- $P = 2^p$ : Page size (bytes)/页大小 (字节)

## ■ 虚拟页划分/Components of the virtual address (VA)

- TLBI: TLB index/TLB索引
- TLBT: TLB tag/TLB标记
- VPO: Virtual page offset/虚拟页偏移量
- VPN: Virtual page number /虚拟页编号

## ■ 物理页划分/Components of the physical address (PA)

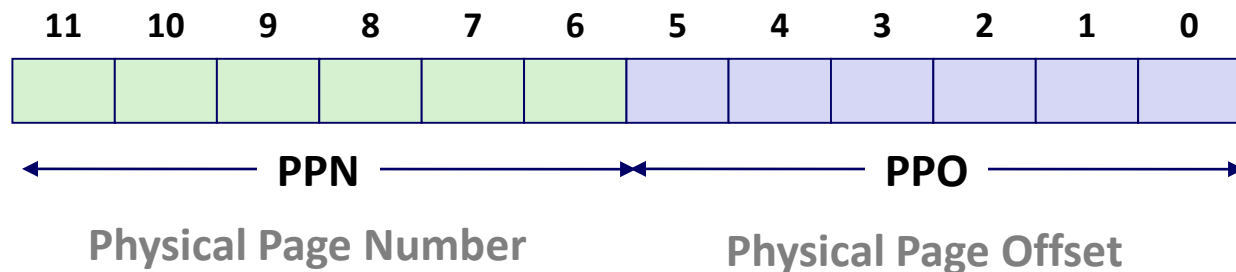
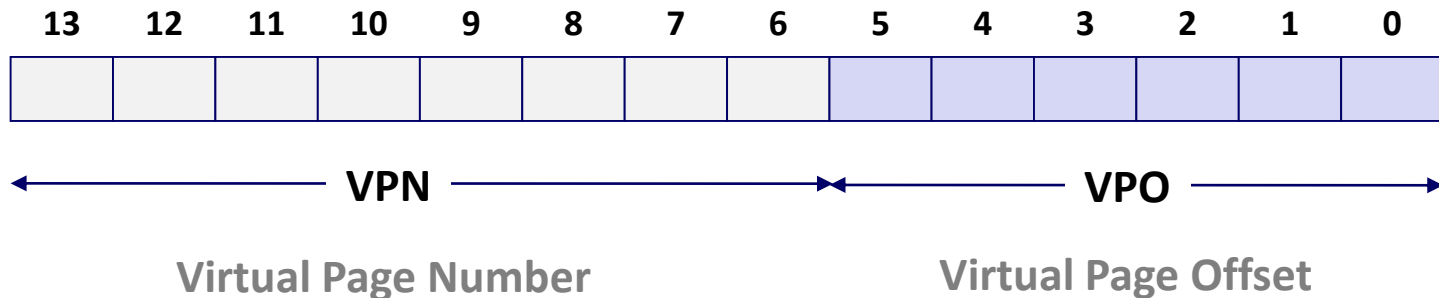
- PPO: Physical page offset (same as VPO)/物理页偏移量
- PPN: Physical page number/物理页号
- CO: Byte offset within cache line/Cache行中的偏移量
- CI: Cache index/Cache索引
- CT: Cache tag/Cache标记



# 简单的内存系统示例/Simple Memory System Example

## ■ 寻址/Addressing

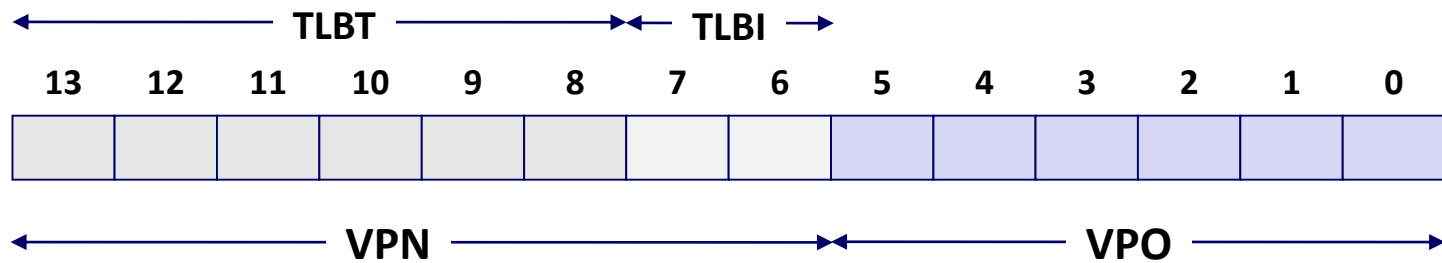
- 14位虚拟地址/14-bit virtual addresses
- 12位物理地址/12-bit physical address
- 页大小为64字节/Page size = 64 bytes





# 1. 简单内存系统TLB/Simple Memory System TLB

- 16条记录/16 entries
- 4路组相联/4-way associative



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

## 2. 简单内存系统页表/Simple Memory System Page Table



只显示了256的前16个/Only show first 16 entries (out of 256)

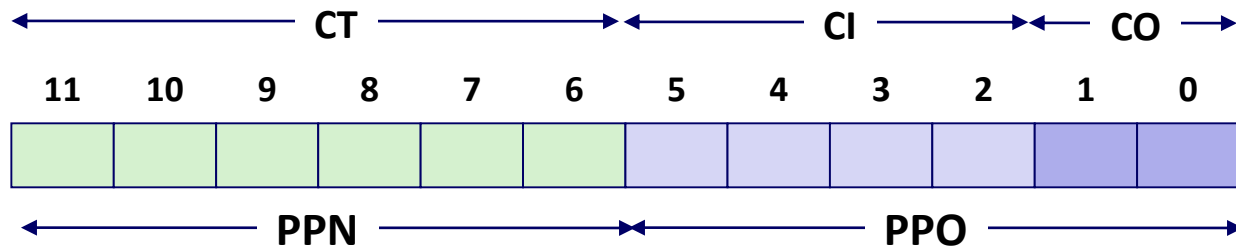
<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
00	28	1
01	—	0
02	33	1
03	02	1
04	—	0
05	16	1
06	—	0
07	—	0

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
08	13	1
09	17	1
0A	09	1
0B	—	0
0C	—	0
0D	2D	1
0E	11	1
0F	0D	1



### 3. 简单内存系统Cache/Simple Memory System Cache

- 16行，4字节Cache块大小/16 lines, 4-byte block size
- 物理寻址/Physically addressed
- 直接映射/Direct mapped



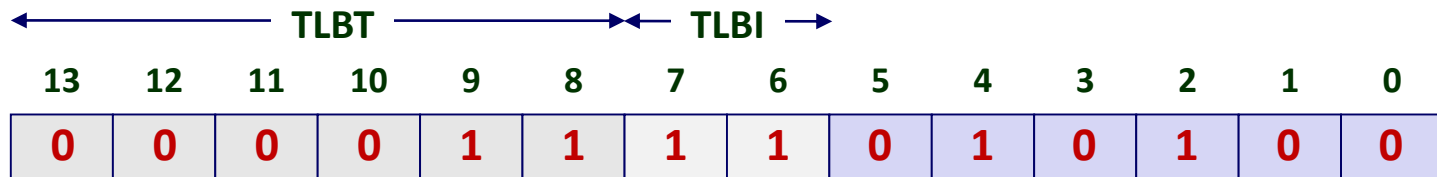
<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
0	19	1	99	11	23	11
1	15	0	—	—	—	—
2	1B	1	00	02	04	08
3	36	0	—	—	—	—
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	—	—	—	—
7	16	1	11	C2	DF	03

<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
8	24	1	3A	00	51	89
9	2D	0	—	—	—	—
A	2D	1	93	15	DA	3B
B	0B	0	—	—	—	—
C	12	0	—	—	—	—
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	—	—	—	—



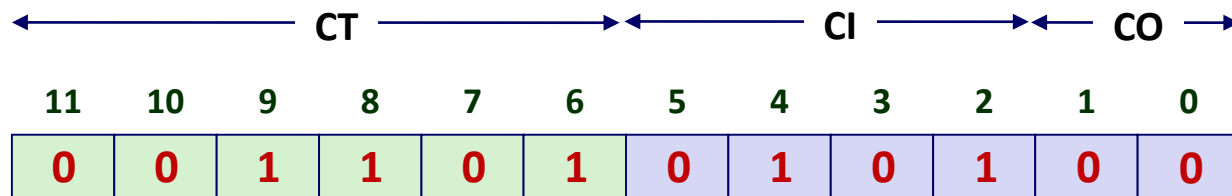
# 地址翻译示例#1/Address Translation Example #1

虚拟地址/Virtual Address: 0x03D4



VPN 0x0F TLBI 0x3 TLBT 0x03 TLB Hit? Y Page Fault? N PPN: 0x0D

物理地址/Physical Address



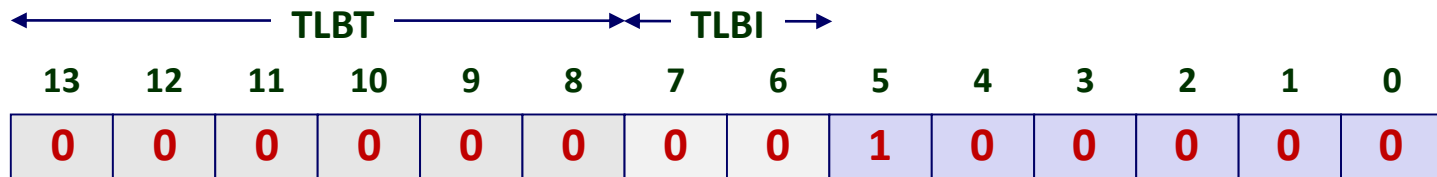
CO 0 CI 0x5 CT 0x0D Hit? Y Byte: 0x36





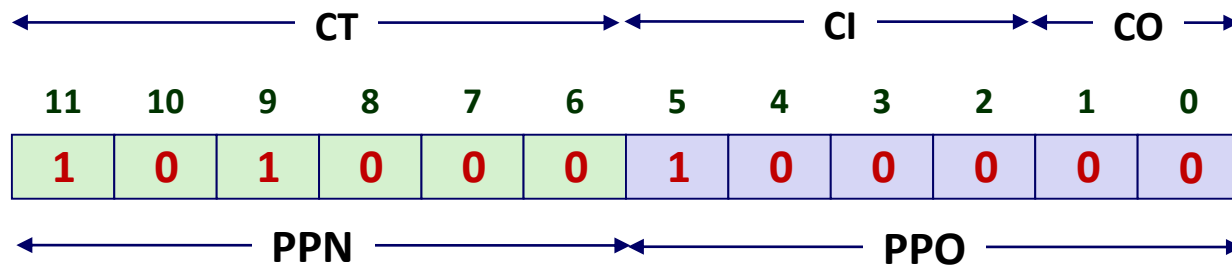
## 地址翻译示例#2/Address Translation Example #2

虚拟地址/Virtual Address: 0x0020



VPN 0x00    TLBI 0    TLBT 0x00    TLB Hit? N    Page Fault? N    PPN: 0x28

物理地址/Physical Address

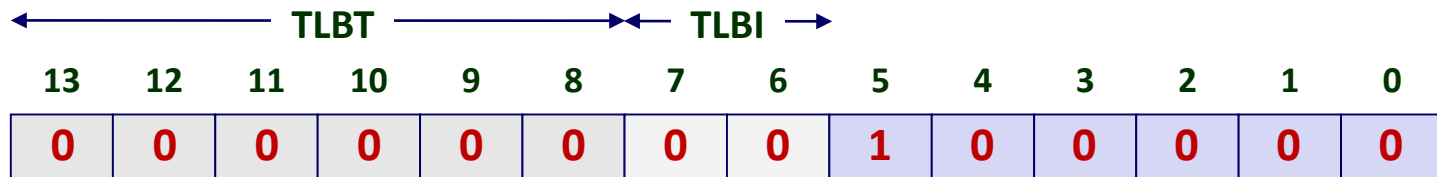


CO 0    CI 0x8    CT 0x28    Hit? N    Byte: Mem



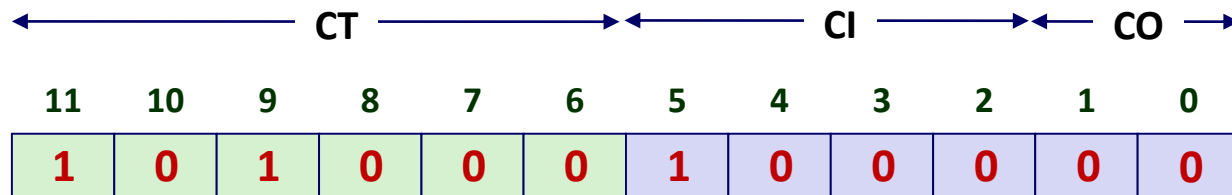
# 地址翻译示例#3/Address Translation Example #3

虚拟地址/Virtual Address: 0x0020



VPN 0x00    TLBI 0    TLBT 0x00    TLB Hit? N    Page Fault? N    PPN: 0x28

物理地址/Physical Address



CO 0    CI 0x8    CT 0x28    Hit? N    Byte: Mem



# 内容提纲/Today

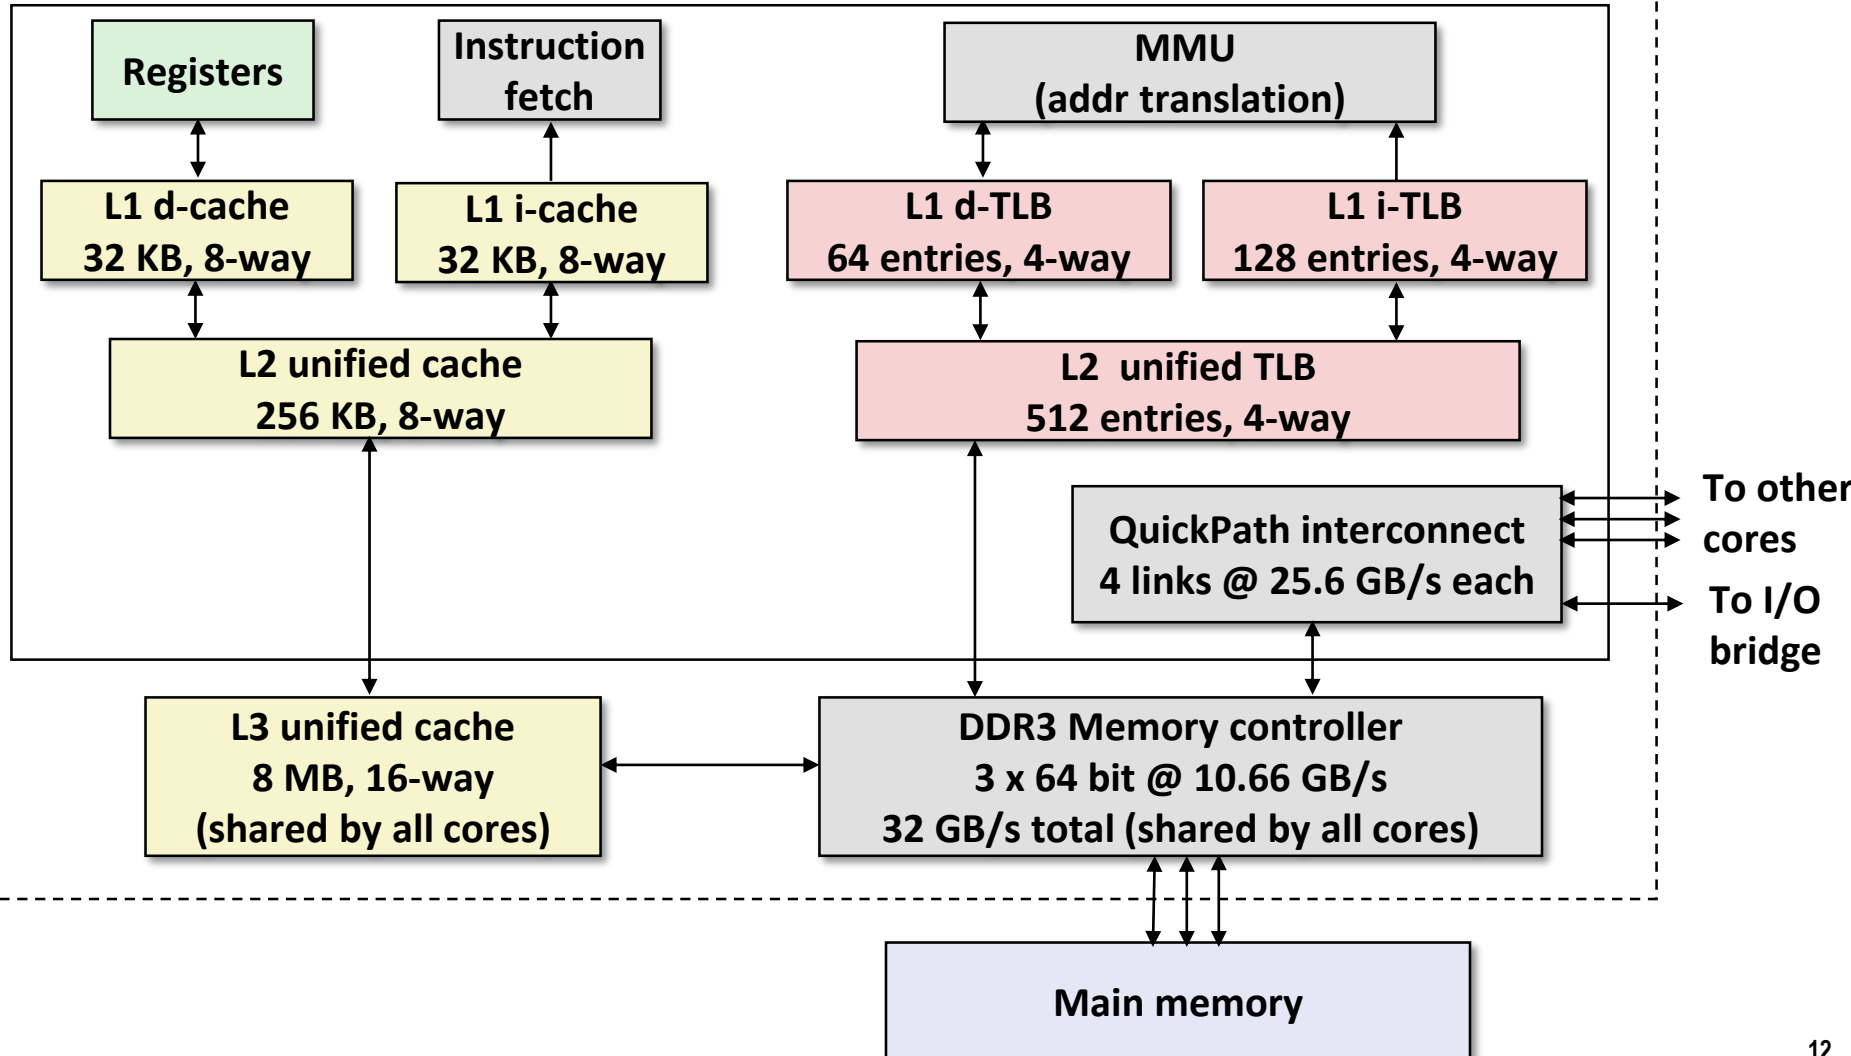
- 简单内存系统示例/Simple memory system example
- 实例：Core i7/Linux内存系统/Case study: Core i7/Linux memory system
- 内存映射/Memory mapping



# Intel Core i7存储系统/Intel Core i7 Memory System

## Processor package

Core x4





# 符号回顾/Review of Symbols

## ■ 基本参数/Basic Parameters

- $N = 2^n$ : Number of addresses in virtual address space/虚拟内存空间的地址数量
- $M = 2^m$ : Number of addresses in physical address space/物理内存空间的地址数量
- $P = 2^p$ : Page size (bytes)/页大小 (字节)

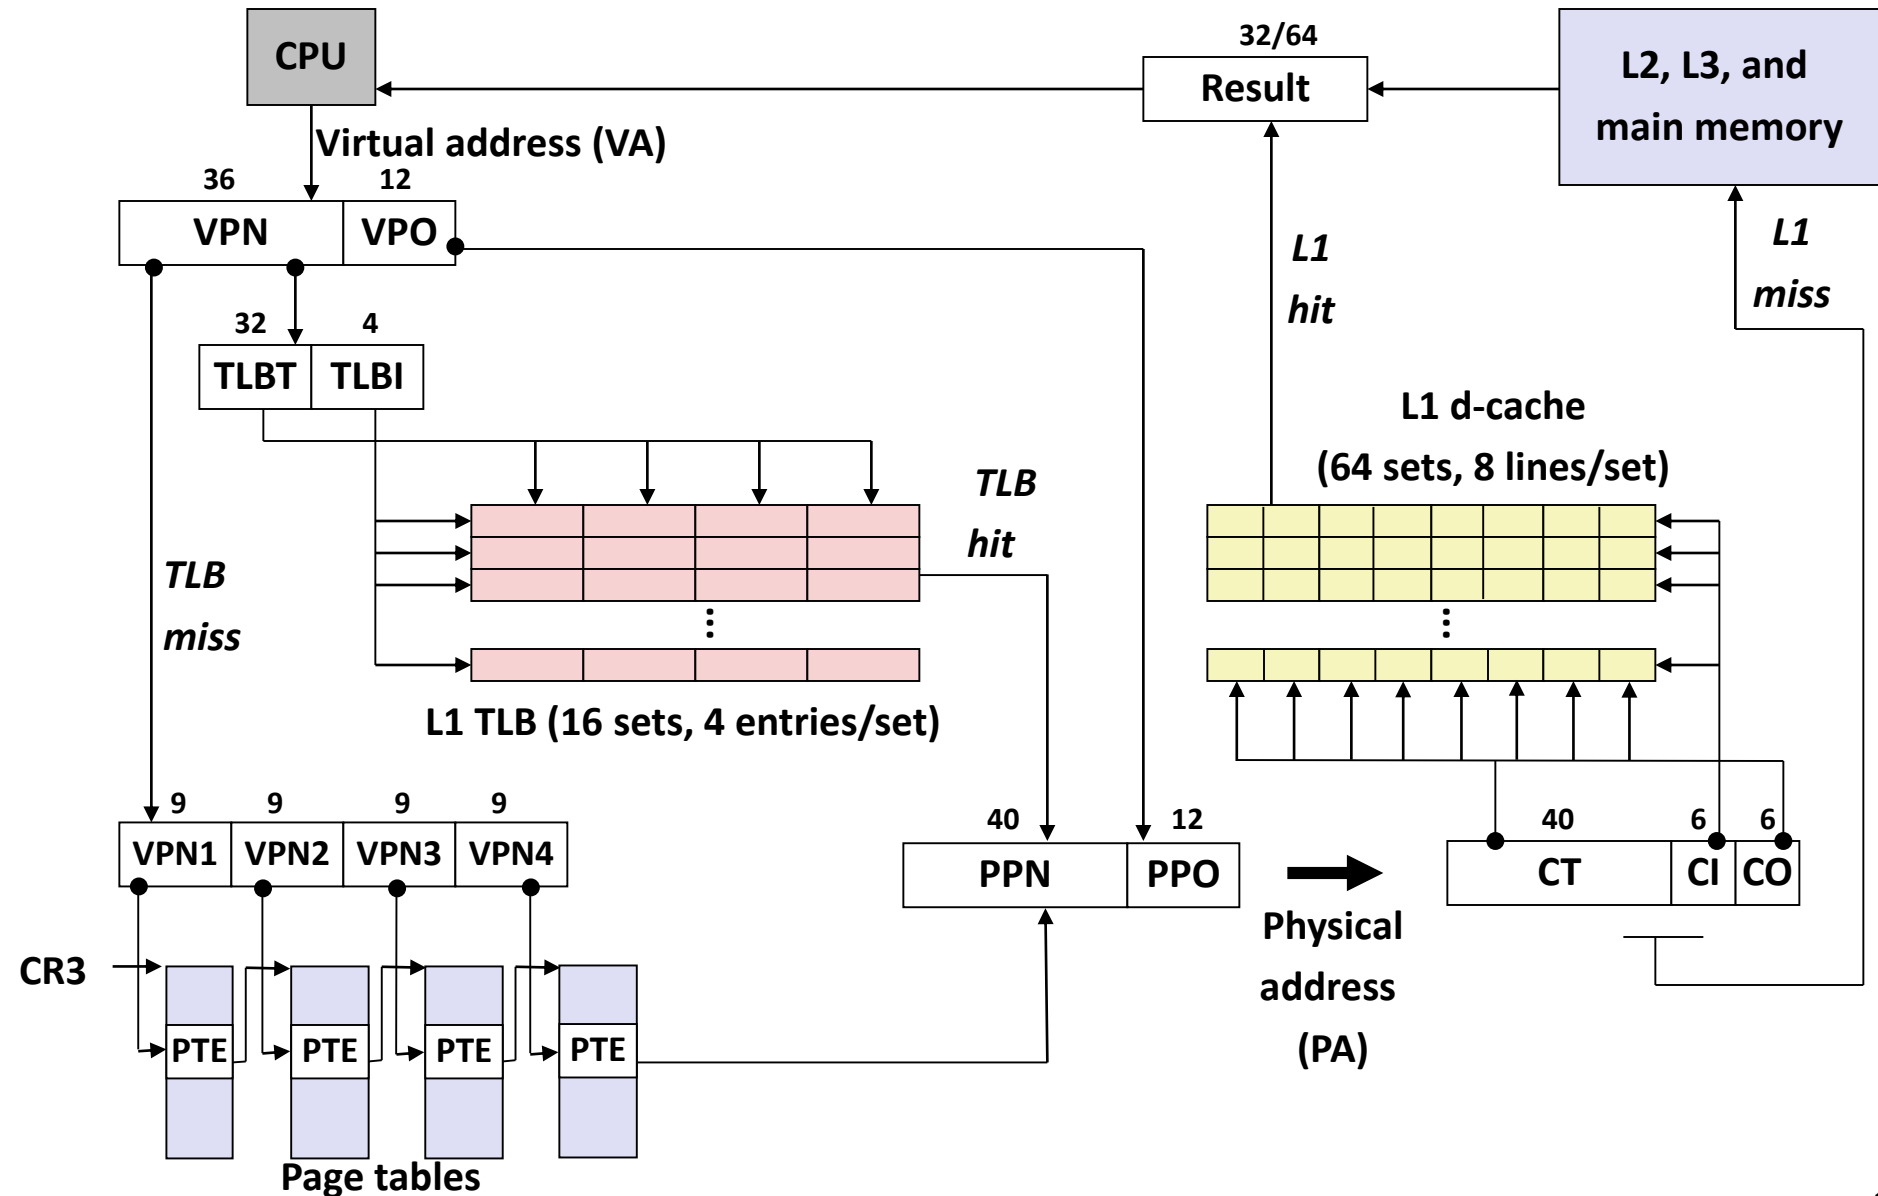
## ■ 虚拟页划分/Components of the virtual address (VA)

- TLBI: TLB index/TLB索引
- TLBT: TLB tag/TLB标记
- VPO: Virtual page offset/虚拟页偏移量
- VPN: Virtual page number /虚拟页编号

## ■ 物理页划分/Components of the physical address (PA)

- PPO: Physical page offset (same as VPO)/物理页偏移量
- PPN: Physical page number/物理页号
- CO: Byte offset within cache line/Cache行中的偏移量
- CI: Cache index/Cache索引
- CT: Cache tag/Cache标记

# 端到端Core i7地址翻译/End-to-end Core i7 Address Translation





# Core i7 1-3级页表记录/Core i7 Level 1-3 Page Table Entries

63	62	52	51	12				11	9	8	7	6	5	4	3	2	1	0
XD	Unused	页表物理及地址/ Page table physical base address				Unused				G	PS		A	CD	WT	U/S	R/W	P=1
操作系统可用（页表位于磁盘上） Available for OS (page table location on disk)																		P=0

**每个条目对应一个4k子页表，主要的域包括/Each entry references a 4K child page table. Significant fields:**

**P:**子页表是否在物理内存/Child page table present in physical memory (1) or not (0).

**R/W:** 只读或者读写权限标记位/Read-only or read-write access permission for all reachable pages.

**U/S:** 用户或特权（内核）模式标记位/user or supervisor (kernel) mode access permission for all reachable pages.

**WT:** 子页表的写透或者写回Cache策略/Write-through or write-back cache policy for the child page table.

**A:** 引用标记/Reference bit (由MMU读写时设置，软件清除/set by MMU on reads and writes, cleared by software).

**PS:** 页面大小，4KB或者4MB/Page size either 4 KB or 4 MB (defined for Level 1 PTEs only).

**Page table physical base address:** 物理页表地址的高40位(强制页表按照4KB对齐)/40 most significant bits of physical page table address (forces page tables to be 4KB aligned)

**XD:** 禁止或允许取指操作/Disable or enable instruction fetches from all pages reachable from this PTE.



# Core i7 4级页表条目/Core i7 Level 4 Page Table Entries

63	62	52	51	12 11				9	8	7	6	5	4	3	2	1	0
XD	Unused	物理页基址/ Page physical base address				Unused		G			D	A	CD	WT	U/S	R/W	P=1
操作系统可见（内存页位于磁盘） Available for OS (page location on disk)																	P=0

**Each entry references a 4K child page. Significant fields:**

**P:**子页表是否在物理内存/Child page table present in physical memory (1) or not (0).

**R/W:** 只读或者读写权限标记位/Read-only or read-write access permission for all reachable pages.

**U/S:** 用户或特权（内核）模式标记位/user or supervisor (kernel) mode access permission for all reachable pages.

**WT:** 子页表的写透或者写回Cache策略/Write-through or write-back cache policy for the child page table.

**A:** 引用标记/Reference bit (由MMU读写时设置，软件清除/set by MMU on reads and writes, cleared by software).

**D:** 脏位/Dirty bit (写操作时由MMU设置，软件清除/set by MMU on writes, cleared by software)

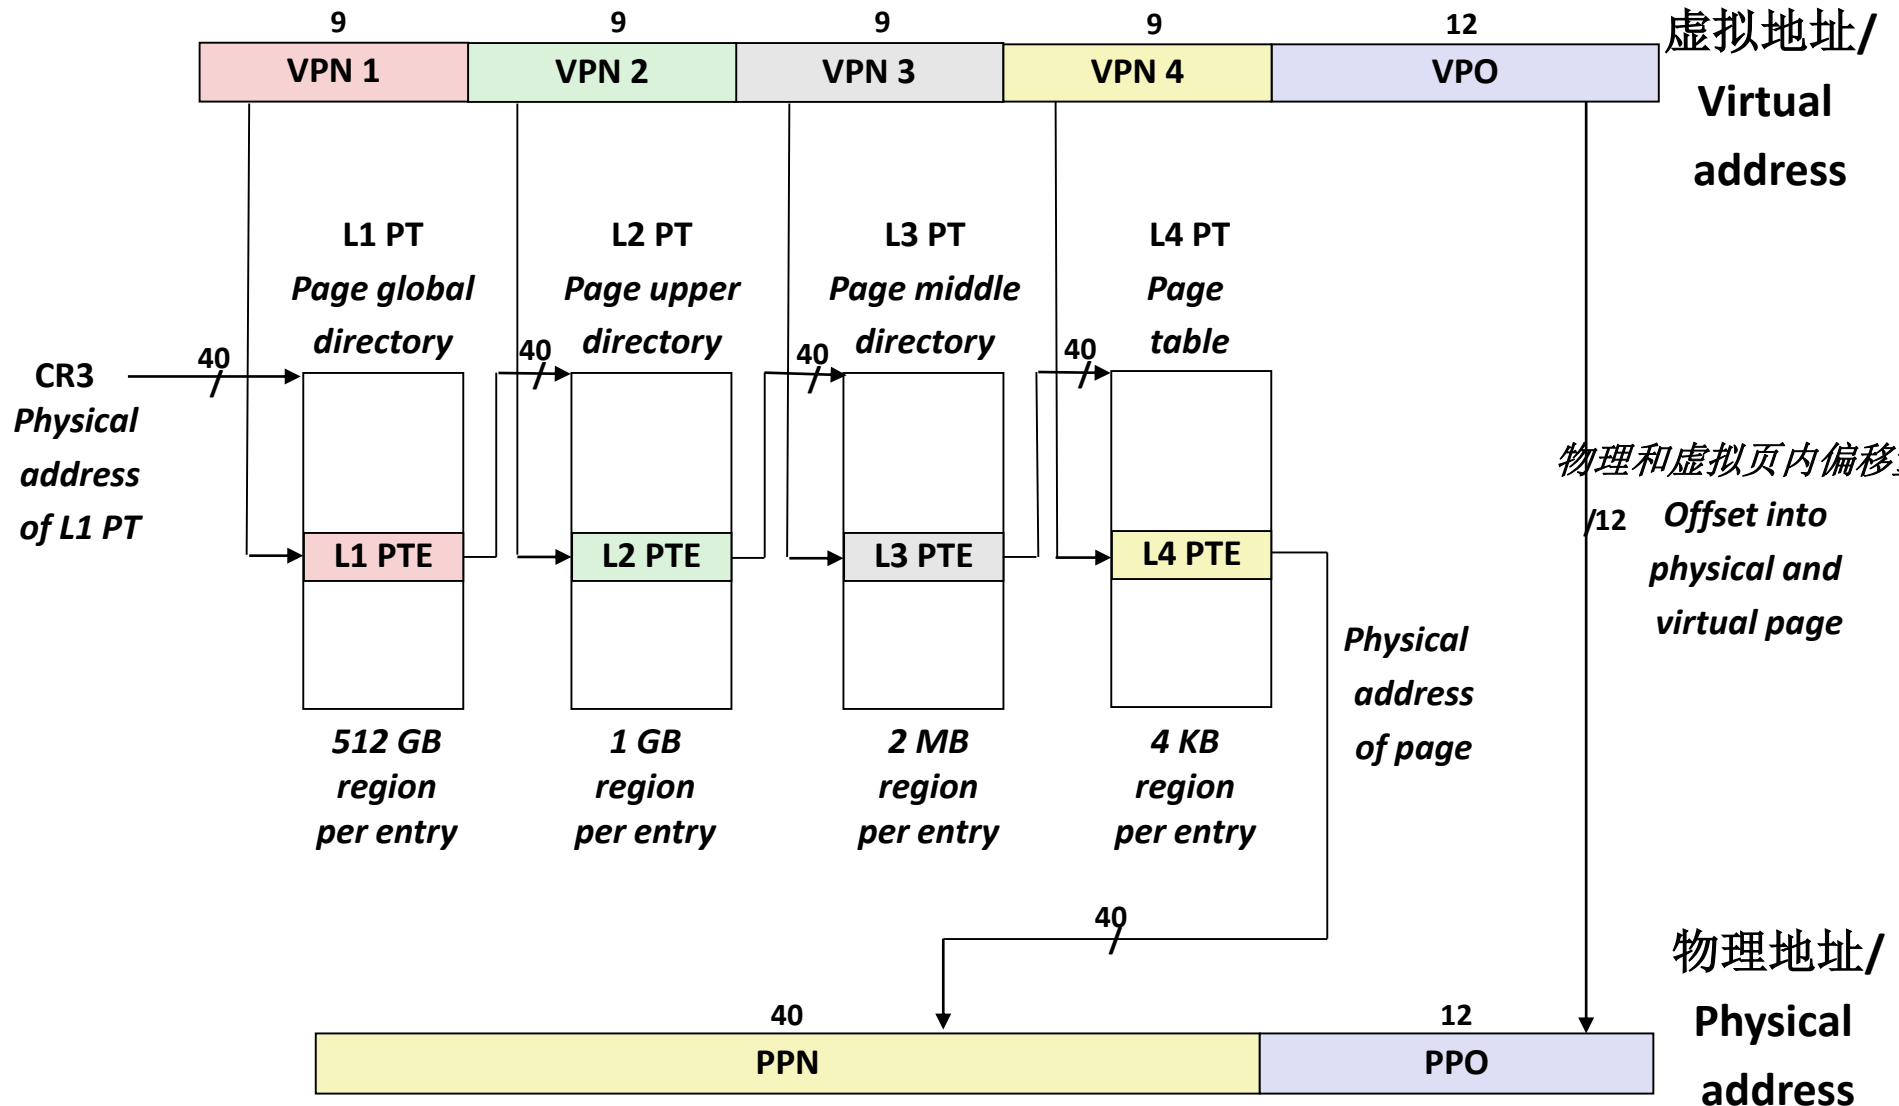
**Page physical base address:**物理页表地址的高40位(强制页表按照4KB对齐)/ 40 most significant bits of physical page address (forces pages to be 4KB aligned)

**XD:**禁止或允许取指操作/ Disable or enable instruction fetches from this page.



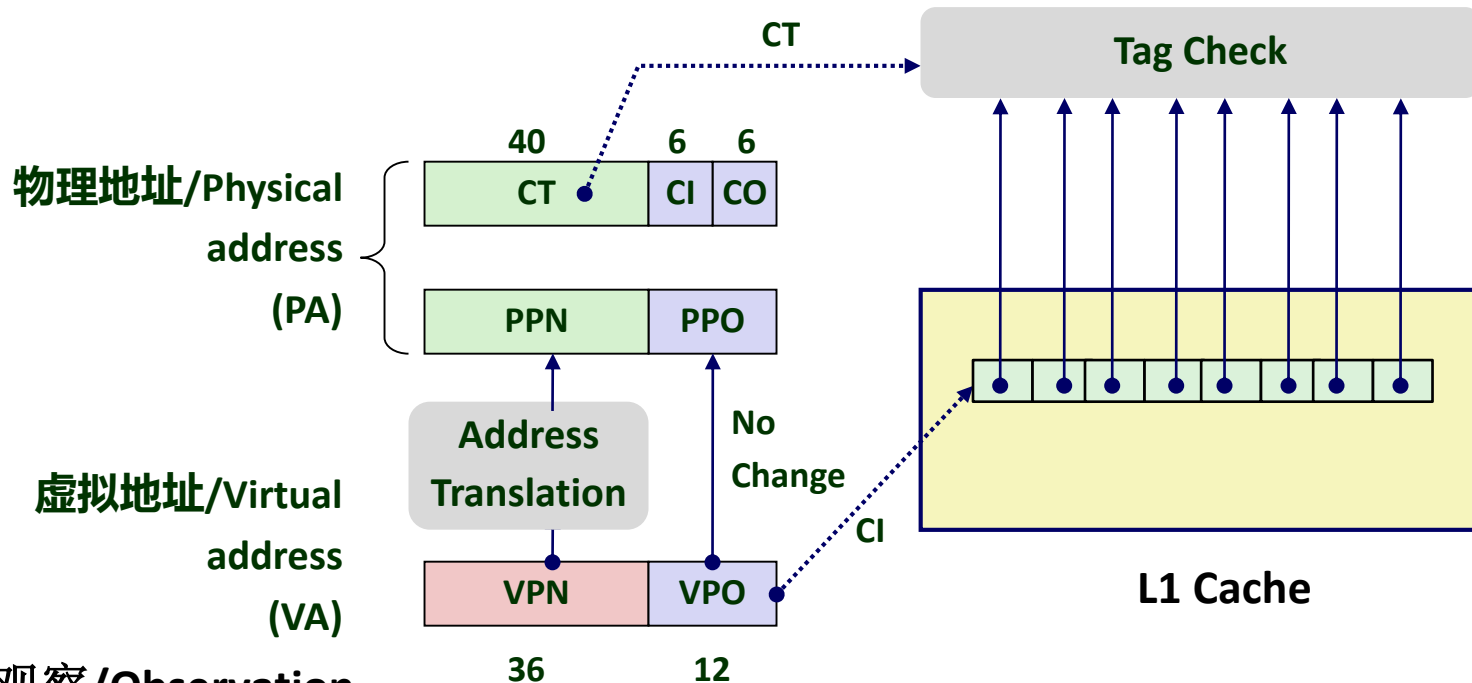


# Core i7页表翻译/Core i7 Page Table Translation





# L1访问加速小技巧/Cute Trick for Speeding Up L1 Access

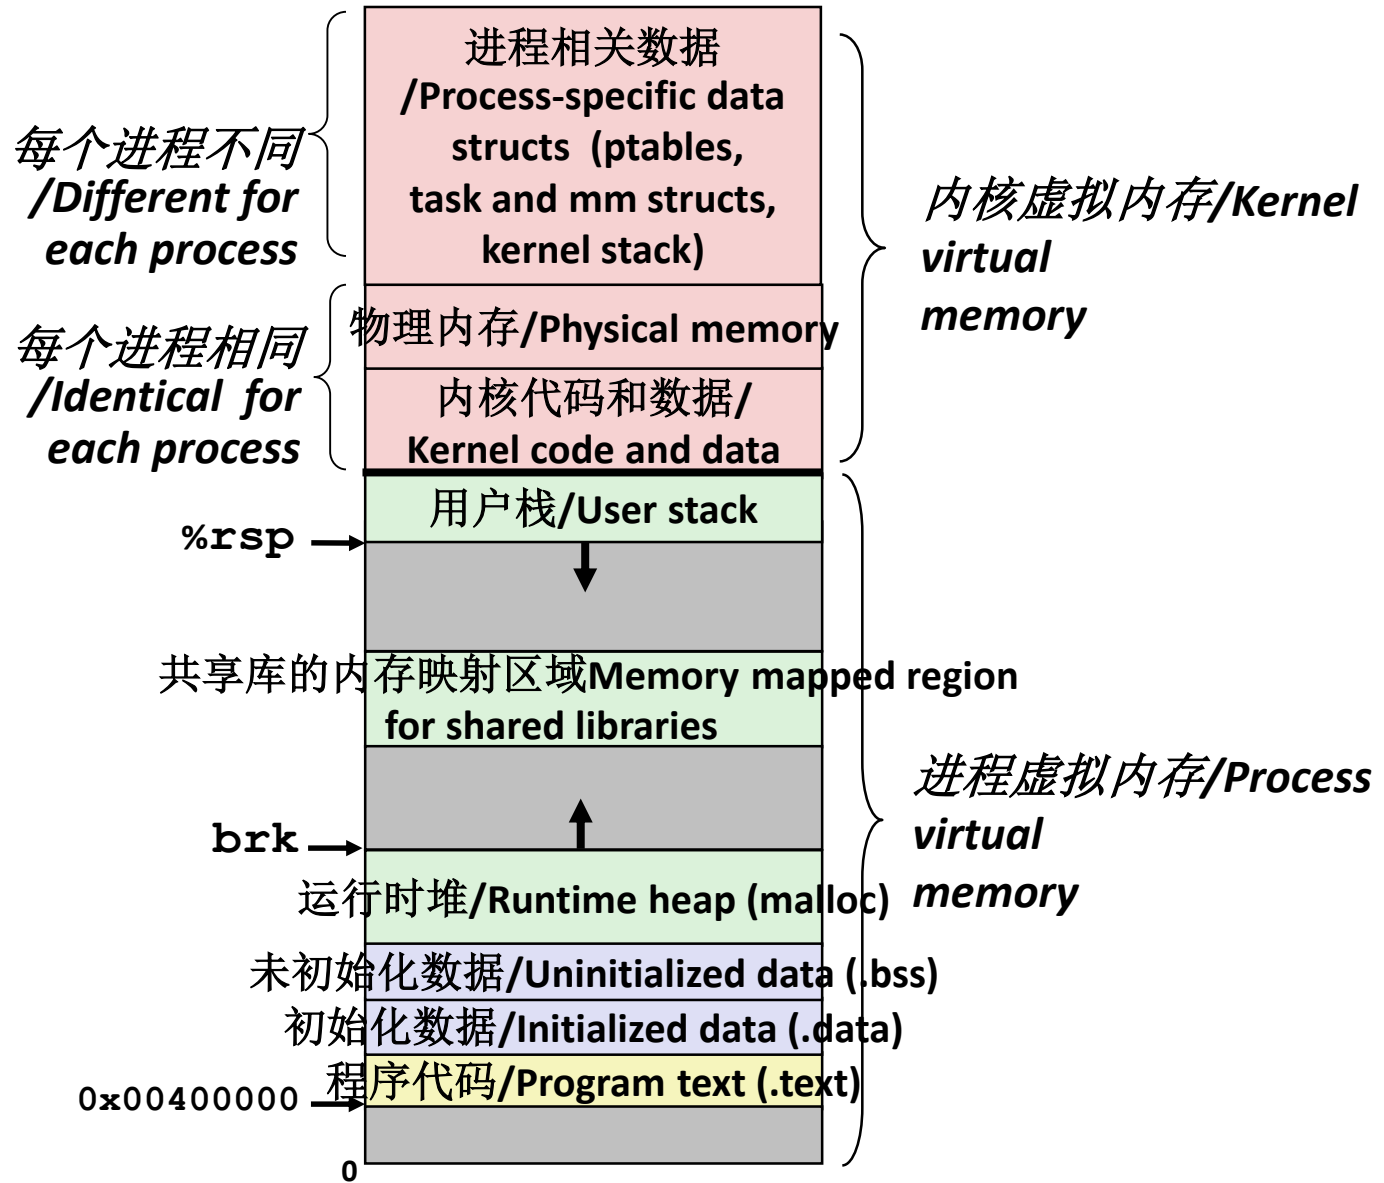


## ■ 观察/Observation

- 虚拟地址和物理地址中用于Cache索引的位是相同的/Bits that determine CI identical in virtual and physical address
- 地址翻译的同时可以进行Cache索引/Can index into cache while address translation taking place
- 通常情况下TLB会命中, PPN (Cache标记) 接下来会可用/Generally we hit in TLB, so PPN bits (CT bits) available next
- 虚拟索引, 物理标记/“Virtually indexed, physically tagged”
- Cache大小设计需要注意才能这样并行做/Cache carefully sized to make this possible

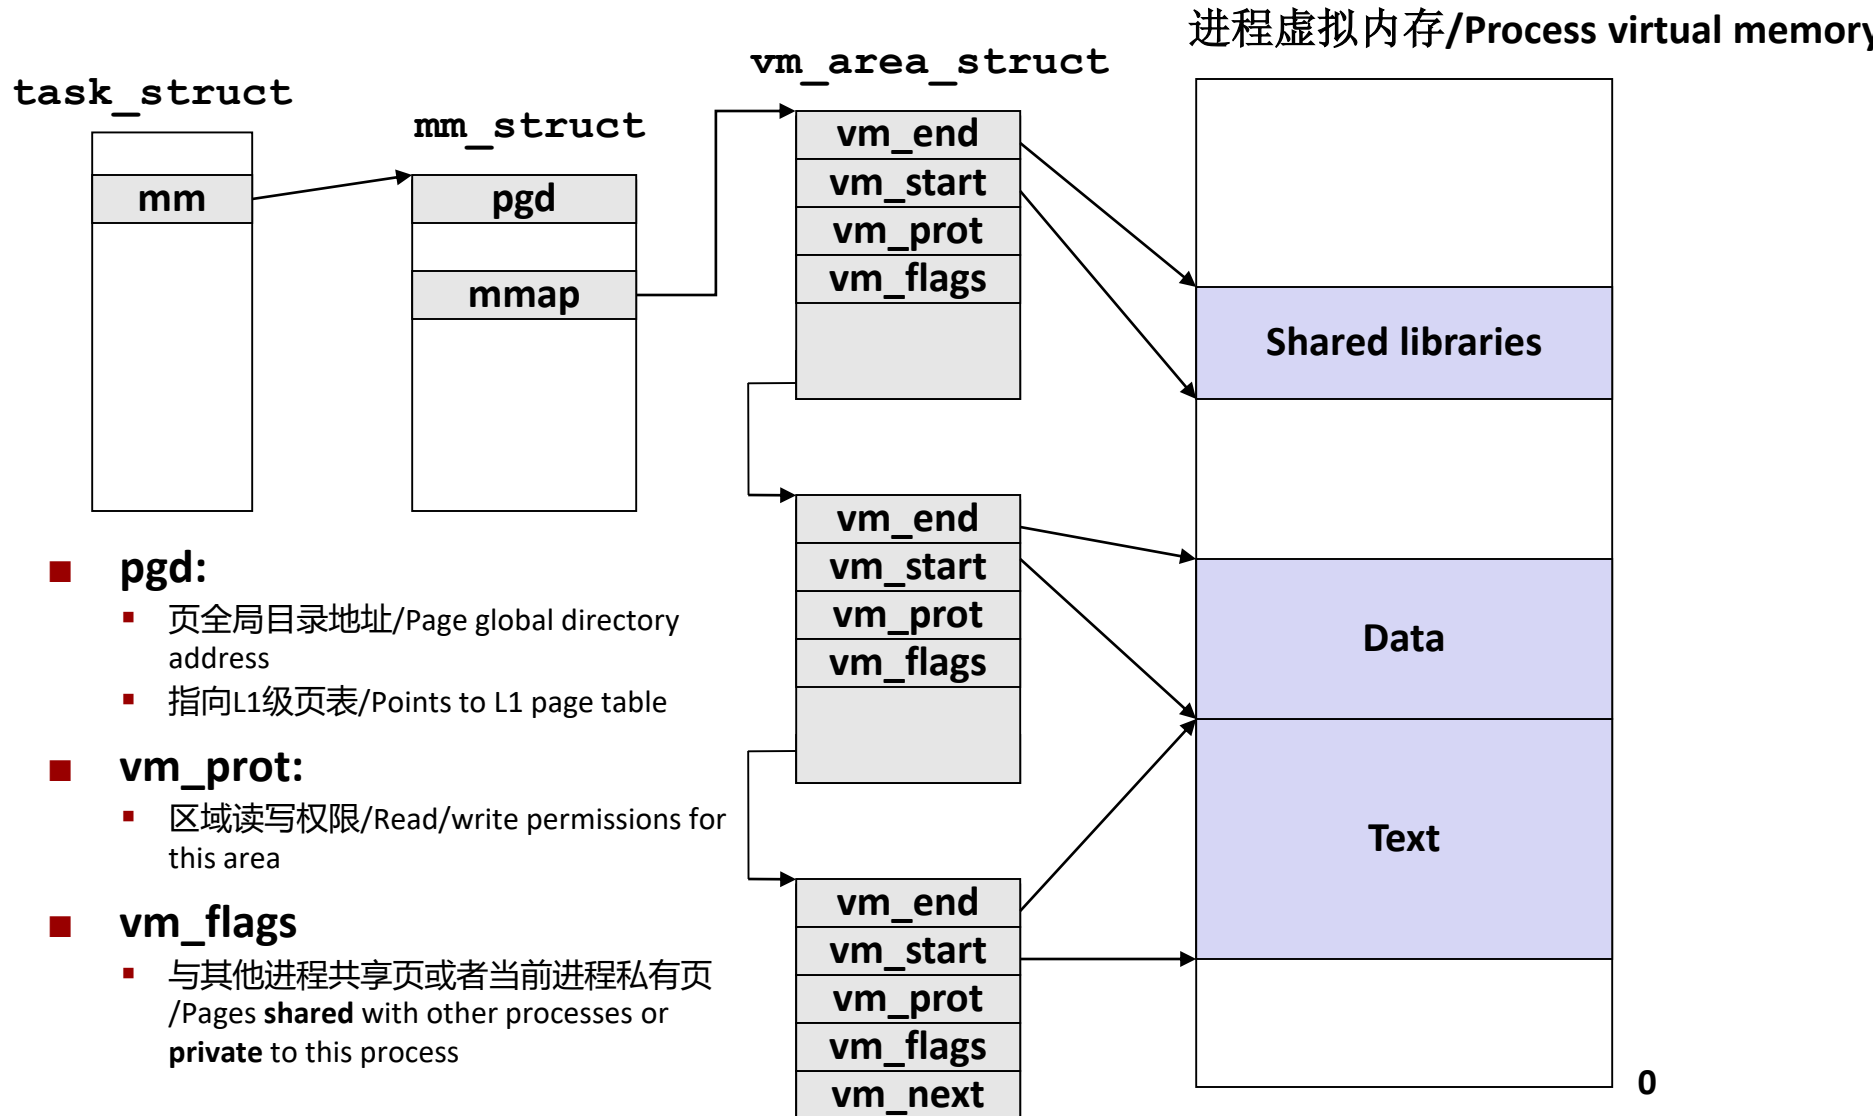


# Linux进程的虚拟地址空间/Virtual Address Space of a Linux Process



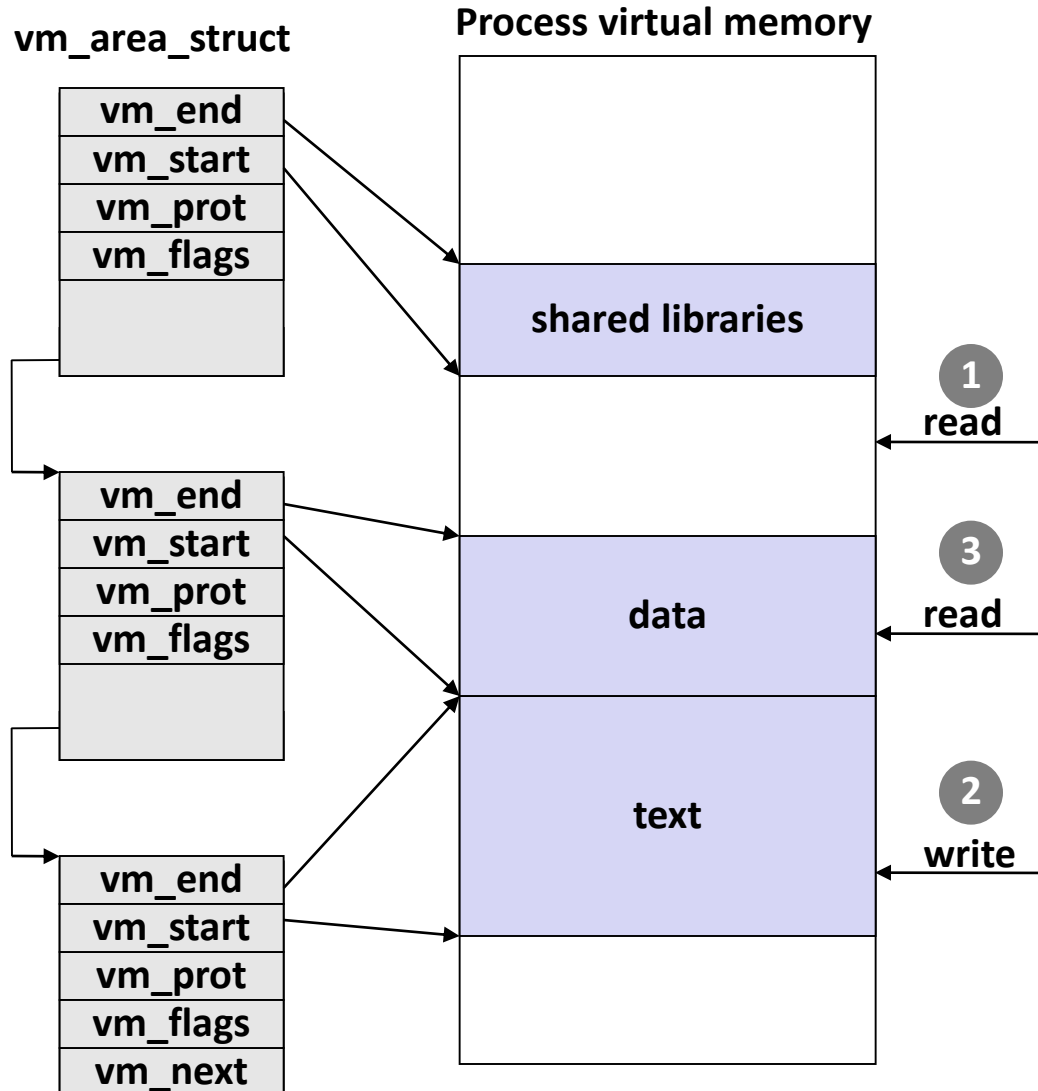


# Linux将虚拟内存组织为一些区域的集合/Linux Organizes VM as Collection of “Areas”





# Linux中的缺页中断处理/Linux Page Fault Handling



**段错误/Segmentation fault:**  
访问不存在的页/  
accessing a non-existing page

**普通缺页中断/  
Normal page fault**

**保护异常/Protection exception:**  
例如，对只读页进行写操作  
/e.g., violating permission by  
writing to a read-only page (Linux  
reports as Segmentation fault)



# 内容提纲/Today

- 简单内存系统示例/Simple memory system example
- 实例：Core i7/Linux内存系统/Case study: Core i7/Linux memory system
- 内存映射/Memory mapping



# 内存映射/Memory Mapping

- VM区域初始化是将其与磁盘对象相关联/VM areas initialized by associating them with disk objects.
  - 这一过程称为内存映射/Process is known as **memory mapping**.
- 区域可以备份/Area can be **backed by** (例如从以下获得初始值/i.e., get its initial values from) :
  - 磁盘上的**常规文件/Regular file** on disk (例如一个可执行目标文件/e.g., an executable object file)
    - 通过文件的节初始化页中数据/Initial page bytes come from a section of a file
  - **匿名文件/Anonymous file** (e.g., nothing)
    - 第一次缺页时分配一个填充为0的物理页/First fault will allocate a physical page full of 0's (**demand-zero page**)
    - 页面被写之后就和其他页一样/Once the page is written to (**dirtied**), it is like any other page
- 脏页会在内存和一个特殊的交换文件之间来回拷贝/Dirty pages are copied back and forth between memory and a special **swap file**.



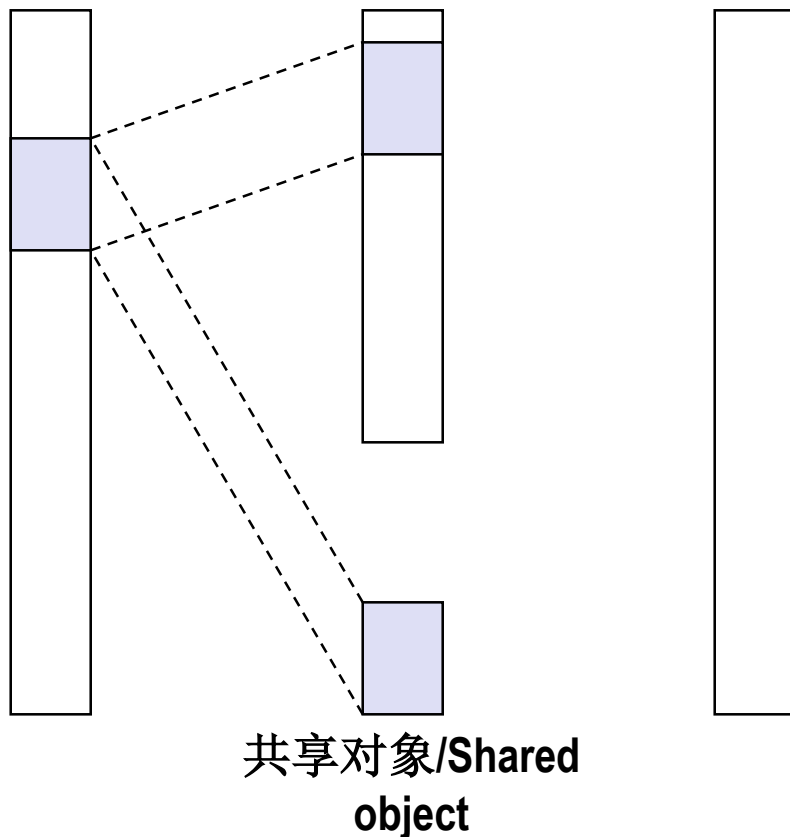
# 共享重回顾：共享对象/Sharing Revisited: Shared Objects

进程1虚拟内存/  
Process 1  
virtual memory

物理内存/  
Physical  
memory

进程2虚拟内存/  
Process 2  
virtual memory

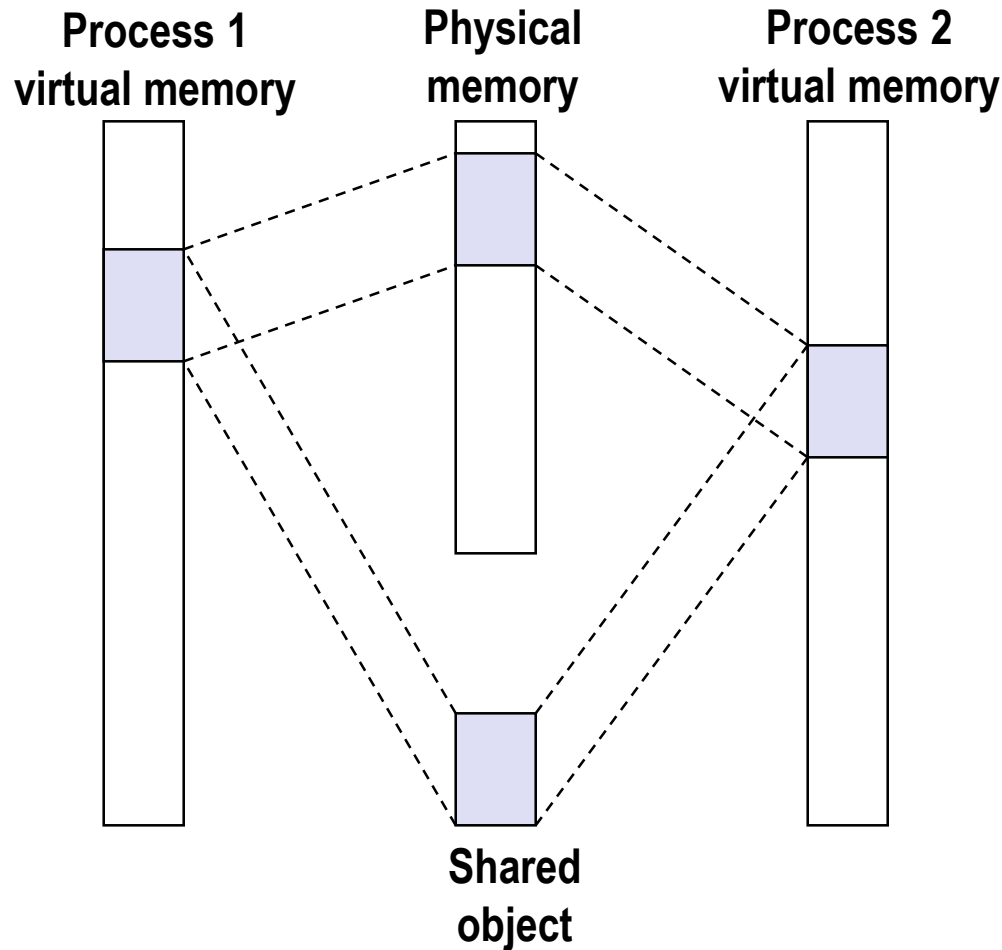
- 进程1映射共享对象/Process 1 maps the shared object.







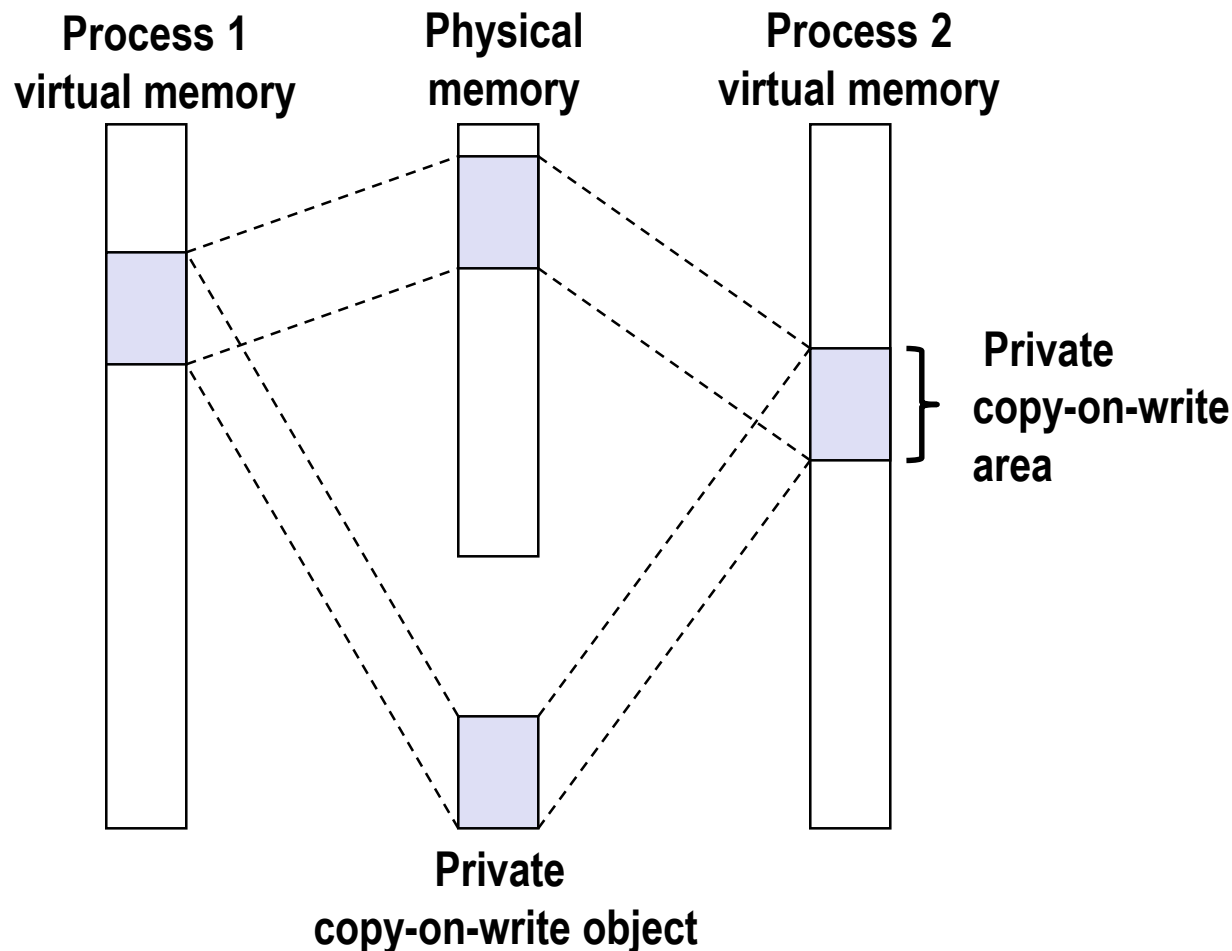
# 共享重回顾：共享对象/Sharing Revisited: Shared Objects



- 进程2映射共享对象/Process 2 maps the shared object.
- 注意虚拟地址如何不同/Notice how the virtual addresses can be different.



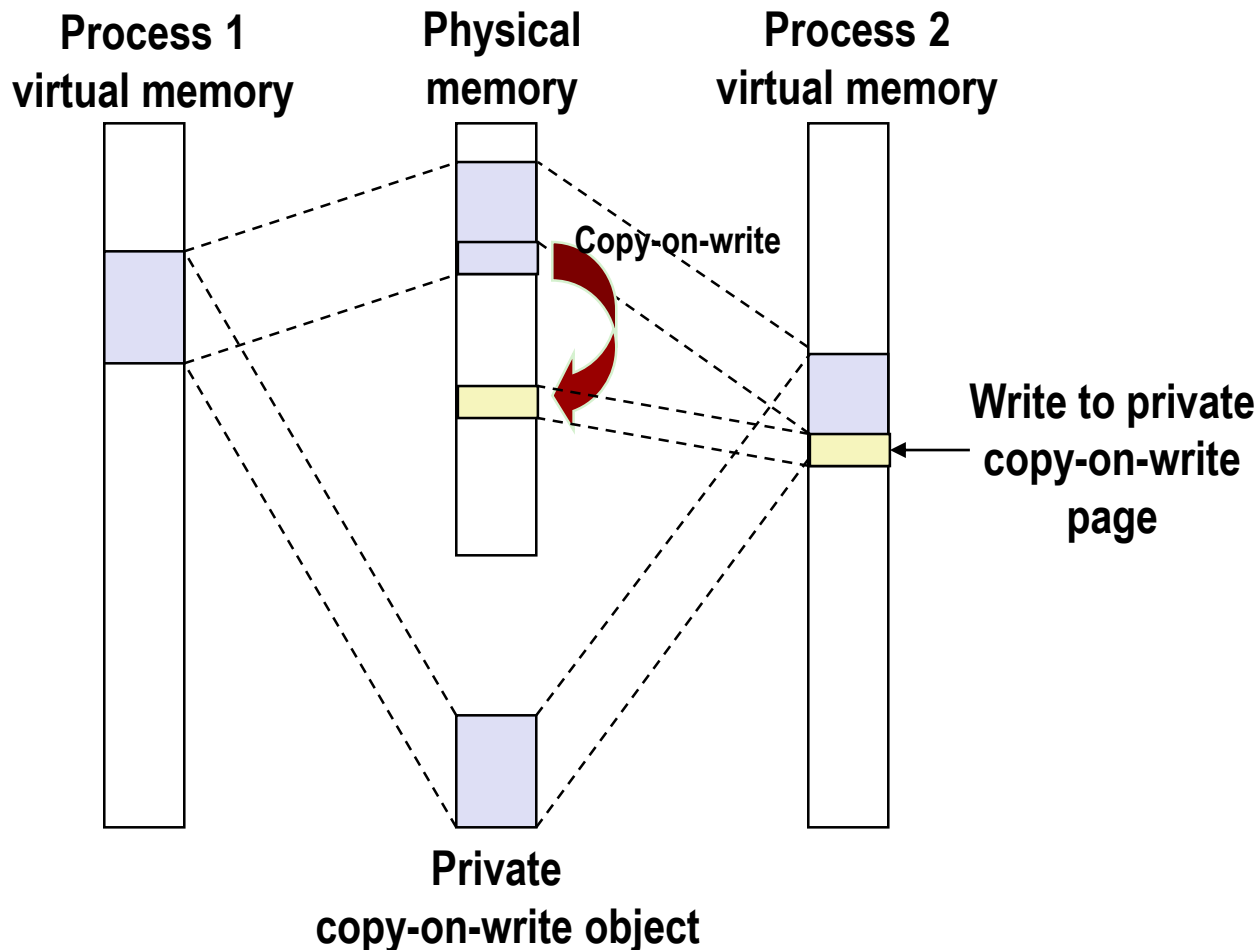
## 共享重回顾:私有写时拷贝对象/Sharing Revisited: Private Copy-on-write (COW) Objects



- 两个进程映射了一个私有写时拷贝对象  
/Two processes mapping a **private copy-on-write (COW)** object.
- 区域被标记为私有写时拷贝/Area flagged as private copy-on-write
- 私有区域的PTE被标记位只读/ PTEs in private areas are flagged as read-only



## 共享重回顾:私有写时拷贝对象/ Sharing Revisited: Private Copy-on-write (COW) Objects



- 写私有页会触发保护异常/Instruction writing to private page triggers protection fault.
- 处理程序创建一个新的R/W页/Handler creates new R/W page.
- 处理程序返回后重新执行指令/Instruction restarts upon handler return.
- 尽可能延迟拷贝操作/Copying deferred as long as possible!

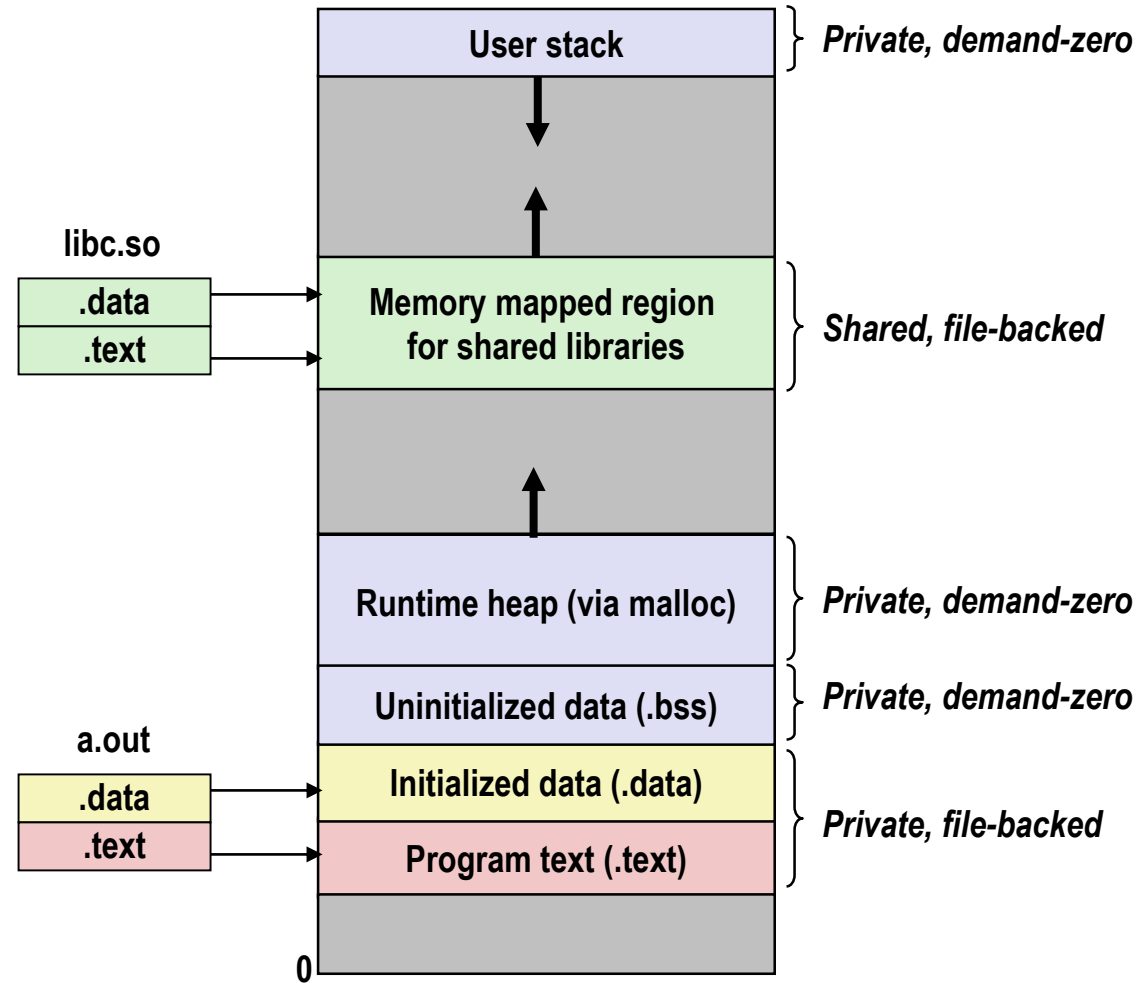


# fork函数重回顾/The fork Function Revisited

- **VM和内存映射解释了fork如何为每个进程设置私有空间/VM and memory mapping explain how fork provides private address space for each process.**
- **为新进程创建虚拟地址/To create virtual address for new new process**
  - 创建完全与现有的完全相同的内存数据结构/Create exact copies of current `mm_struct`, `vm_area_struct`, and page tables.
  - 每个进程都将其标记为只读/Flag each page in both processes as read-only
  - 在两个进程空间中的`vm_area_struct` 设置为COW/Flag each `vm_area_struct` in both processes as private COW
- **返回时，每个进程有完全相同的虚拟内存/On return, each process has exact copy of virtual memory**
- **后续写操作会因为COW创建新的页/Subsequent writes create new pages using COW mechanism.**



# execve重回顾/The execve Function Revisited



- 在现有进程用 `execve` 加载并运行一个新的程序 `a.out` / To load and run a new program `a.out` in the current process using `execve`:
- 释放旧区域的相关数据结构 / Free `vm_area_struct`'s and page tables for old areas
- 创建新区域的相关数据结构 / Create `vm_area_struct`'s and page tables for new areas
  - Programs and initialized data backed by object files.
  - `.bss` and stack backed by anonymous files.
- 设置 PC 到 `.text` / Set PC to entry point in `.text`
  - Linux will fault in code and data pages as needed.



# 用户级内存映射/User-Level Memory Mapping

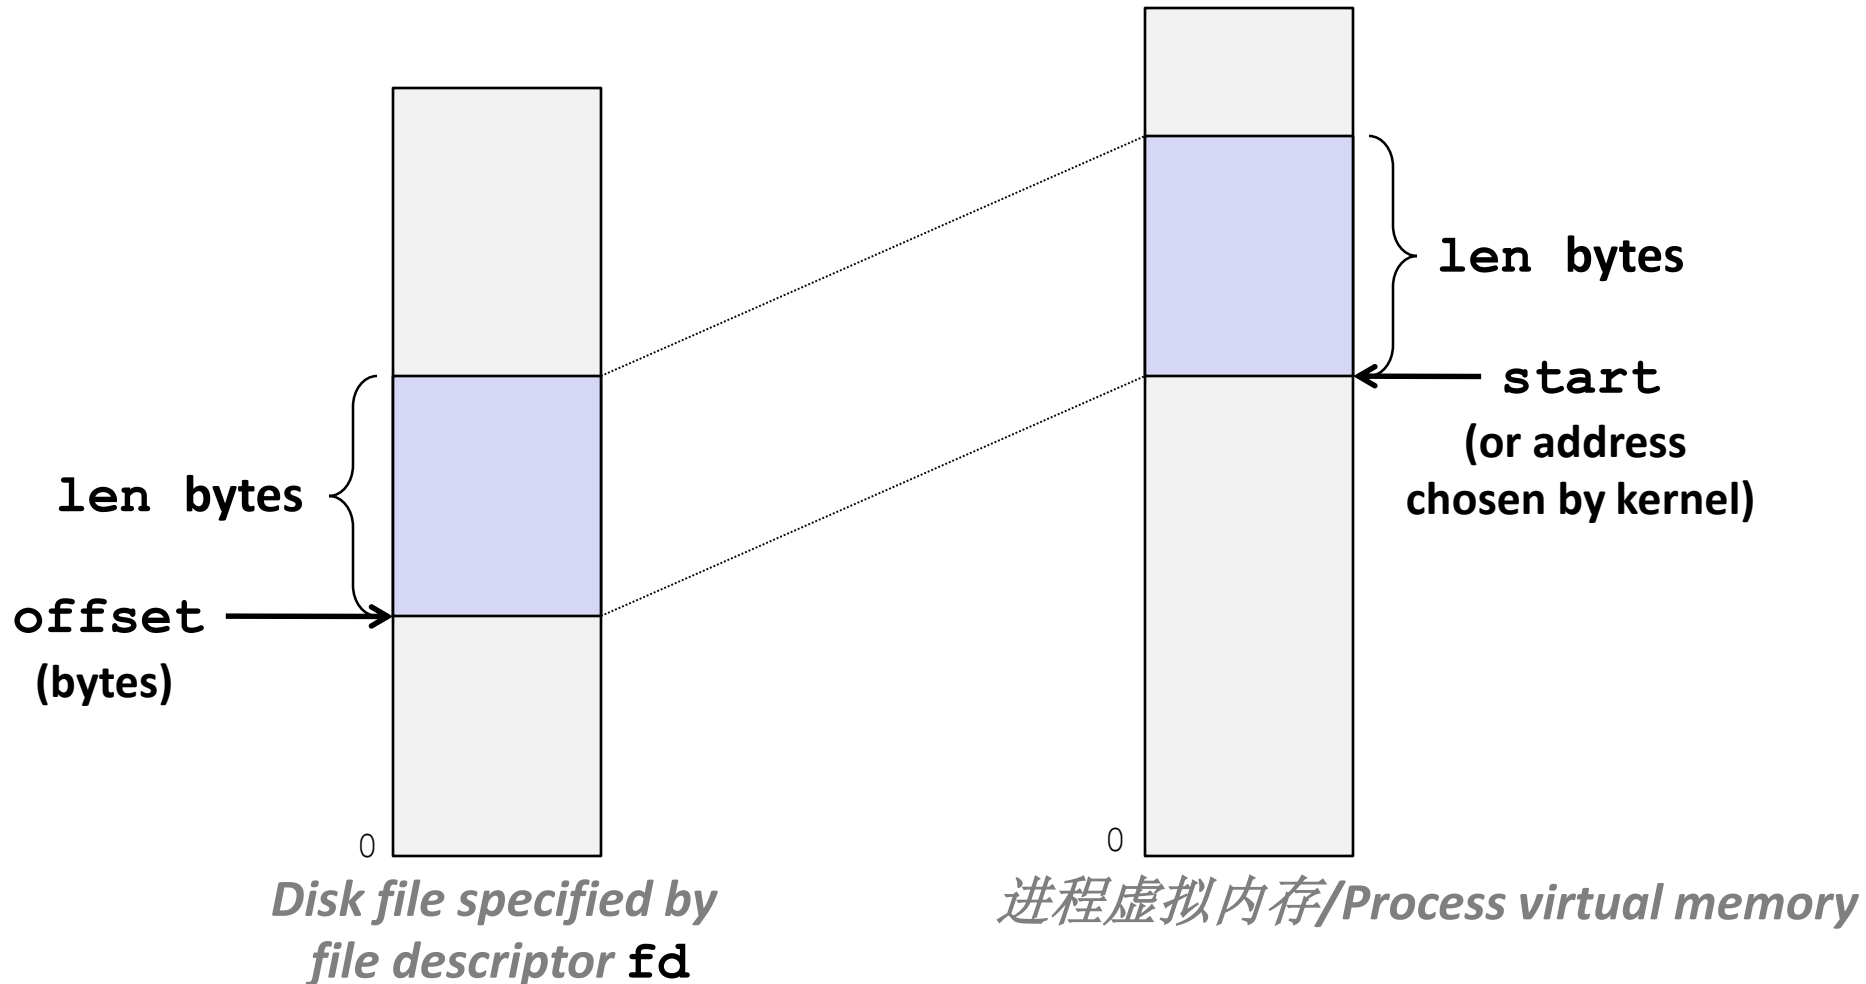
```
void *mmap(void *start, int len,  
           int prot, int flags, int fd, int offset)
```

- 将文件描述符`fd`中偏移量`offset`开始的长度为`len`的字节映射到地址`start`/Map `len` bytes starting at offset `offset` of the file specified by file description `fd`, preferably at address `start`
  - `start`: may be 0 for “pick an address”/有可能是0
  - `prot`: `PROT_READ`, `PROT_WRITE`, ...
  - `flags`: `MAP_ANON`, `MAP_PRIVATE`, `MAP_SHARED`, ...
- 返回一个映射区域的开始地址（有可能不是`start`）/Return a pointer to start of mapped area (may not be `start`)



# 用户级内存映射/User-Level Memory Mapping

```
void *mmap(void *start, int len,  
           int prot, int flags, int fd, int offset)
```





# 示例：使用mmap拷贝文件/Example: Using mmap to Copy Files

- 不用传输数据到用户空间来，就可以将一个文件拷贝到stdout/Copying a file to stdout without transferring data to user space .

```
#include "csapp.h"

void mmapcopy(int fd, int size)
{

    /* Ptr to memory mapped area */
    char *bufp;

    bufp = Mmap(NULL, size,
                PROT_READ,
                MAP_PRIVATE,
                fd, 0);
    Write(1, bufp, size);
    return;
}
```

mmapcopy.c

```
/* mmapcopy driver */
int main(int argc, char **argv)
{
    struct stat stat;
    int fd;

    /* Check for required cmd line arg */
    if (argc != 2) {
        printf("usage: %s <filename>\n",
               argv[0]);
        exit(0);
    }

    /* Copy input file to stdout */
    fd = Open(argv[1], O_RDONLY, 0);
    Fstat(fd, &stat);
    mmapcopy(fd, stat.st_size);
    exit(0);
}
```

mmapcopy.c