



Thread-Level Parallelism

线程级并行

100076202: 计算机系统导论

任课教师:

计卫星 宿红毅 张艳

原作者:

Randal E. Bryant and David R. O'Hallaron



Carnegie
Mellon
University



内容提纲/Today

- **并行计算硬件/Parallel Computing Hardware**
 - 多核/Multicore
 - 一个芯片上多个独立的处理器/Multiple separate processors on single chip
 - 超线程/Hyperthreading
 - 单个核上执行多个线程/Efficient execution of multiple threads on single core
- **线程级并行/Thread-Level Parallelism**
 - 将程序分为多个独立的任务/Splitting program into independent tasks
 - Example 1: Parallel summation
 - 分治并行/Divide-and conquer parallelism
 - Example 2: Parallel quicksort
- **一致性模型/Consistency Models**
 - 当多个线程读写共享状态时会发生什么/What happens when multiple threads are reading & writing shared state

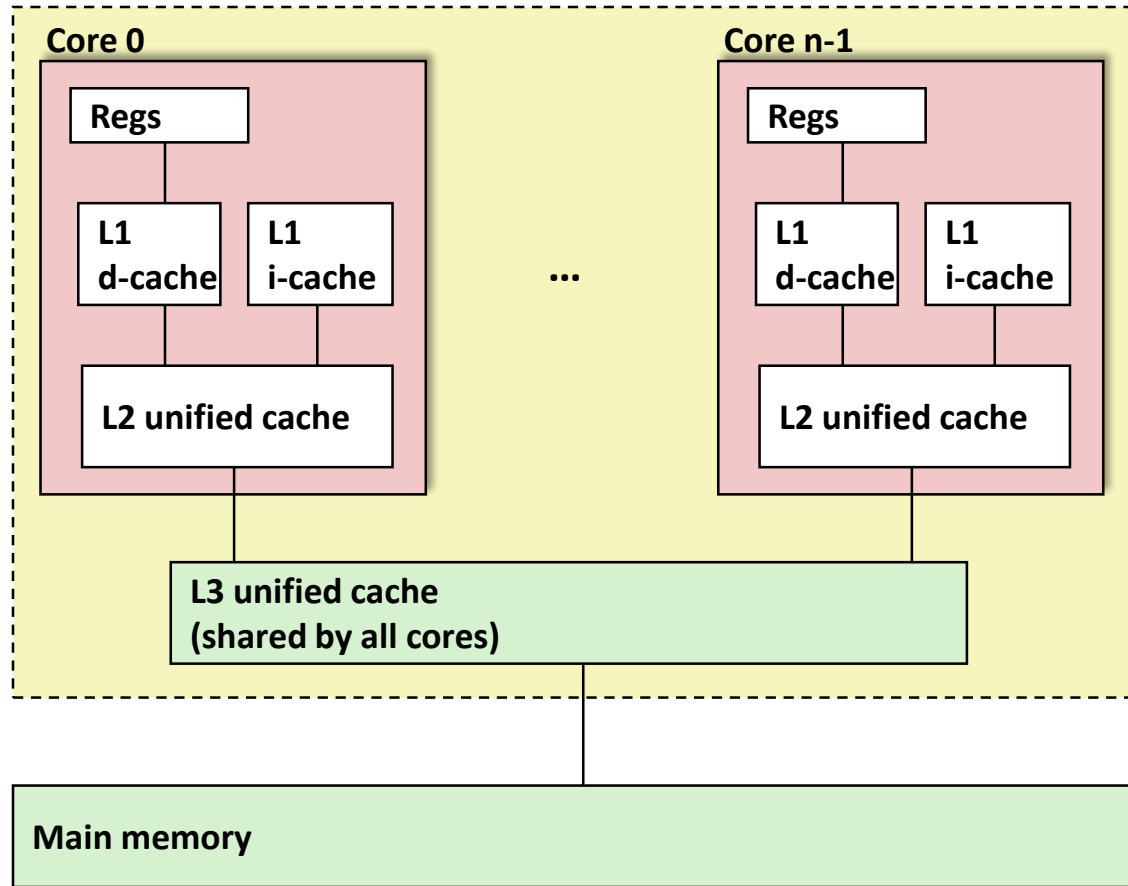


挖掘并行执行/**Exploiting parallel execution**

- **目前我们还是在用线程处理I/O延迟/So far, we've used threads to deal with I/O delays**
 - 例如为每个客户端设置一个线程以防止互相延迟/e.g., one thread per client to prevent one from delaying another
- **多核处理器/超线程CPU提供了其他的可能/Multi-core/Hyperthreaded CPUs offer another opportunity**
 - 将任务分布给线程并行执行/Spread work over threads executing in parallel
 - 如果有很多独立的任务则自动实现/Happens automatically, if many independent tasks
 - 例如有很多程序或者服务很多客户端/e.g., running many applications or serving many clients
 - 也可以编写代码实现一个任务的加速运行/Can also write code to make one big task go faster
 - 按照多个子任务并行组织/by organizing it as multiple parallel sub-tasks



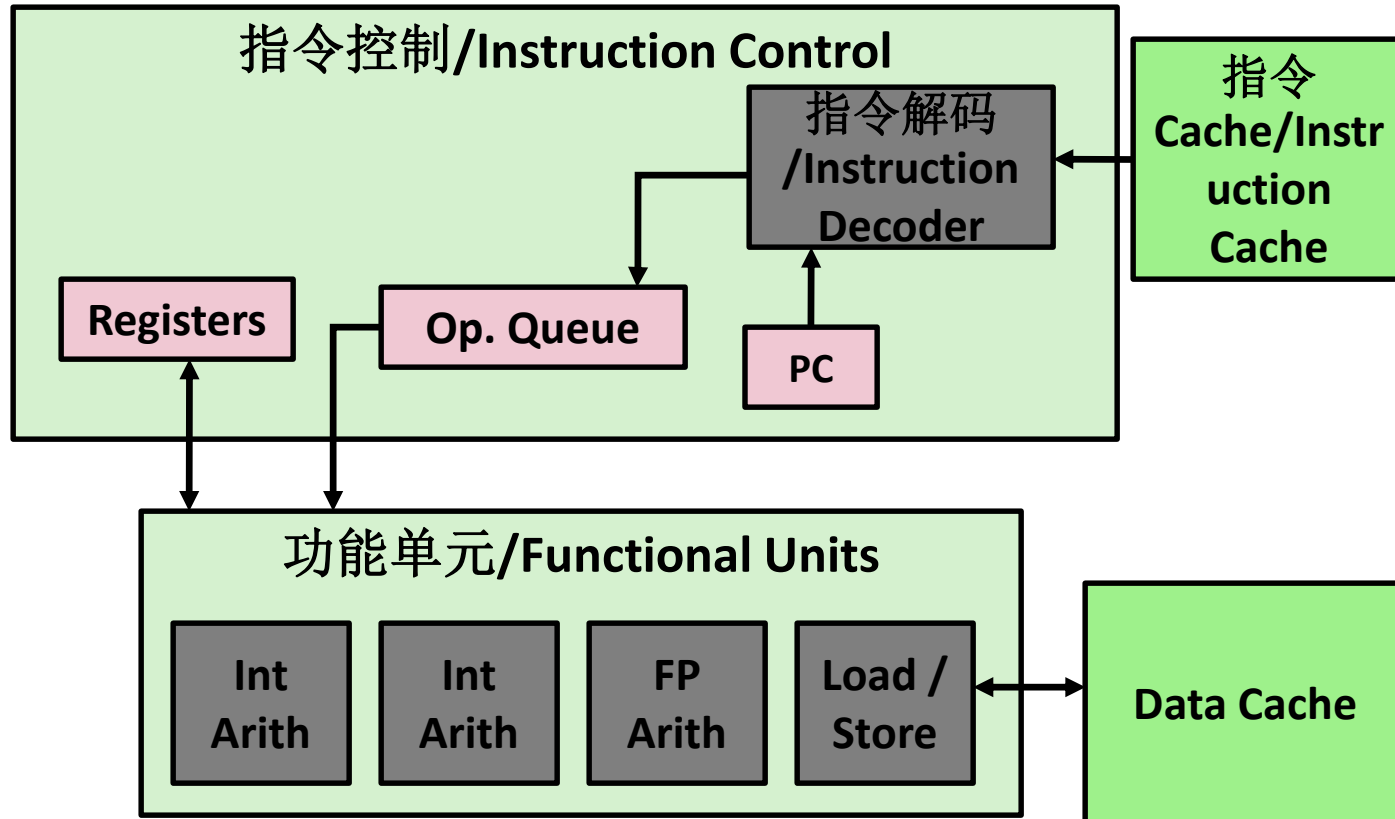
典型多核处理器/Typical Multicore Processor



- 多个处理器在运行过程中对内存有一致的视图/Multiple processors operating with coherent view of memory

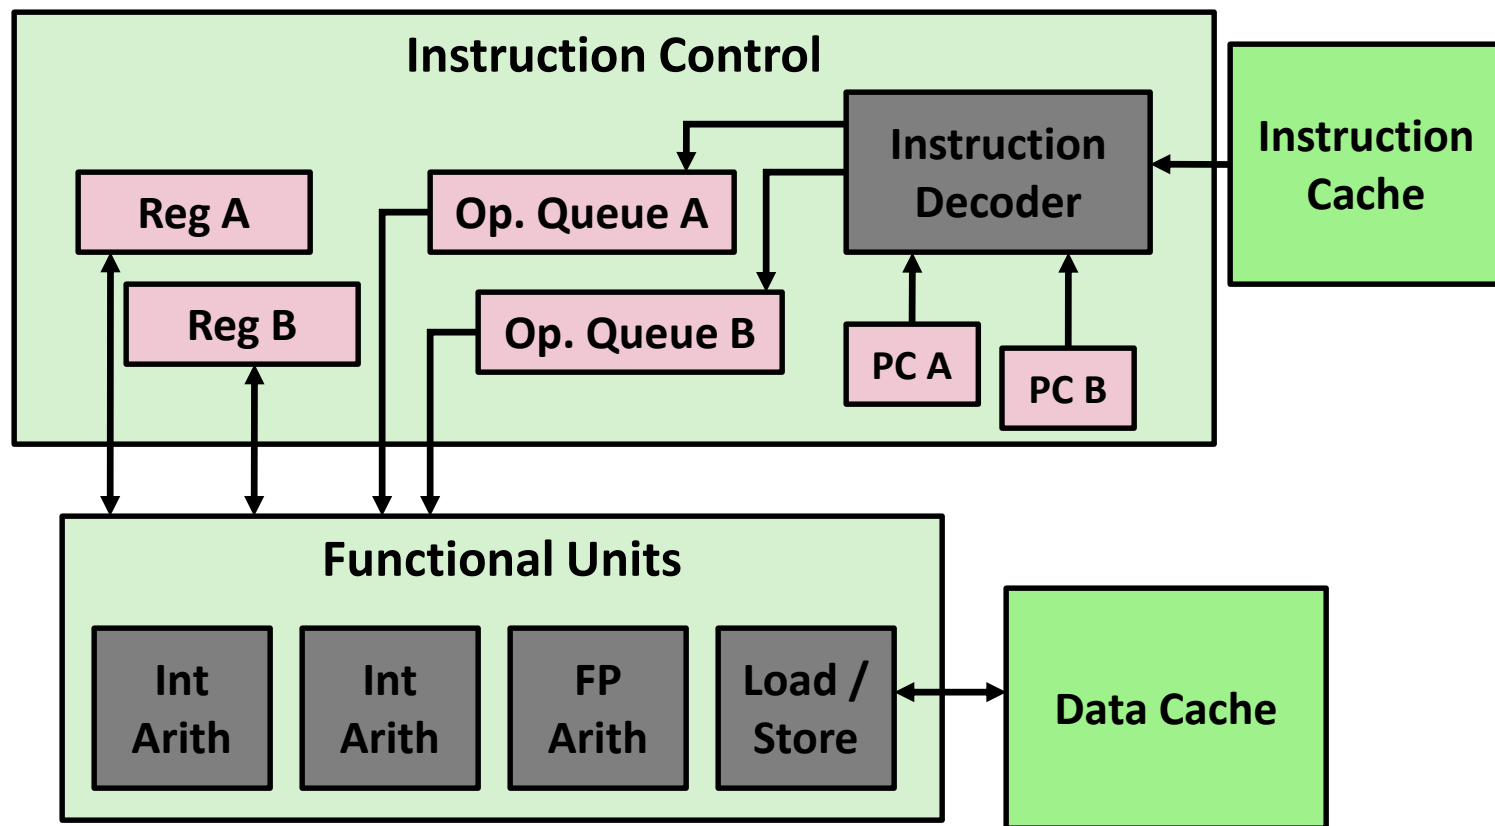


乱序处理器结构/Out-of-Order Processor Structure



- 指令控制自动将程序转为操作流/Instruction control dynamically converts program into stream of operations
- 操作映射到功能单元上并行执行/Operations mapped onto functional units to execute in parallel

超线程实现/Hyperthreading Implementation



- 复制指令控制以处理K个指令流/Replicate enough instruction control to process K instruction streams
- 所有寄存器有K个副本/K copies of all registers
- 共享功能单元/Share functional units



测试机器/Benchmark Machine

- 从/`proc/cpuinfo`获取机器信息/Get data about machine from `/proc/cpuinfo`
- 机器Shark/Shark Machines
 - Intel Xeon E5520 @ 2.27 GHz
 - Nehalem, ca. 2010
 - 8 Cores
 - Each can do 2x hyperthreading



例1:并行求和/Example 1: Parallel Summation

- 对 $0, \dots, n-1$ 求和/Sum numbers $0, \dots, n-1$
 - Should add up to $((n-1)*n)/2$
- 将 $1, \dots, n-1$ 划分为 t 个区间/Partition values $1, \dots, n-1$ into t ranges
 - 每个区间有 $\lfloor n/t \rfloor$ 个值/ $\lfloor n/t \rfloor$ values in each range
 - 每个线程处理一个区间/Each of t threads processes 1 range
 - 简单起见, 假设 n 是 t 的整数倍/For simplicity, assume n is a multiple of t
- 考虑一下多个线程在不同的区间上工作的不同方式
/Let's consider different ways that multiple threads might work on their assigned ranges in parallel



尝试1/First attempt: psum-mutex

- 简单方法：线程将求和结果合并到受mutex保护的全局变量上/Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
void *sum_mutex(void *vargp); /* Thread routine */

/* Global shared variables */
long gsum = 0;                /* Global sum */
long nelems_per_thread;      /* Number of elements to sum */
sem_t mutex;                  /* Mutex to protect global sum */

int main(int argc, char **argv)
{
    long i, nelems, log_nelems, nthreads, myid[MAXTHREADS];
    pthread_t tid[MAXTHREADS];

    /* Get input arguments */
    nthreads = atoi(argv[1]);
    log_nelems = atoi(argv[2]);
    nelems = (1L << log_nelems);
    nelems_per_thread = nelems / nthreads;
    sem_init(&mutex, 0, 1);
```

psum-mutex.c



psum-mutex (cont)

- Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
/* Create peer threads and wait for them to finish */
for (i = 0; i < nthreads; i++) {
    myid[i] = i;
    Pthread_create(&tid[i], NULL, sum_mutex, &myid[i]);
}
for (i = 0; i < nthreads; i++)
    Pthread_join(tid[i], NULL);

/* Check final answer */
if (gsum != (nelems * (nelems-1))/2)
    printf("Error: result=%ld\n", gsum);

exit(0);
}
```

psum-mutex.c



psum-mutex Thread Routine/线程函数

- Simplest approach: Threads sum into a global variable protected by a semaphore mutex.

```
/* Thread routine for psum-mutex.c */
void *sum_mutex(void *vargp)
{
    long myid = *((long *)vargp);          /* Extract thread ID */
    long start = myid * nelems_per_thread; /* Start element index */
    long end = start + nelems_per_thread;  /* End element index */
    long i;

    for (i = start; i < end; i++) {
        P(&mutex);
        gsum += i;
        V(&mutex);
    }
    return NULL;
}
```

psum-mutex.c



psum-mutex Performance/性能

- Shark machine with 8 cores, $n=2^{31}$

Threads (Cores)	1 (1)	2 (2)	4 (4)	8 (8)	16 (8)
psum-mutex (secs)	51	456	790	536	681

- 意外的结果/Nasty surprise:
 - 单个线程非常慢/Single thread is very slow
 - 核越多越慢/Gets slower as we use more cores



尝试2:/Next Attempt: psum-array

- 不同的线程归并到不同的数组元素/Peer thread *i* sums into global array element `psum[i]`
- 主线程等待其他线程完成，并对数组元素进行求和/Main waits for theads to finish, then sums elements of `psum`
- 消除了基于mutex的同步/Eliminates need for mutex synchronization

```
/* Thread routine for psum-array.c */
void *sum_array(void *vargp)
{
    long myid = *((long *)vargp);          /* Extract thread ID */
    long start = myid * nelems_per_thread; /* Start element index */
    long end = start + nelems_per_thread;  /* End element index */
    long i;

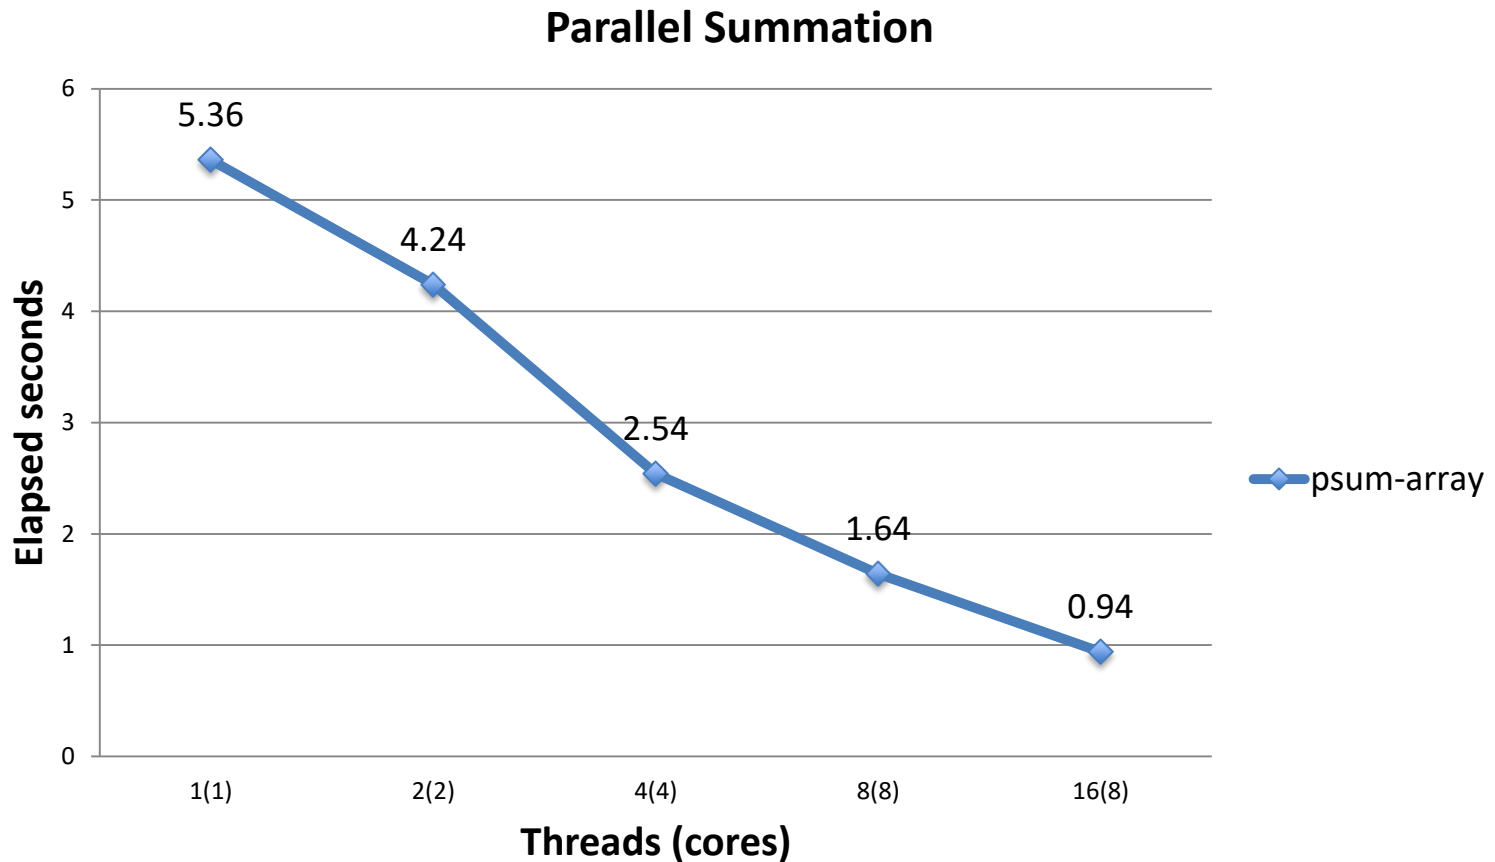
    for (i = start; i < end; i++) {
        psum[myid] += i;
    }
    return NULL;
}
```

psum-array.c



psum-array Performance/性能

- 比psum-mutex快一个量级 / Orders of magnitude faster than psum-mutex





尝试3:/Next Attempt: psum-local

- 每个线程求和归并到局部变量/Reduce memory references by having peer thread i sum into a local variable (register)

```
/* Thread routine for psum-local.c */
void *sum_local(void *vargp)
{
    long myid = *((long *)vargp);          /* Extract thread ID */
    long start = myid * nelems_per_thread; /* Start element index */
    long end = start + nelems_per_thread;  /* End element index */
    long i, sum = 0;

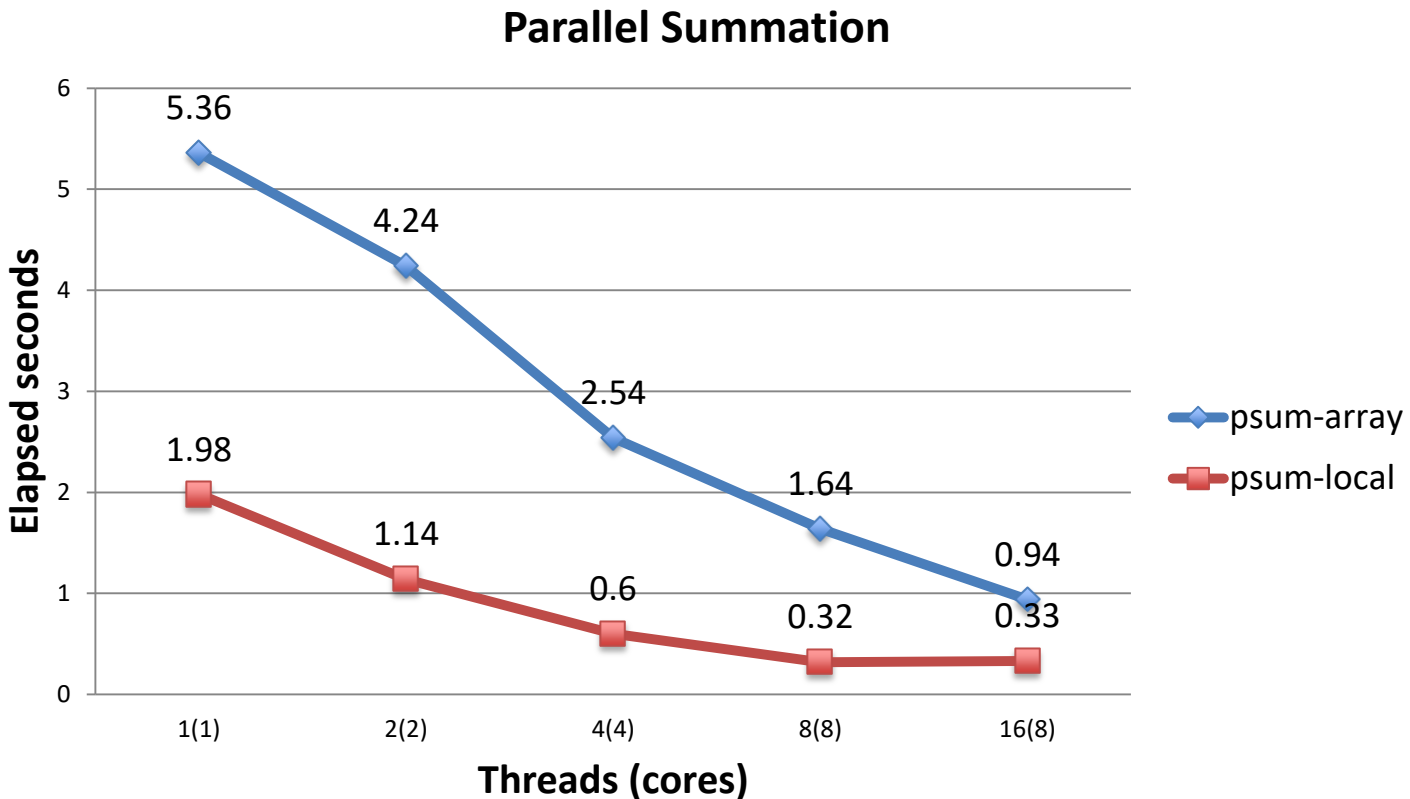
    for (i = start; i < end; i++) {
        sum += i;
    }
    psum[myid] = sum;
    return NULL;
}
```

psum-local.c



psum-local Performance/性能

- 比psum-array性能有了大幅提升/Significantly faster than psum-array





表征并行程序性能/Characterizing Parallel Program Performance

- p 表示处理器核数, T_k 表示使用 k 个核运行的时间/ p processor cores, T_k is the running time using k cores
- 加速比定义/Def. **Speedup**: $S_p = T_1 / T_p$
 - S_p 表示相对加速比, 如果 T_1 是并行版本代码在1个核上的运行时间/ S_p is *relative speedup* if T_1 is running time of parallel version of the code running on 1 core.
 - S_p 表示绝对加速比, 如果 T_1 是串行版本代码在1个核上的运行时间/ S_p is *absolute speedup* if T_1 is running time of sequential version of code running on 1 core.
 - 绝对加速比能够更加真实的表示并行加速收益/Absolute speedup is a much truer measure of the benefits of parallelism.
- 并行效率定义/Def. **Efficiency**: $E_p = S_p / p = T_1 / (pT_p)$
 - 是(0, 100]之间的一个百分比/Reported as a percentage in the range (0, 100].
 - 测度的是并行带来的额外开销/Measures the overhead due to parallelization



psum-local性能/Performance of psum-local

Threads (t)	1	2	4	8	16
Cores (p)	1	2	4	8	8
Running time (T_p)	1.98	1.14	0.60	0.32	0.33
Speedup (S_p)	1	1.74	3.30	6.19	6.00
Efficiency (E_p)	100%	87%	82%	77%	75%

- 并行效率还可以，但是不是很好/Efficiencies OK, not great
- 我们的例子比较容易并行/Our example is easily parallelizable
- 实际代码更加难以并行/Real codes are often much harder to parallelize
 - 例如后面的quicksort例子/e.g., parallel quicksort later in this lecture



阿姆达尔定律/Amdahl's Law

- Gene Amdahl (Nov. 16, 1922 – Nov. 10, 2015)
- 描述了并行化的困难/Captures the difficulty of using parallelism to speed things up.
- 问题概述/Overall problem
 - T Total sequential time required/串行运行时间
 - p Fraction of total that can be sped up ($0 \leq p \leq 1$)/可并行加速比例
 - k Speedup factor/加速因子
- 最终的性能/Resulting Performance
 - $T_k = pT/k + (1-p)T$
 - 并行部分被加速 k 倍/Portion which can be sped up runs k times faster
 - 串行部分保持不动/Portion which cannot be sped up stays the same
 - 最短时间/Least possible running time:
 - $k = \infty$
 - $T_\infty = (1-p)T$



阿姆达尔定律举例/Amdahl's Law Example

■ 问题概述/Overall problem

- $T = 10$ Total time required/总的运行时间
- $p = 0.9$ Fraction of total which can be sped up/可并行部分加速比
- $k = 9$ Speedup factor/加速因子

■ 最终性能/Resulting Performance

- $T_9 = 0.9 * 10/9 + 0.1 * 10 = 1.0 + 1.0 = 2.0$
- 最短时间/Least possible running time:
 - $T_{\infty} = 0.1 * 10.0 = 1.0$

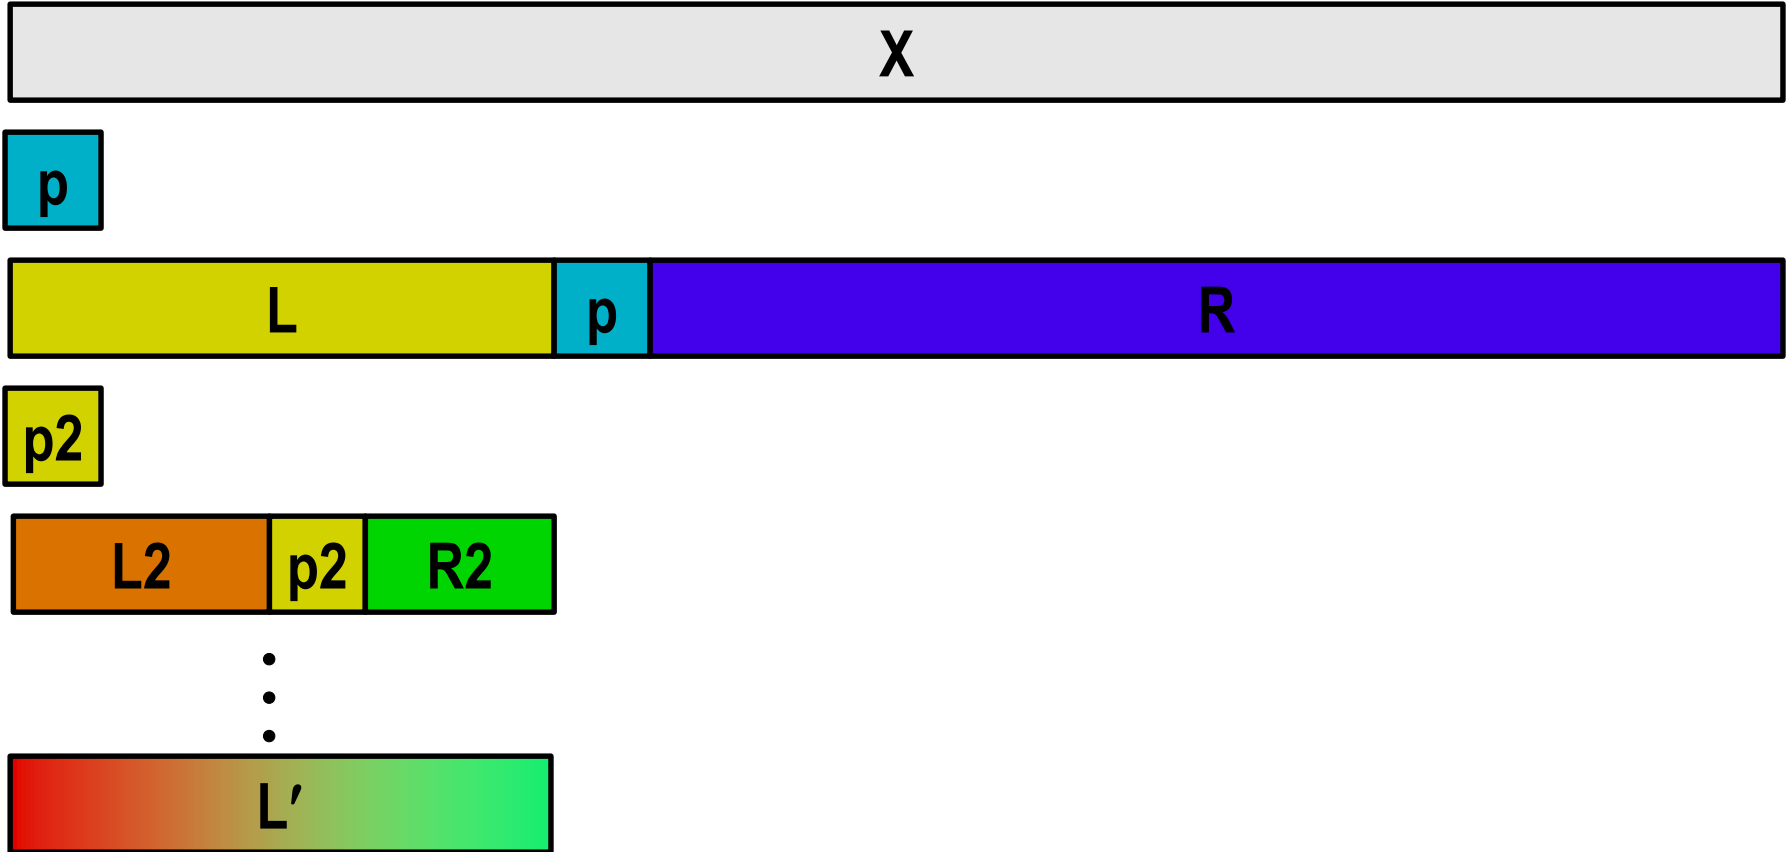


一个更复杂的例子/A More Substantial Example: Sort

- 对N个随机数排序/Sort set of N random numbers
- 多个可能的算法/Multiple possible algorithms
 - 使用quicksort的并行版本/Use parallel version of quicksort
- 对X数集的串行quicksort排序/Sequential quicksort of set of values X
 - 从X中选择枢轴p/Choose “pivot” p from X
 - 对X划分/Rearrange X into
 - L: Values $\leq p$
 - R: Values $\geq p$
 - 对L递归排序形成L'/Recursively sort L to get L'
 - 对R递归排序形成R'/Recursively sort R to get R'
 - 返回/Return L' : p : R'

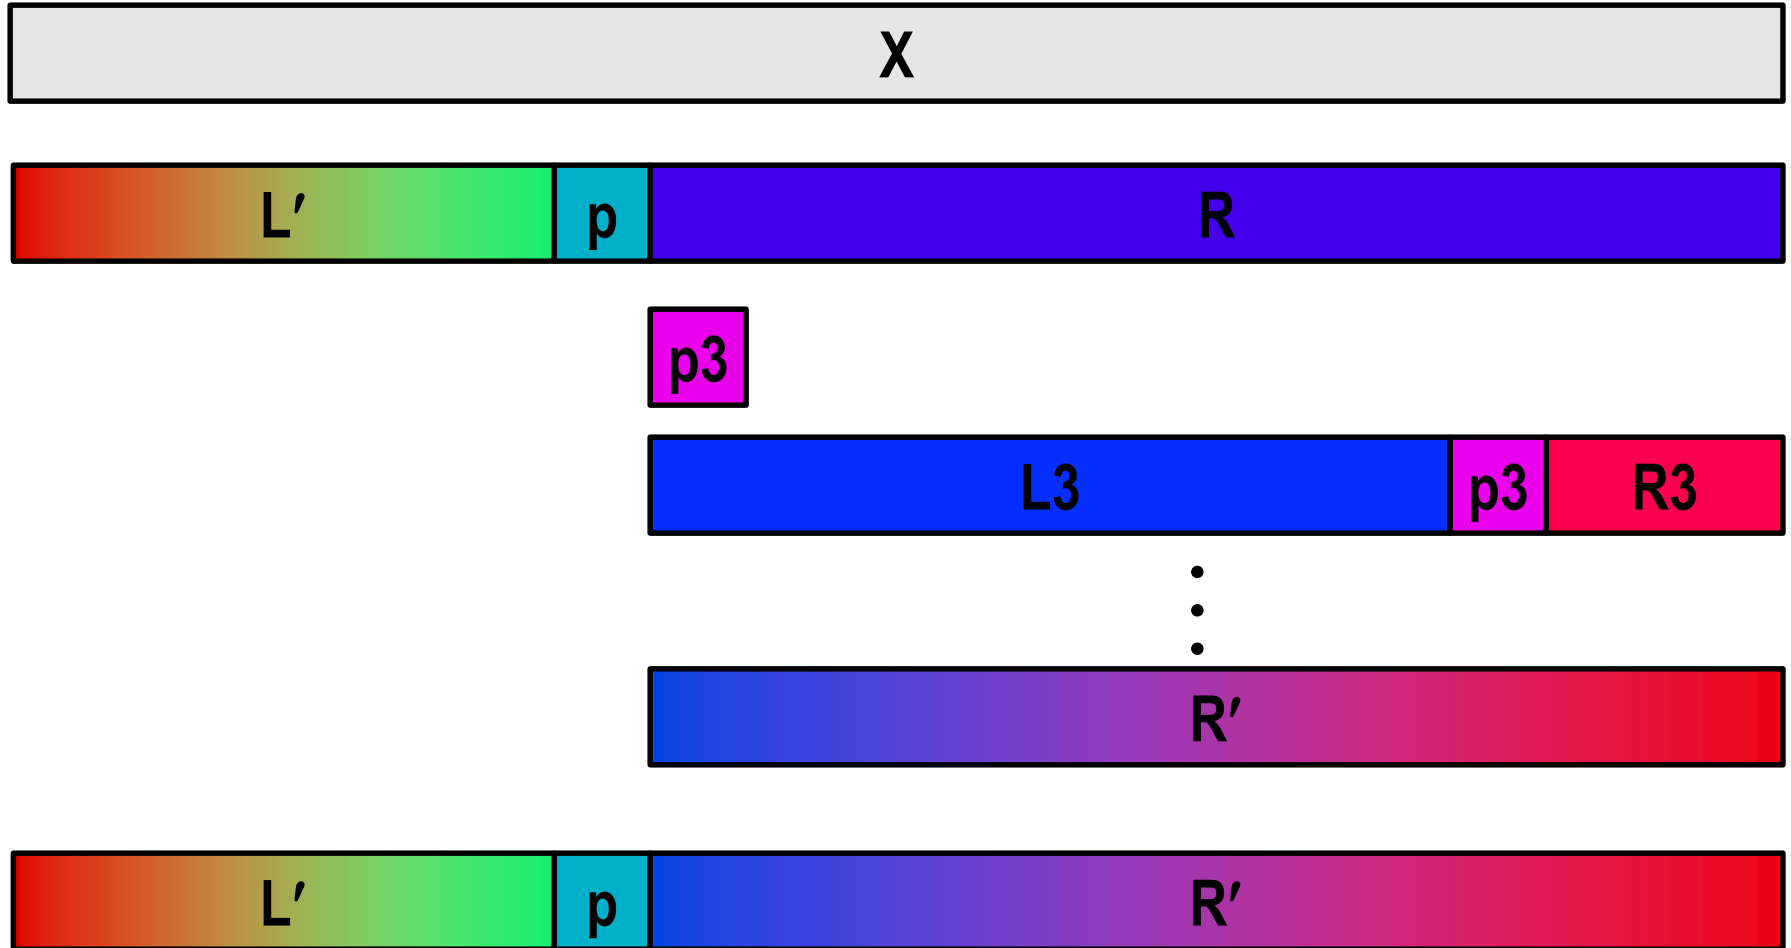


串行Quicksort 可视化/Sequential Quicksort Visualized





串行Quicksort 可视化/Sequential Quicksort Visualized





串行Quicksort 代码/Sequential Quicksort Code

```
void qsort_serial(data_t *base, size_t nele) {
    if (nele <= 1)
        return;
    if (nele == 2) {
        if (base[0] > base[1])
            swap(base, base+1);
        return;
    }

    /* Partition returns index of pivot */
    size_t m = partition(base, nele);
    if (m > 1)
        qsort_serial(base, m);
    if (nele-1 > m+1)
        qsort_serial(base+m+1, nele-m-1);
}
```

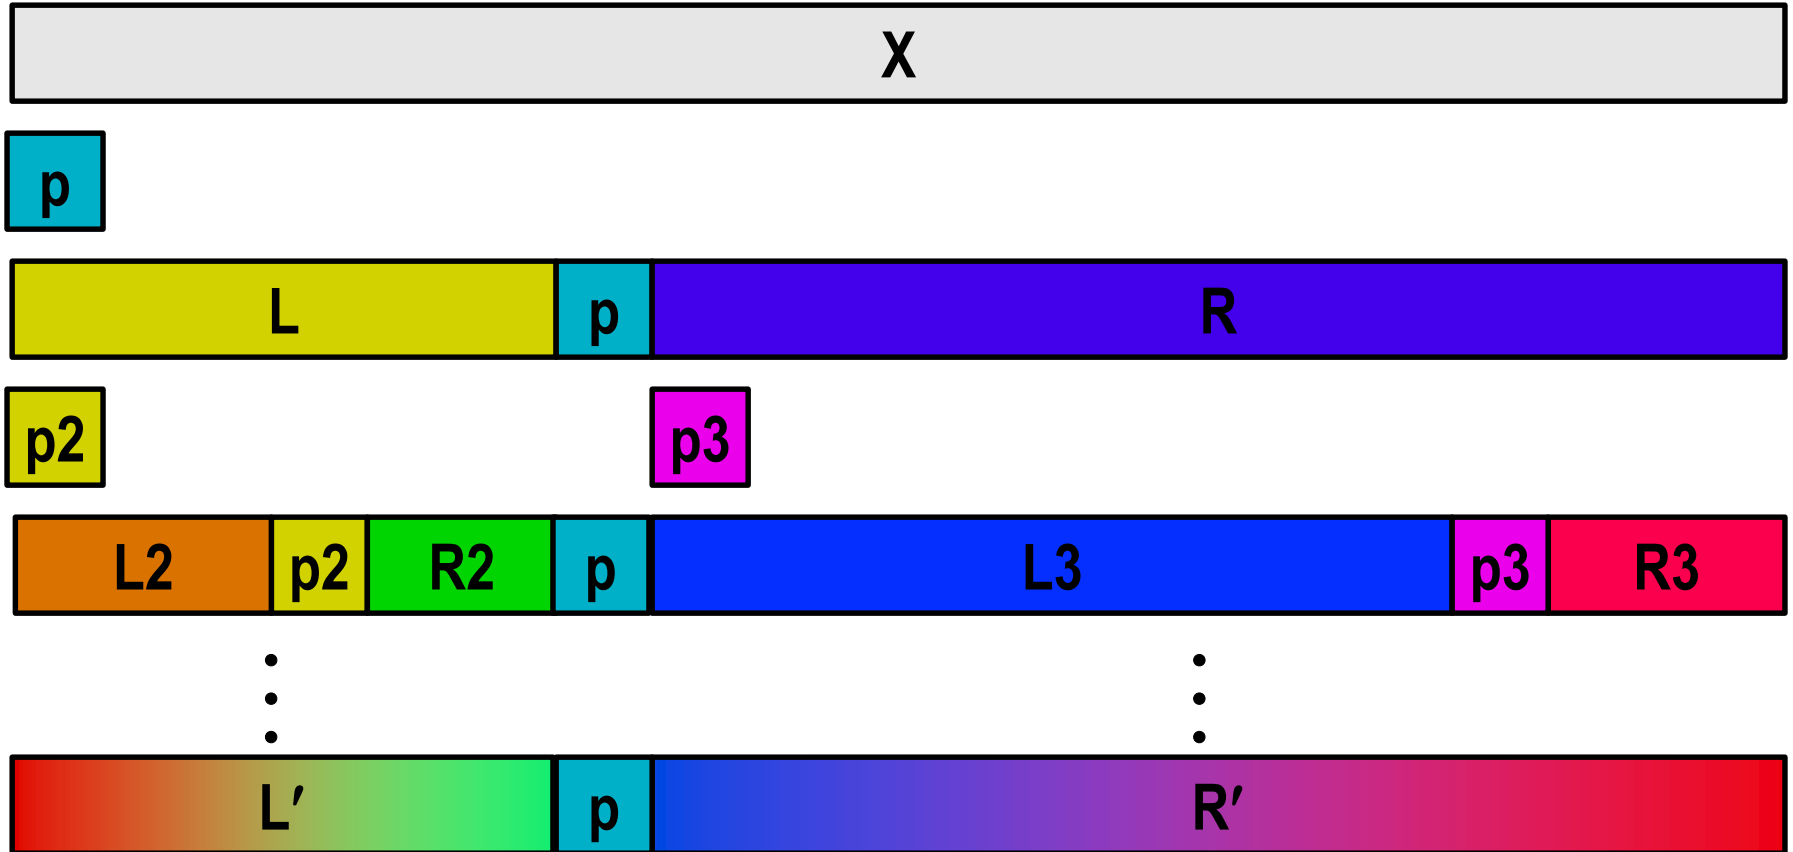
- 从base开始对nele个元素排序/Sort nele elements starting at base
 - 如果L或者R多于一个元素则递归排序/Recursively sort L or R if has more than one element



并行Quicksort/Parallel Quicksort

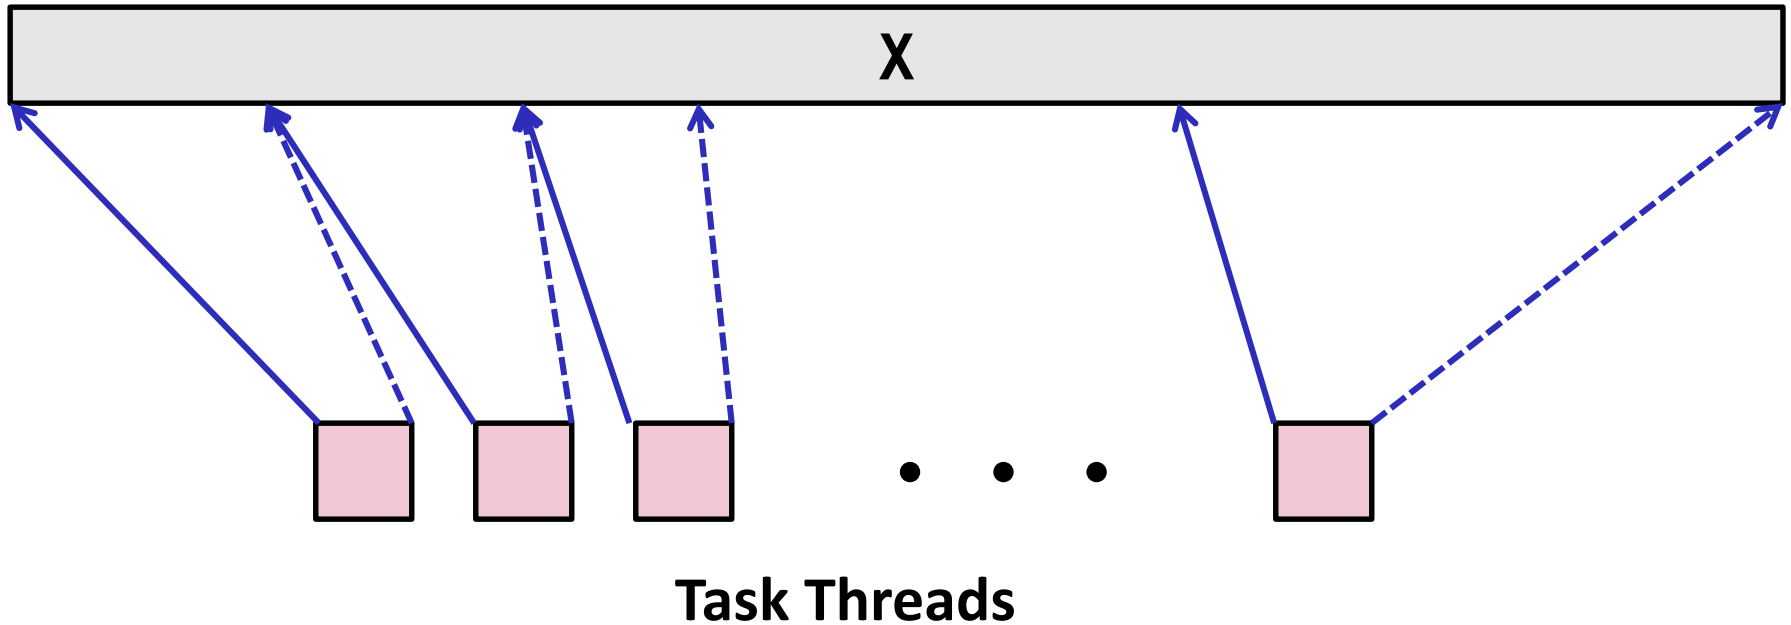
- 对数集X进行并行排序/Parallel quicksort of set of values X
 - 如果 $N \leq N_{\text{thresh}}$, 则进行串行排序/If $N \leq N_{\text{thresh}}$, do sequential quicksort
 - 否则/Else
 - 从X中选择枢轴p/Choose “pivot” p from X
 - 将X划分为/Rearrange X into
 - L: Values $\leq p$
 - R: Values $\geq p$
 - 递归创建独立线程进行排序/Recursively spawn separate threads
 - 对L进行排序形成L'/Sort L to get L'
 - 对R进行排序形成R'/Sort R to get R'
 - 返回/Return $L' : p : R'$

并行Quicksort 可视化/Parallel Quicksort Visualized





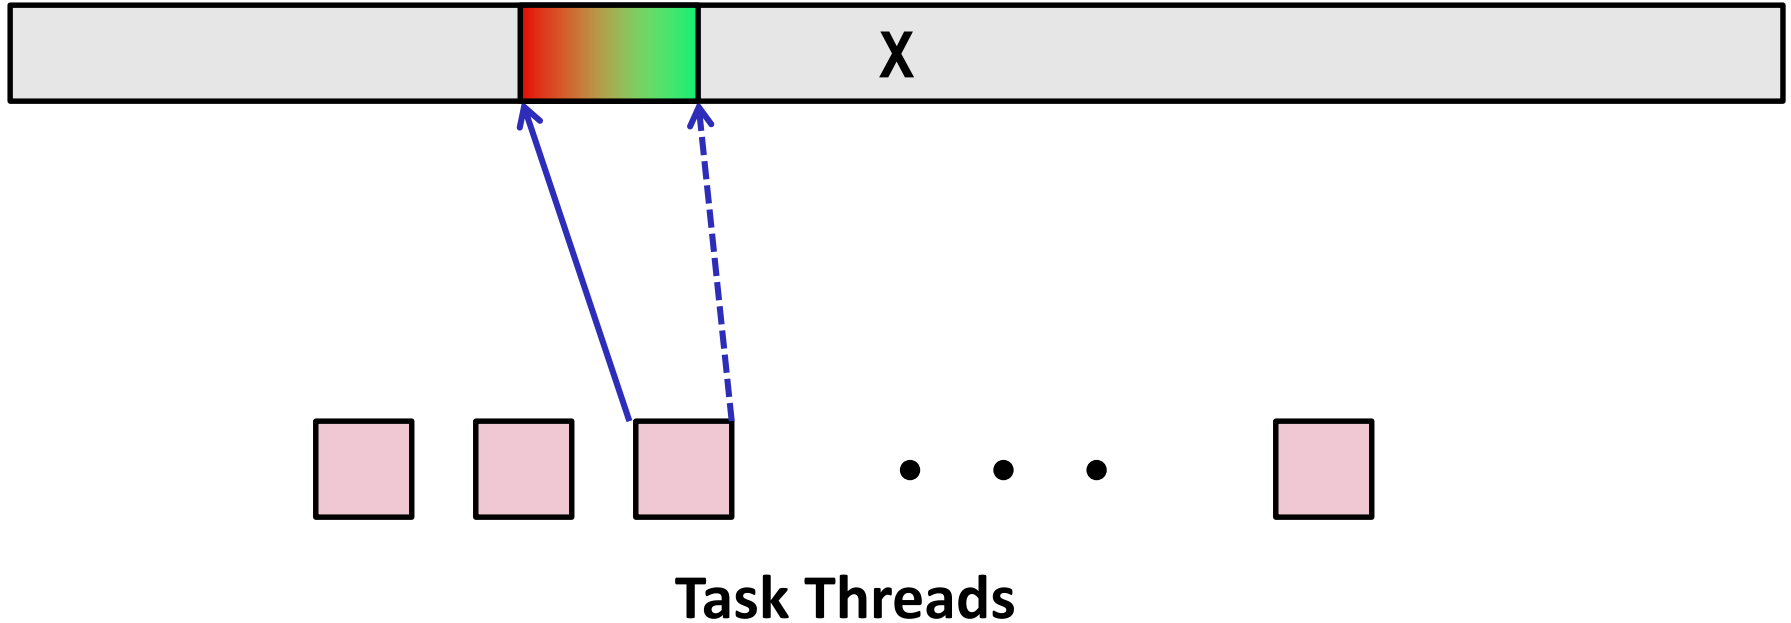
排序任务的线程结构/Thread Structure: Sorting Tasks



- 任务:对子区间进行排序/Task: Sort subrange of data
 - 描述为/Specify as:
 - **base**: Starting address/开始地址
 - **nele**: Number of elements in subrange/子区间元素数量
- 按照独立线程运行/Run as separate thread



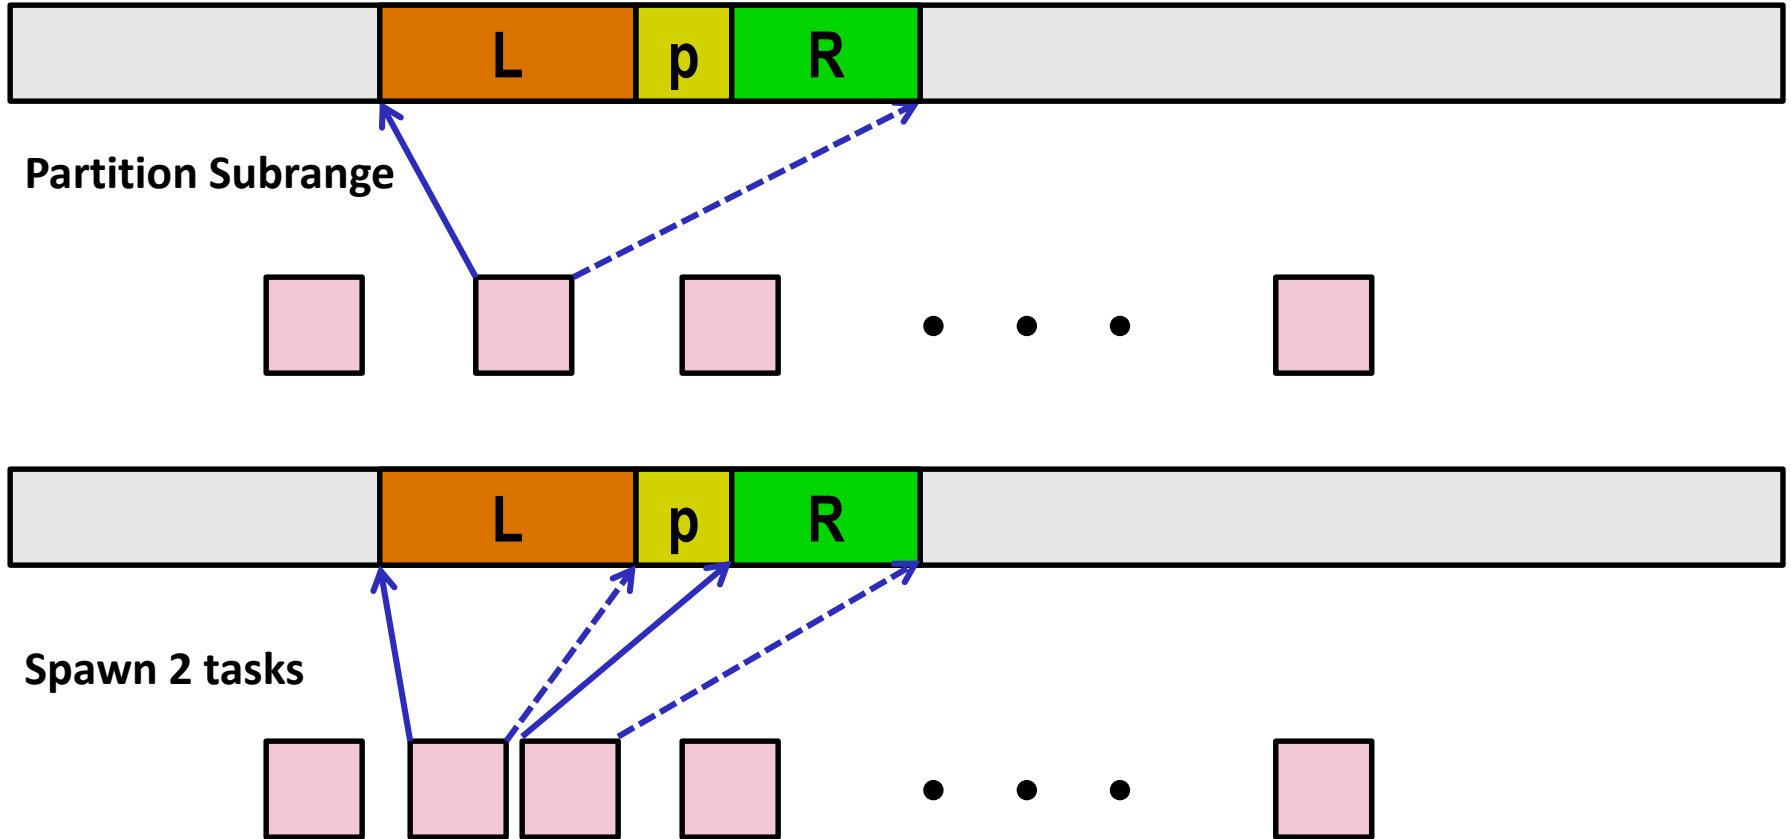
小规模排序任务运行/Small Sort Task Operation



- 使用串行quicksort 对子区间排序/Sort subrange using serial quicksort



大规模排序任务操作/ Large Sort Task Operation





顶层函数（简化版）/Top-Level Function (Simplified)

```
void tqsort(data_t *base, size_t nele) {  
    init_task(nele);  
    global_base = base;  
    global_end = global_base + nele - 1;  
    task_queue_ptr tq = new_task_queue();  
    tqsort_helper(base, nele, tq);  
    join_tasks(tq);  
    free_task_queue(tq);  
}
```

- 创建数据结构/Sets up data structures
- 递归调用排序函数/Calls recursive sort routine
- 持续对线程进行合并/Keeps joining threads until none left
- 释放数据结构Frees data structures



递归排序函数（简化版） Recursive sort routine (Simplified)

```
/* Multi-threaded quicksort */
static void tqsort_helper(data_t *base, size_t nele,
                           task_queue_ptr tq) {
    if (nele <= nele_max_sort_serial) {
        /* Use sequential sort */
        qsort_serial(base, nele);
        return;
    }
    sort_task_t *t = new_task(base, nele, tq);
    spawn_task(tq, sort_thread, (void *) t);
}
```

- 小区间：串行排序/Small partition: Sort serially
- 大区间：创建新的排序任务/Large partition: Spawn new sort task



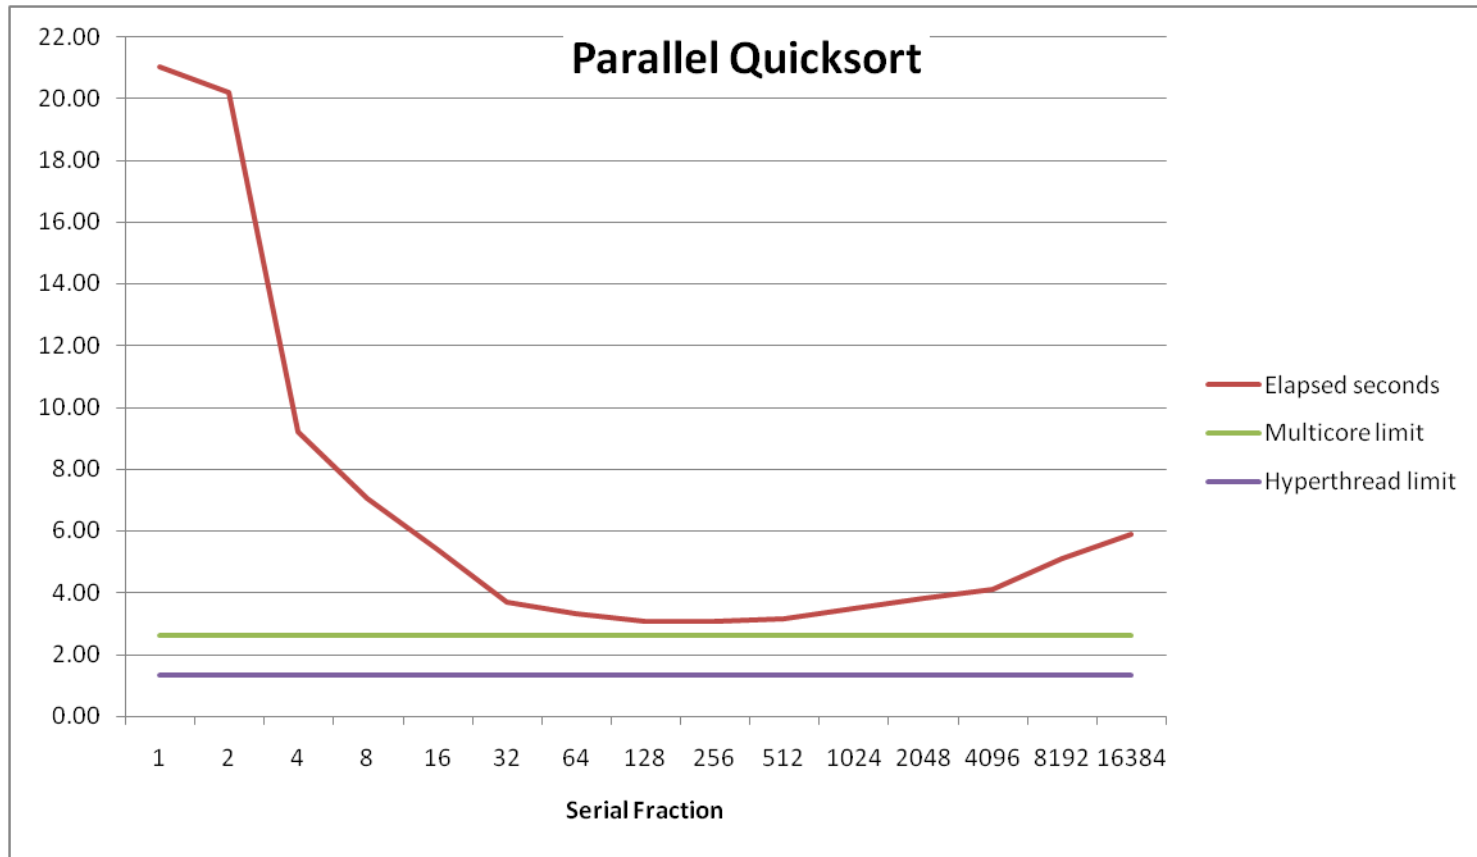
排序任务线程（简化版） Sort task thread (Simplified)

```
/* Thread routine for many-threaded quicksort */
static void *sort_thread(void *vargp) {
    sort_task_t *t = (sort_task_t *) vargp;
    data_t *base = t->base;
    size_t nele = t->nele;
    task_queue_ptr tq = t->tq;
    free(vargp);
    size_t m = partition(base, nele);
    if (m > 1)
        tqsort_helper(base, m, tq);
    if (nele-1 > m+1)
        tqsort_helper(base+m+1, nele-m-1, tq);
    return NULL;
}
```

- 获得任务参数/Get task parameters
- 进行划分/Perform partitioning step
- 对每个划分调用递归排序函数/Call recursive sort routine on each partition



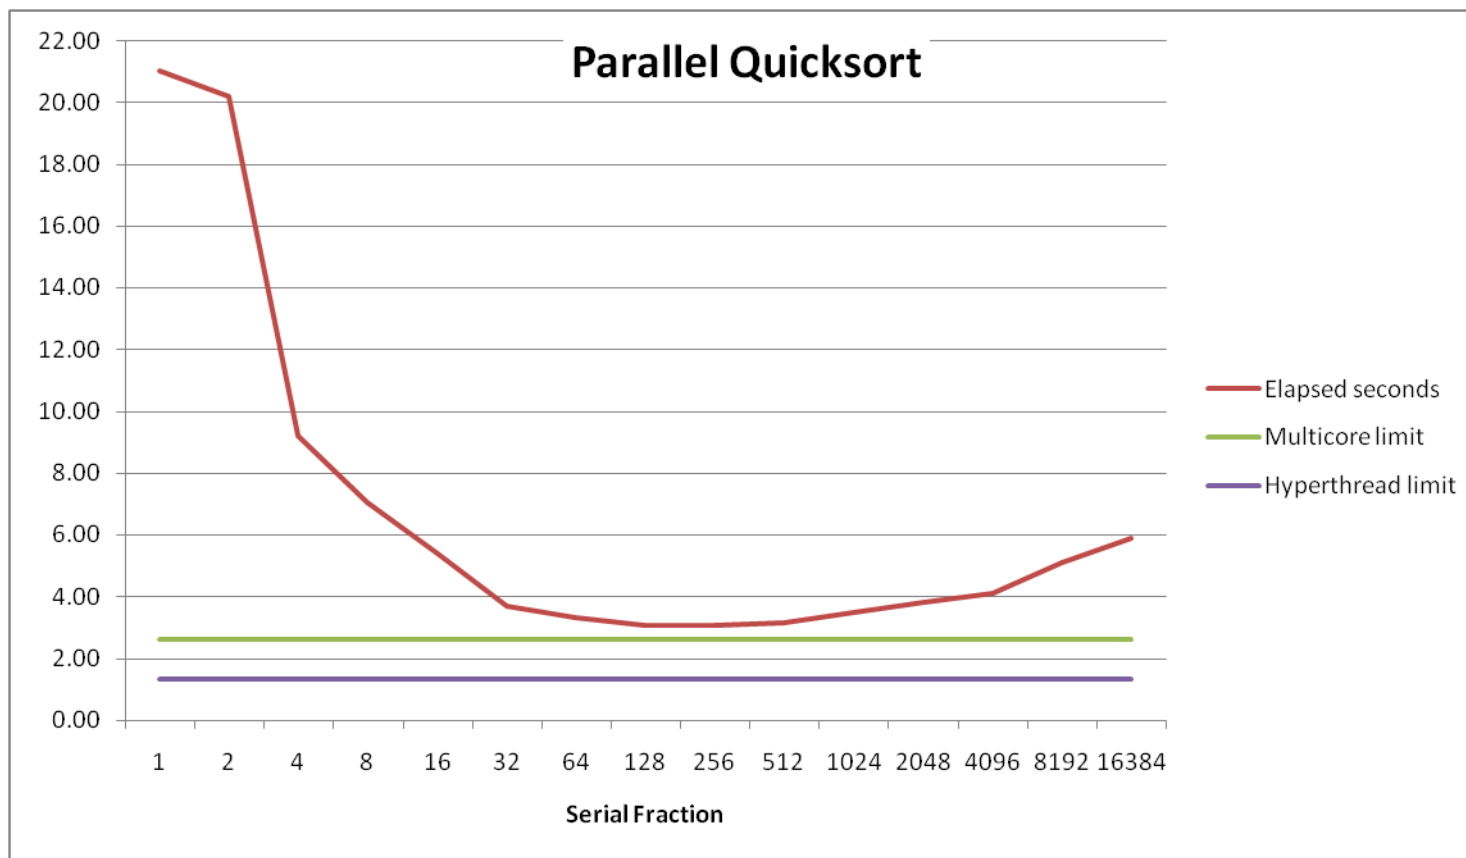
并行Quicksort性能/Parallel Quicksort Performance



- 串行占比:输入中串行排序的占比/Serial fraction: Fraction of input at which do serial sort
- 对 2^{27} 进行排序/Sort 2^{27} (134,217,728) random values
- 最好的加速比/Best speedup = 6.84X



并行Quicksort性能/ Parallel Quicksort Performance



- 对于大部分占比都有比较好的性能/Good performance over wide range of fraction values
 - F太小:并行度不够/F too small: Not enough parallelism
 - F太大:线程开销较大+线程栈空间不够/F too large: Thread overhead + run out of thread memory



阿姆达尔定律和并行Quicksort/Amdahl's Law & Parallel Quicksort

■ 串行瓶颈/Sequential bottleneck

- 顶层划分:无加速/Top-level partition: No speedup
- 第二层: $\leq 2X$ 加速比/Second level: $\leq 2X$ speedup
- 第k层: $\leq 2^{k-1}X$ 加速比/ k^{th} level: $\leq 2^{k-1}X$ speedup

■ 启发/Implications

- 小规模并行具有比较好的性能/Good performance for small-scale parallelism
- 需要对划分进行并行以实现更大规模的并行/Would need to parallelize partitioning step to get large-scale parallelism
 - 基于采样的并行排序/Parallel Sorting by Regular Sampling
 - H. Shi & J. Schaeffer, J. Parallel & Distributed Computing, 1992



划分并行/Parallelizing Partitioning Step



p

Parallel partitioning based on global p



Reassemble into partitions





并行划分的经验/Experience with Parallel Partitioning

- 无法获得加速比/Could not obtain speedup
- 原因分析：太多数据拷贝/Speculate: Too much data copying
 - 无法在原有数组内完成/Could not do everything within source array
 - 创建临时空间以重新整合划分/Set up temporary space for reassembling partition

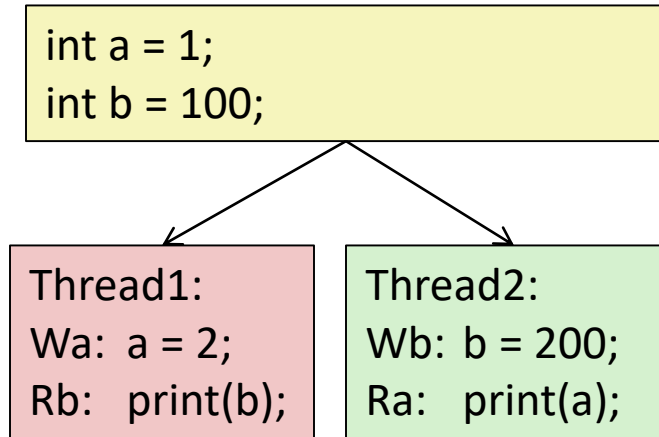


获得的教训/Lessons Learned

- **必须有并行化策略/Must have parallelization strategy**
 - 划分为K个独立的部分/Partition into K independent parts
 - 分治/Divide-and-conquer
- **内存循环不应该有同步/Inner loops must be synchronization free**
 - 同步操作开销过高/Synchronization operations very expensive
- **时刻记住阿尔达姆定律/Beware of Amdahl's Law**
 - 串行代码可能成为瓶颈/Serial code can become bottleneck
- **你可以的/You can do it!**
 - 实现一定程度的并行并不困难/Achieving modest levels of parallelism is not difficult
 - 构建实验框架并测试不同的策略/Set up experimental framework and test multiple strategies



内存顺序一致性模型/Memory Consistency



线程一致性限制/Thread consistency constraints

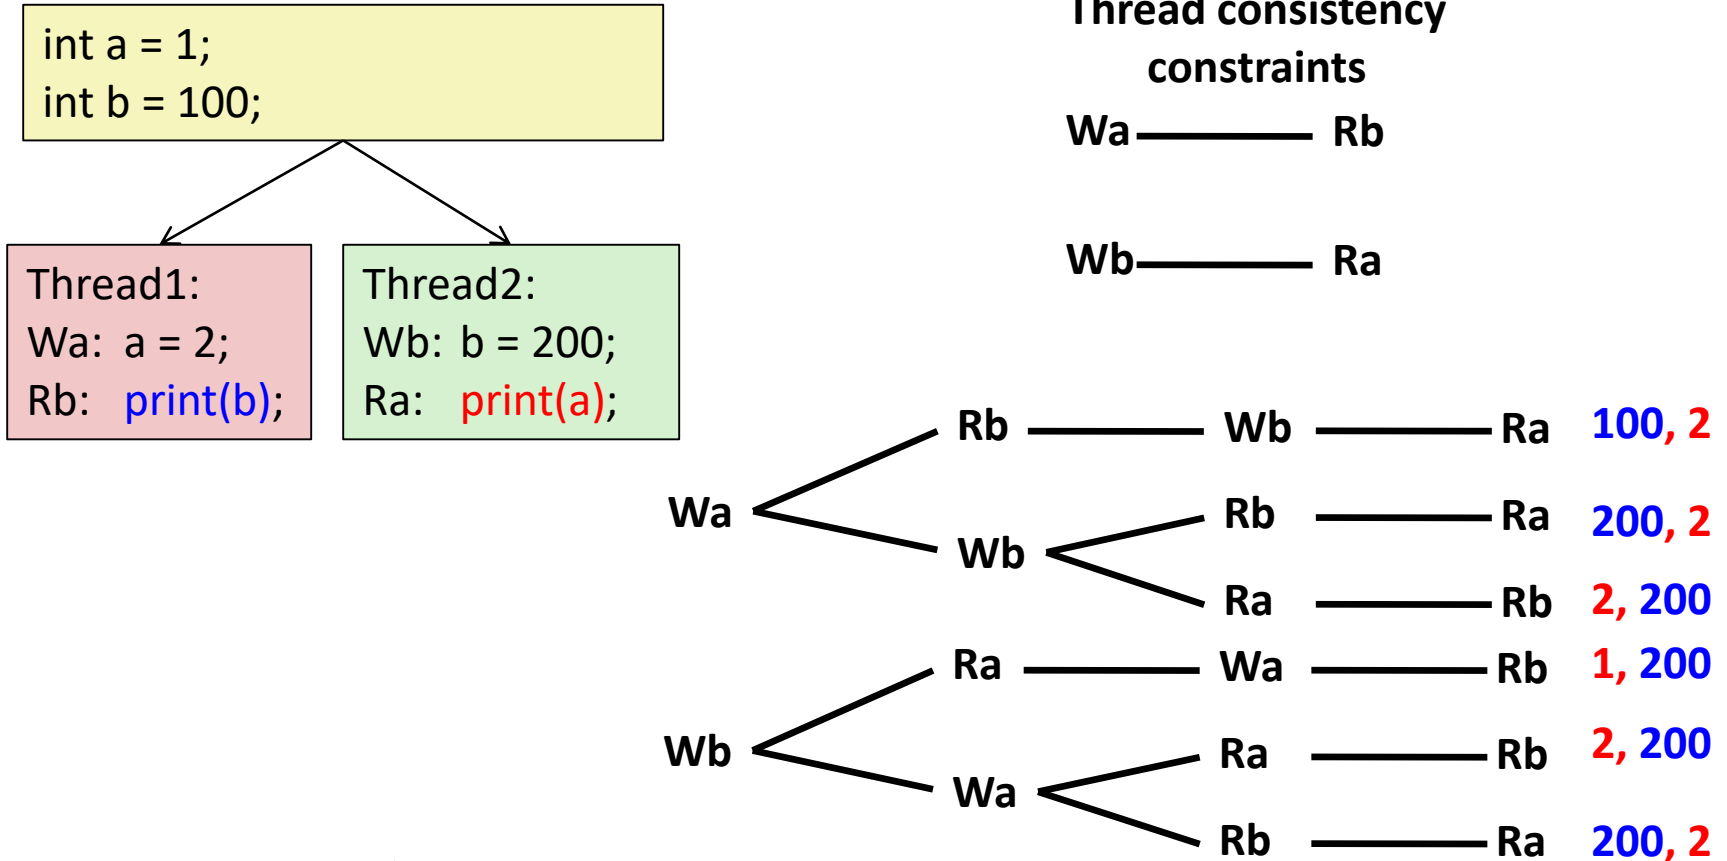
Wa → Rb

Wb → Ra

- 打印出来的值可能有哪些? /What are the possible values printed?
 - 依赖于内存顺序一致性模型/Depends on memory consistency model
 - 是硬件如何实现并发访问的抽象模型/Abstract model of how hardware handles concurrent accesses
- 串行一致性模型/Sequential consistency
 - 线程内满足程序序/Overall effect consistent with each individual thread
 - 其他任意交叉/Otherwise, arbitrary interleaving



串行顺序一致性举例/Sequential Consistency Example



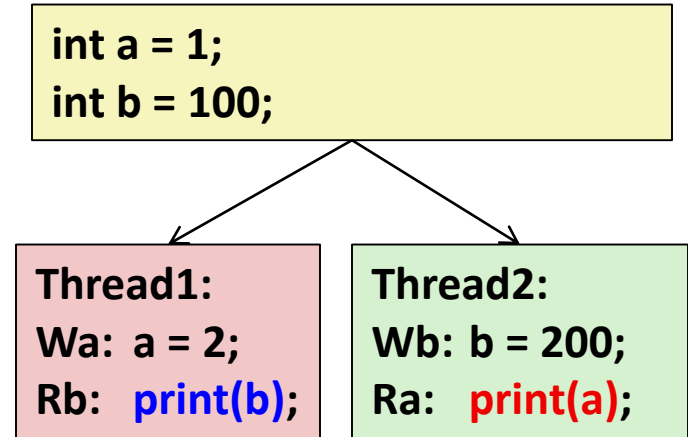
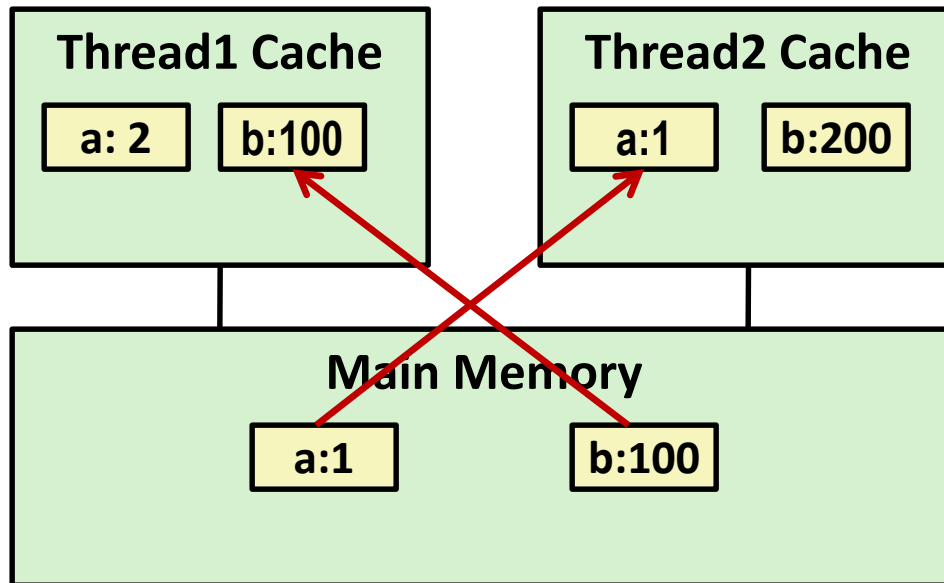
■ 不可能的输出/Impossible outputs

- 100, 1 and 1, 100
- 需要在Wa和Wb之前到达Ra和Rb/Would require reaching both Ra and Rb before Wa and Wb



没有Cache一致性协议时/Non-Coherent Cache Scenario

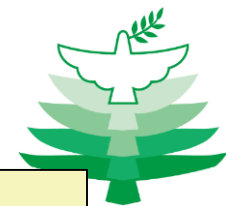
- Write-back caches, without coordination between them



print 1

print 100

总线侦听Cache一致性协议/Snoopy Caches



- 对每个Cache块打标签/Tag each cache block with state
block with state

非法/Invalid 不能用/Cannot use value

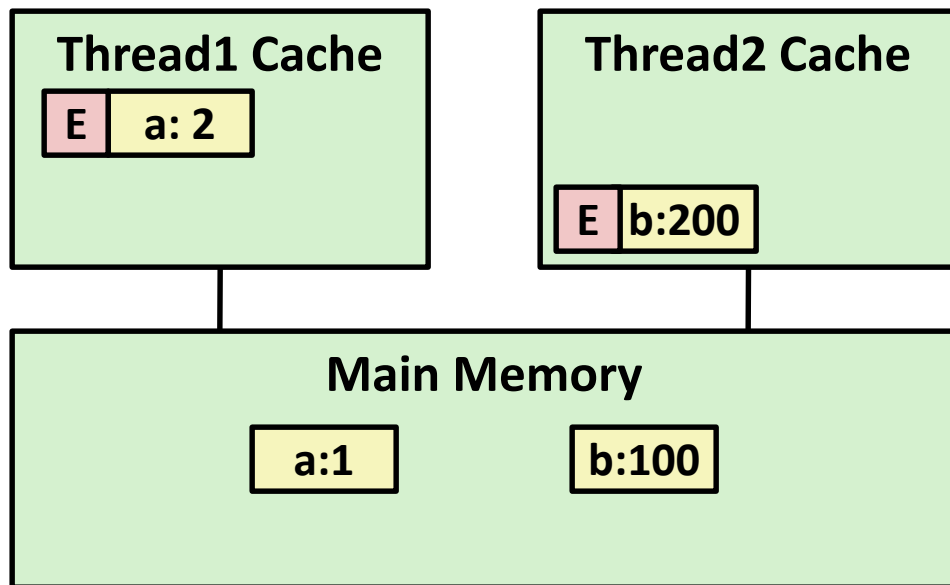
共享/Shared 可读副本/Readable copy

独占/Exclusive 可写副本/Writable copy

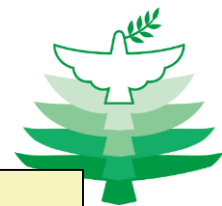
```
int a = 1;  
int b = 100;
```

Thread1:
Wa: a = 2;
Rb: print(b);

Thread2:
Wb: b = 200;
Ra: print(a);



总线侦听Cache一致性协议/ Snoopy Caches



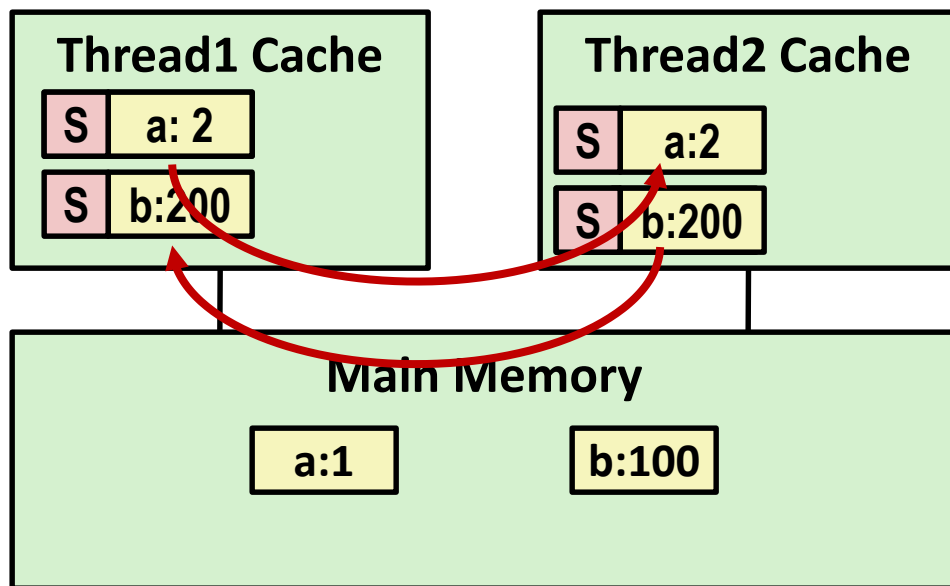
■ Tag each cache block with state

Invalid	Cannot use value
Shared	Readable copy
Exclusive	Writeable copy

```
int a = 1;  
int b = 100;
```

Thread1:
Wa: a = 2;
Rb: **print(b);**

Thread2:
Wb: b = 200;
Ra: **print(a);**



print 2

print 200

- 当cache看到某个对标记位E的块
的请求/When cache sees request
for one of its E-tagged blocks
 - 从自己的Cache提供数据
/Supply value from cache
 - 将标记改为S/Set tag to S